



CSCI-GA.3033-004


# Graphics Processing Units (GPUs): Architecture and Programming

## Lecture 3: CUDA Programming Model

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



# Behind CUDA



# Parallel Computing on a GPU

- GPUs deliver 25 to 200+ GFLOPS on compiled parallel C applications
  - Available in laptops, desktops, and clusters
- GPU parallelism is doubling every year
- Programming model scales transparently
  - Data parallelism
- Programmable in C with CUDA tools
- Multithreaded SPMD model uses application data parallelism and thread parallelism.
  - [SPMD = Single Program Multiple Data]



GeForce 8800



Tesla D870



Tesla S870


# CUDA

- Compute Unified Device Architecture
- Integrated host+device app C program
  - Serial or modestly parallel parts in **host** C code
  - Highly parallel parts in **device** SPMD kernel C code

Serial Code (host)

Parallel **Kernel** (device)


KernelA<<< nBlk, nTid >>>(args);



Serial Code (host)


Parallel Kernel (device)

KernelB<<< nBlk, nTid >>>(args);



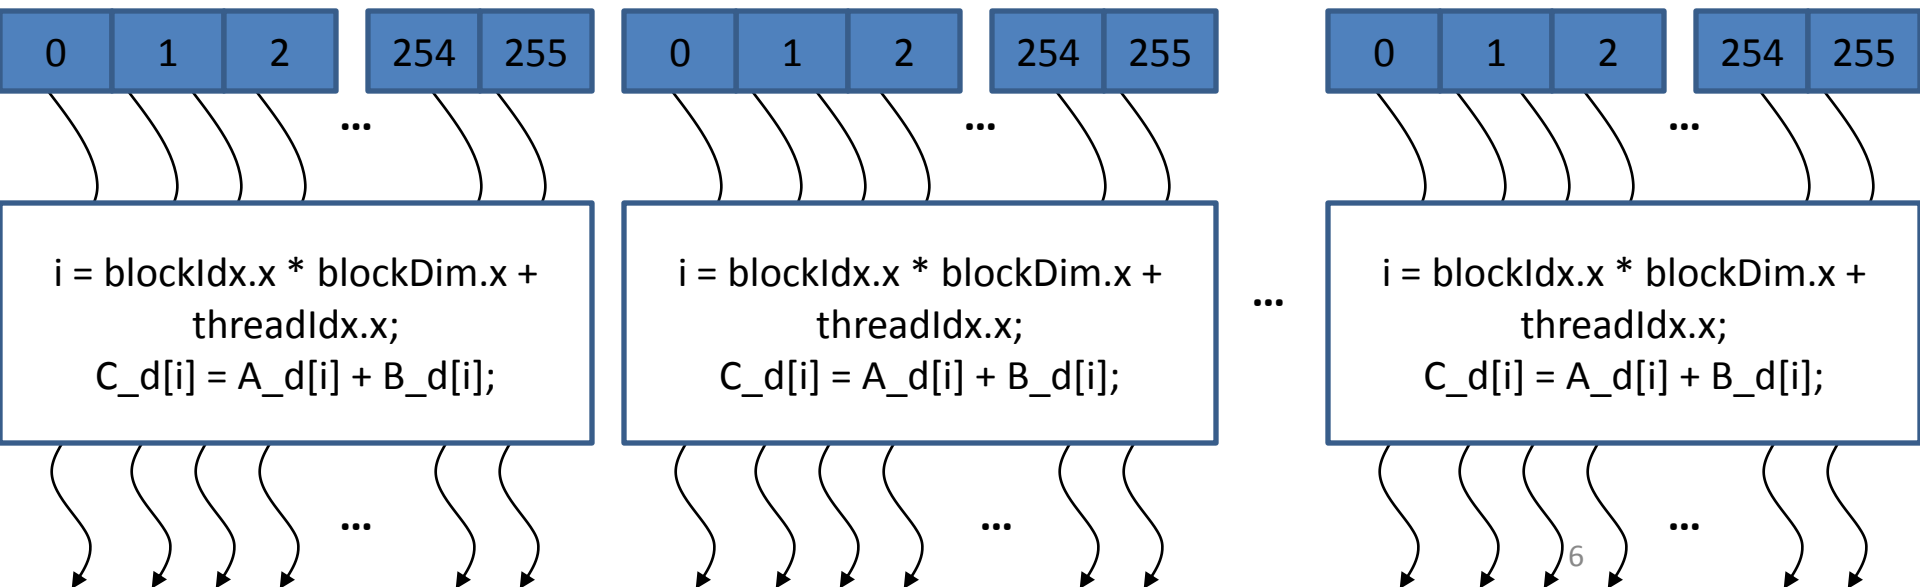
# Parallel Threads

- A CUDA kernel is executed by an array of threads
  - All threads run the same code (the SP in SPMD)
  - Each thread has an ID that it uses to compute memory addresses and make control decisions



# Thread Blocks

- Divide monolithic thread array into multiple blocks
  - Threads within a block cooperate via **shared memory, atomic operations** and **barrier synchronization, ...**
  - Threads in different blocks cannot cooperate



# Kernel

- Launched by the host
- Very similar to a C function
- To be executed on device
- All threads will execute that same code in the kernel.



# Grid

- 1D or 2D (or 3D) organization of a block
- `blockDim.x` and `blockDim.y`
- `gridDim.x` and `gridDim.y`



# Block

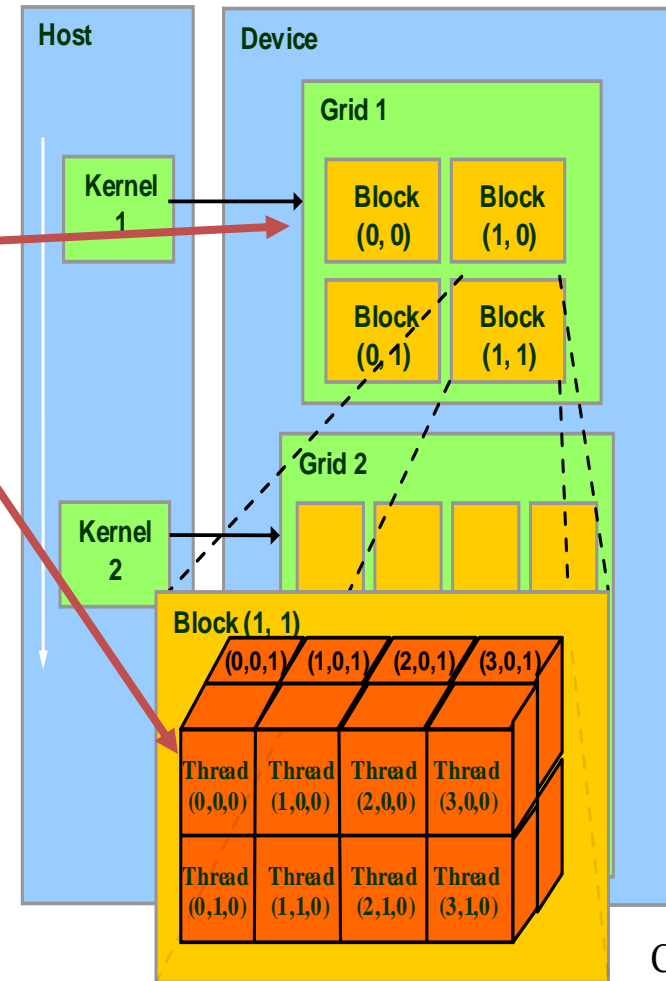
- 1D, 2D, or 3D organization of a block
- Block is assigned to an SM
- `blockIdx.x`, `blockIdx.y`, and `blockIdx.z`



# Thread

# IDs

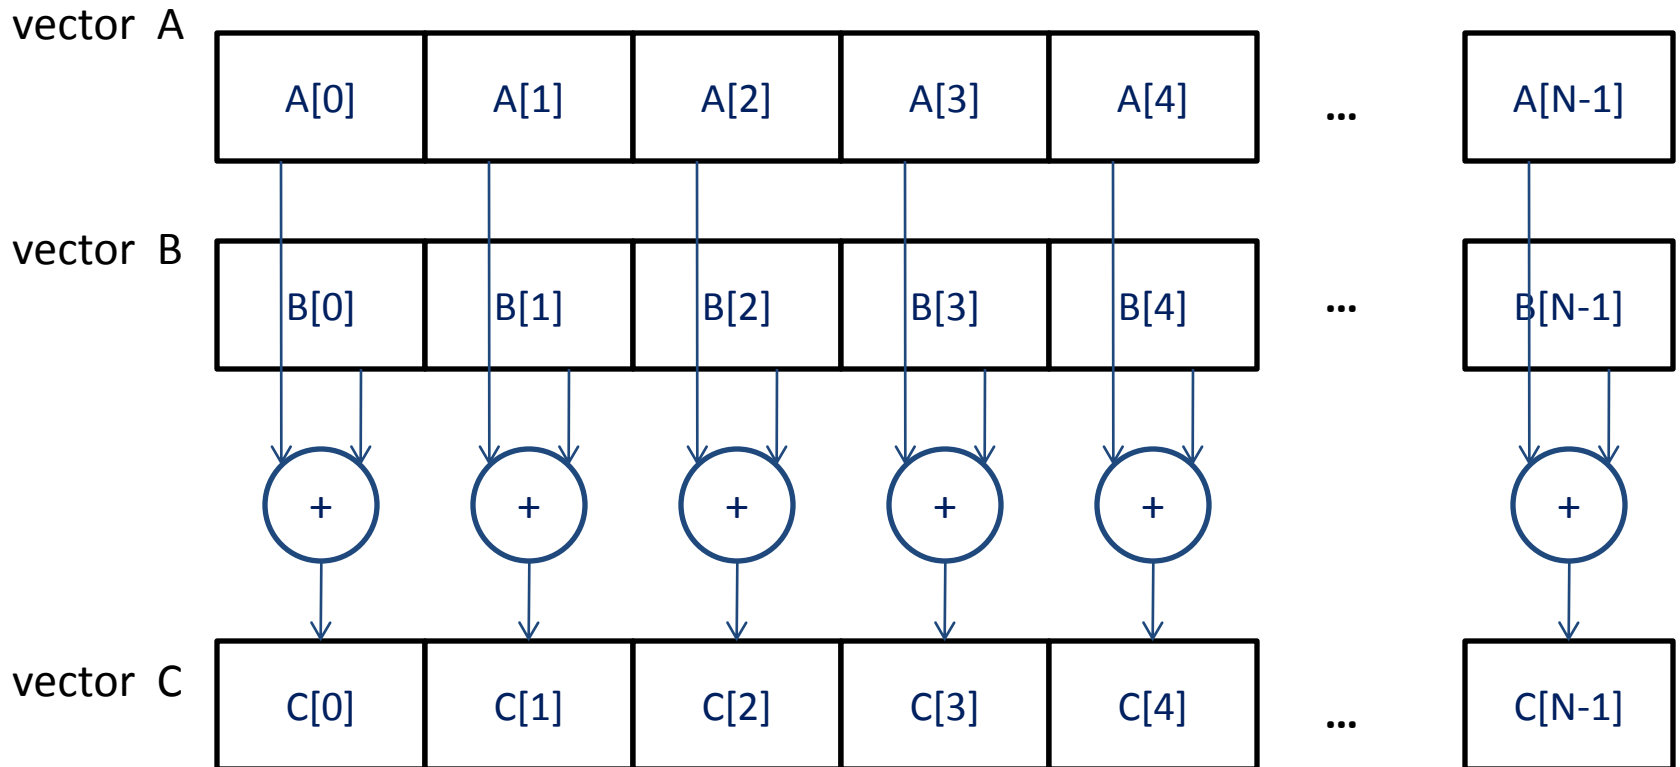
- Each thread uses IDs to decide what data to work on
  - Block ID: 1D or 2D (or 3D)
  - Thread ID: 1D, 2D, or 3D
- Simplifies memory addressing when processing multidimensional data
  - Image processing
  - Solving PDEs on volumes
  - ...



Courtesy: NDVIA



# A Simple Example: Vector Addition



# A Simple Example: Vector Addition


```
// Compute vector sum C = A+B
void vecAdd(float* A, float* B, float* C, int n)
{
    for (i = 0, i < n, i++)
        C[i] = A[i] + B[i];
}
```

GPU friendly!

```
int main()
{
    // Memory allocation for A_h, B_h, and C_h
    // I/O to read A_h and B_h, N elements
    ...
    vecAdd(A_h, B_h, C_h, N);
}
```

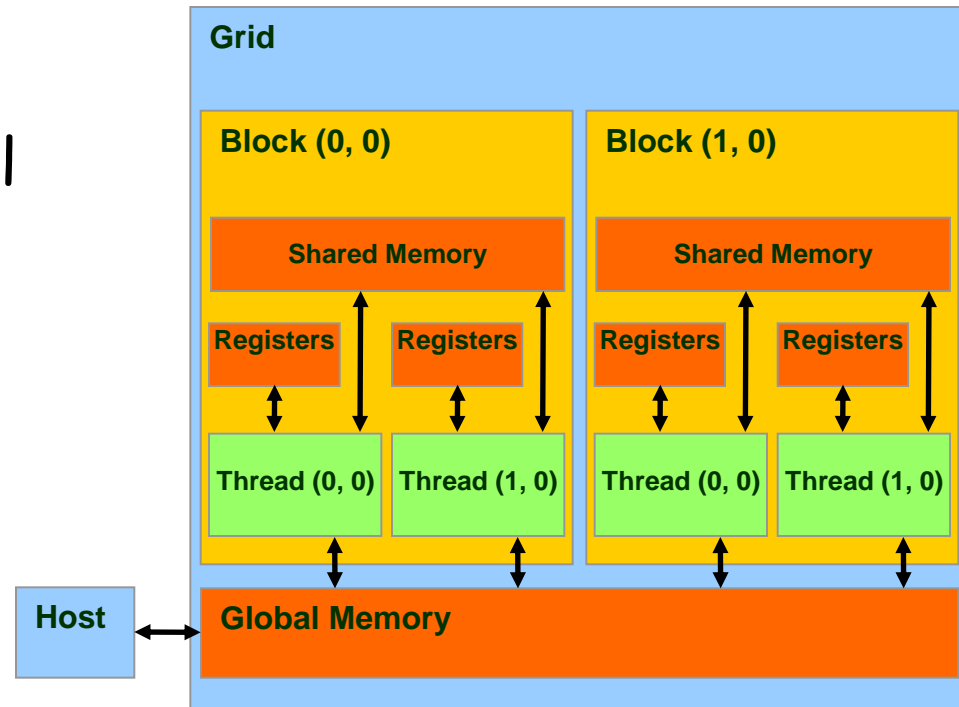
# A Simple Example: Vector Addition

```
#include <cuda.h>  
void vecAdd(float* A, float* B, float* C, int n)  
{  
  int size = n* sizeof(float);  
  float* A_d, B_d, C_d;  
  ...  
1. // Allocate device memory for A, B, and C  
  // copy A and B to device memory  
  
2. // Kernel launch code – to have the device  
  // to perform the actual vector addition  
  
3. // copy C from the device memory  
  // Free device vectors  
}
```



# CUDA Memory Model

- Global memory
  - Main means of communicating R/W Data between **host** and **device**
  - Contents visible to all threads
  - Long latency access
- Device code can:
  - R/W per-thread registers
  - R/W per-grid global memory
- We will cover more later




# CPU & GPU Memory

- In CUDA, host and devices have separate memory spaces.
- If GPU and CPU are on the same chip, then they share memory space → fusion

# CUDA Device Memory Allocation

- **cudaMalloc()**
  - Allocates object in the device Global Memory
  - Requires two parameters
    - Address of a pointer to the allocated object
    - Size of allocated object
- **cudaFree()**
  - Frees object from device Global Memory
    - Pointer to freed object

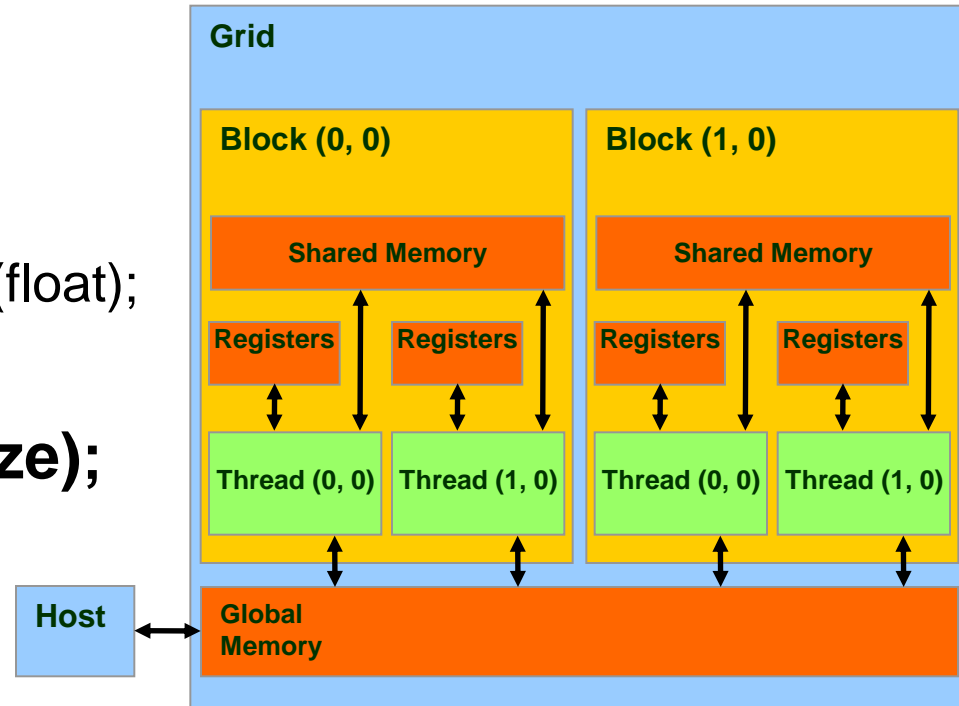


# CUDA Device Memory Allocation

## Example:

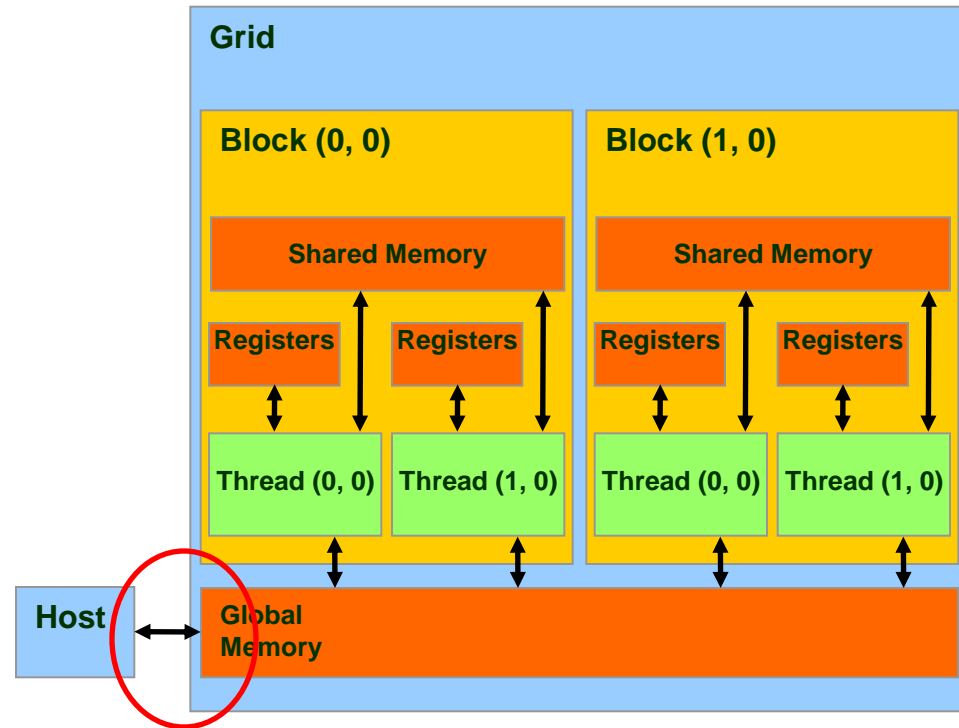
```
WIDTH = 64;  
float * Md;  
int size = WIDTH * WIDTH * sizeof(float);
```

```
cudaMalloc((void**)&Md, size);  
cudaFree(Md);
```



# CUDA Device Memory Allocation

- `cudaMemcpy()`
  - memory data transfer
  - Requires four parameters
    - Pointer to destination
    - Pointer to source
    - Number of bytes copied
    - Type of transfer
      - Host to Host
      - Host to Device
      - Device to Host
      - Device to Device
- Asynchronous transfer



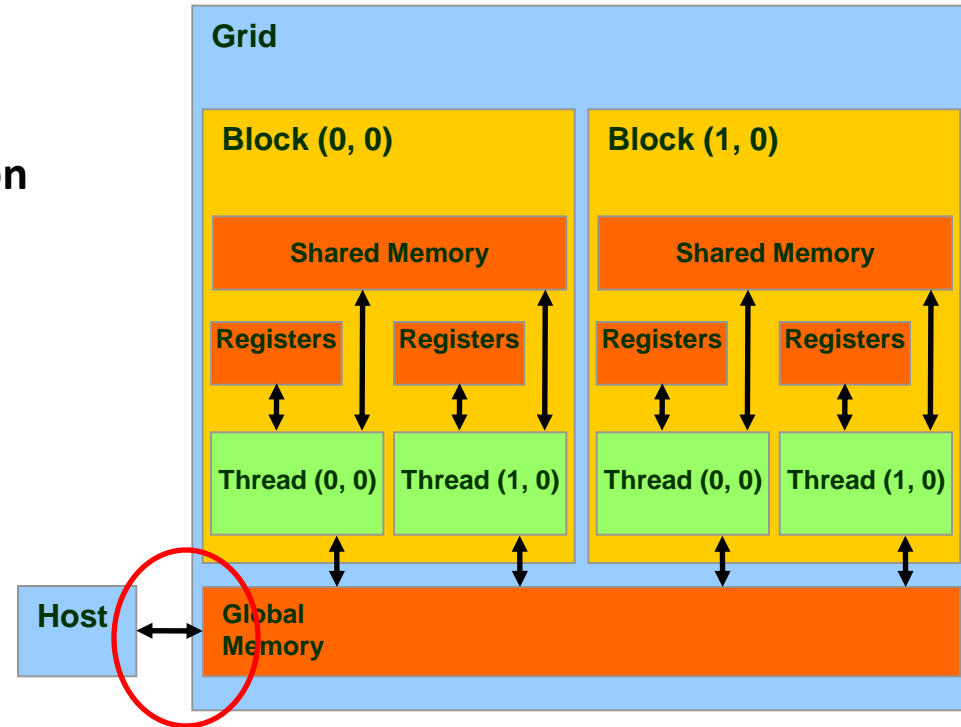

**Important!**

`cudaMemcpy()` **cannot** be used to copy between different GPUs in multi-GPUs system



# CUDA Device Memory Allocation

Example:



```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(M, Md, size, cudaMemcpyDeviceToHost);
```

# A Simple Example: Vector Addition

```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float* A_d, * B_d, * C_d;
```

1. // Transfer A and B to device memory

```
    cudaMalloc((void **) &A_d, size);
    cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &B_d, size);
    cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);
```

```
    // Allocate device memory for C_d
    cudaMalloc((void **) &C_d, size);
```

2. // Kernel invocation code – to be shown later

How to launch a kernel?

...

3. // Transfer C from device to host

```
    cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(A_d); cudaFree(B_d); cudaFree (C_d);
```

```
}
```

```
int vecAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<ceil(n/256),256>>>(A_d, B_d, C_d, n);
}
```

#blocks                      #threads/blks

// Each thread performs one pair-wise addition

```
__global__
void vecAddkernel(float* A_d, float* B_d, float* C_d, int n)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if(i<n) C_d[i] = A_d[i] + B_d[i];
}
```

Unique ID

# Unique ID

## 1D grid of 1D blocks

$\text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$

# Unique ID

## 1D grid of 2D blocks

```
blockIdx.x * blockDim.x * blockDim.y +  
threadIdx.y * blockDim.x +  
threadIdx.x;
```

# Unique ID

## 1D grid of 3D blocks

```
blockIdx.x * blockDim.x * blockDim.y *  
blockDim.z +  
threadIdx.z * blockDim.y * blockDim.x +  
threadIdx.y * blockDim.x +  
threadIdx.x;
```

# Unique ID

## 2D grid of 1D blocks

```
int blockId = blockIdx.y * blockDim.x +  
blockIdx.x;
```

```
int threadId = blockId * blockDim.x +  
threadIdx.x;
```

# Unique ID

## 2D grid of 2D blocks

```
int blockId = blockIdx.x + blockIdx.y *  
gridDim.x;
```

```
int threadId = blockId * (blockDim.x *  
blockDim.y) +  
(threadIdx.y * blockDim.x) +  
threadIdx.x;
```



# Unique ID

## 2D grid of 3D blocks

```
int blockId = blockIdx.x +  
              blockIdx.y * gridDim.x;
```

```
int threadId = blockId * (blockDim.x *  
blockDim.y * blockDim.z) +  
(threadIdx.z * (blockDim.x * blockDim.y))  
+ (threadIdx.y * blockDim.x)  
+ threadIdx.x;
```

# Unique ID

## 3D grid of 1D blocks

```
int blockId = blockIdx.x  
    + blockIdx.y * gridDim.x  
    + gridDim.x * gridDim.y * blockIdx.z;
```

```
int threadId = blockId * blockDim.x +  
    threadIdx.x;
```

# Unique ID

## 3D grid of 2D blocks

```
int blockId = blockIdx.x  
            + blockIdx.y * gridDim.x  
            + gridDim.x * gridDim.y * blockIdx.z;
```

```
int threadId = blockId * (blockDim.x *  
                        blockDim.y)  
              + (threadIdx.y * blockDim.x)  
              + threadIdx.x;
```

# Unique ID

## 3D grid of 3D blocks

```
int blockId = blockIdx.x  
    + blockIdx.y * gridDim.x  
    + gridDim.x * gridDim.y * blockIdx.z;
```

```
int threadId = blockId * (blockDim.x *  
    blockDim.y * blockDim.z) +  
    (threadIdx.z * (blockDim.x * blockDim.y))  
    + (threadIdx.y * blockDim.x)  
    + threadIdx.x;
```

# The *Hello World* of Parallel Programming: Matrix Multiplication

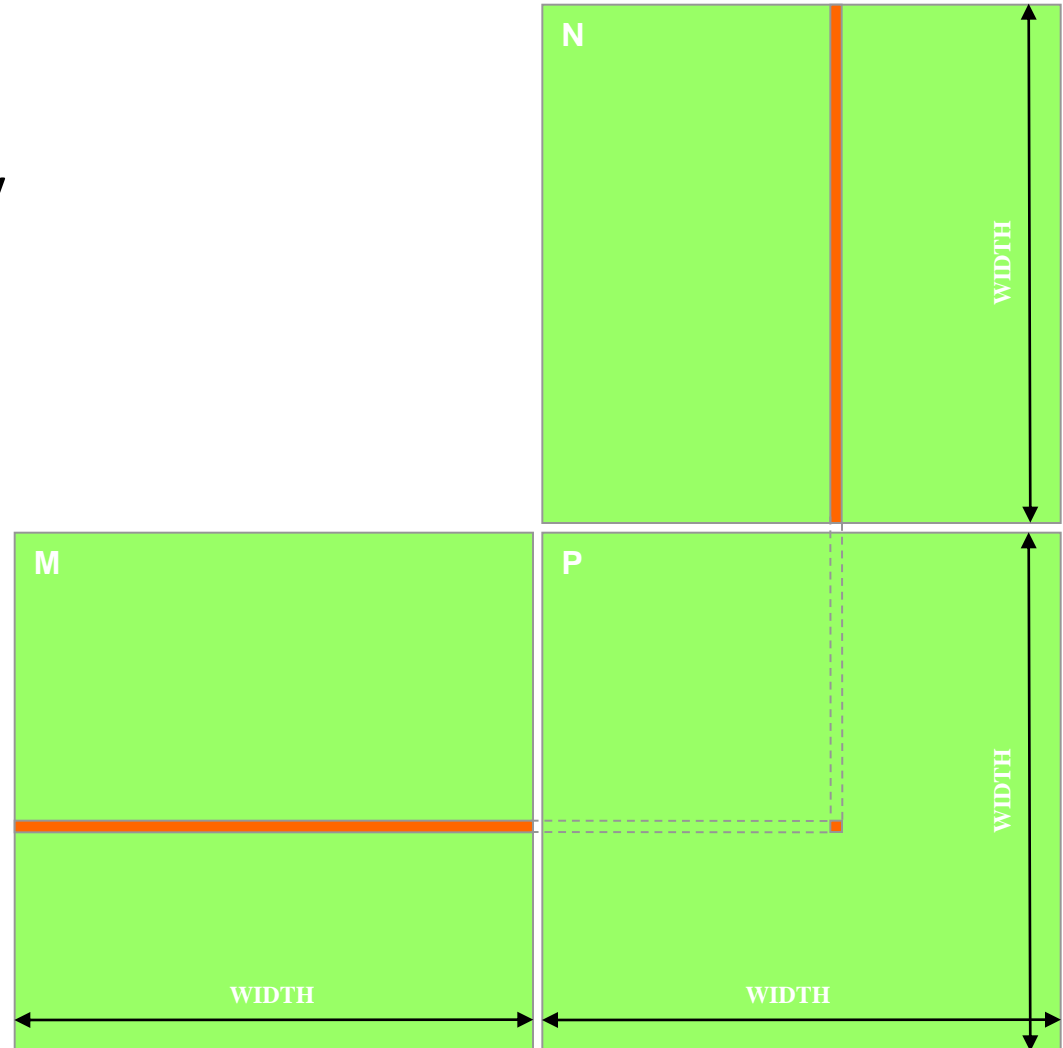
	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

- `__global__` defines a kernel function. Must return `void`
- `__device__` and `__host__` can be used together
- For functions executed on the device:
  - No recursion
  - No static variable declarations inside the function
  - No indirect function calls through pointers

# The *Hello World* of Parallel Programming: Matrix Multiplication

## Data Parallelism:

We can safely perform many arithmetic **operations on** the data structures in a **simultaneous** manner.



# The *Hello World* of Parallel Programming: Matrix Multiplication

$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$
$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$
$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$
$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$

M



$M_{0,0}$	$M_{1,0}$	$M_{2,0}$	$M_{3,0}$	$M_{0,1}$	$M_{1,1}$	$M_{2,1}$	$M_{3,1}$	$M_{0,2}$	$M_{1,2}$	$M_{2,2}$	$M_{3,2}$	$M_{0,3}$	$M_{1,3}$	$M_{2,3}$	$M_{3,3}$
-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------

C adopts row-major placement approach  
when storing 2D matrix in linear memory address.

# The *Hello World* of Parallel Programming: Matrix Multiplication

```
int main(void) {  
1.  // Allocate and initialize the matrices M, N, P  
    // I/O to read the input matrices M and N  
    ....  
  
2.  // M * N on the device  
    MatrixMultiplication(M, N, P, Width);  
  
3.  // I/O to write the output matrix P  
    // Free matrices M, N, P  
    ...  
    return 0;  
}
```

**A Simple main function: executed at the host**



# The *Hello World* of Parallel Programming: Matrix Multiplication

// Matrix multiplication on the (CPU) host

```
void MatrixMulOnHost(float* M, float* N, float* P, int Width)
```

```
{
```

```
    for (int i = 0; i < Width; ++i)
```

```
        for (int j = 0; j < Width; ++j) {
```

```
            double sum = 0;
```

```
            for (int k = 0; k < Width; ++k) {
```

```
                double a = M[i * Width + k];
```

```
                double b = N[k * Width + j];
```

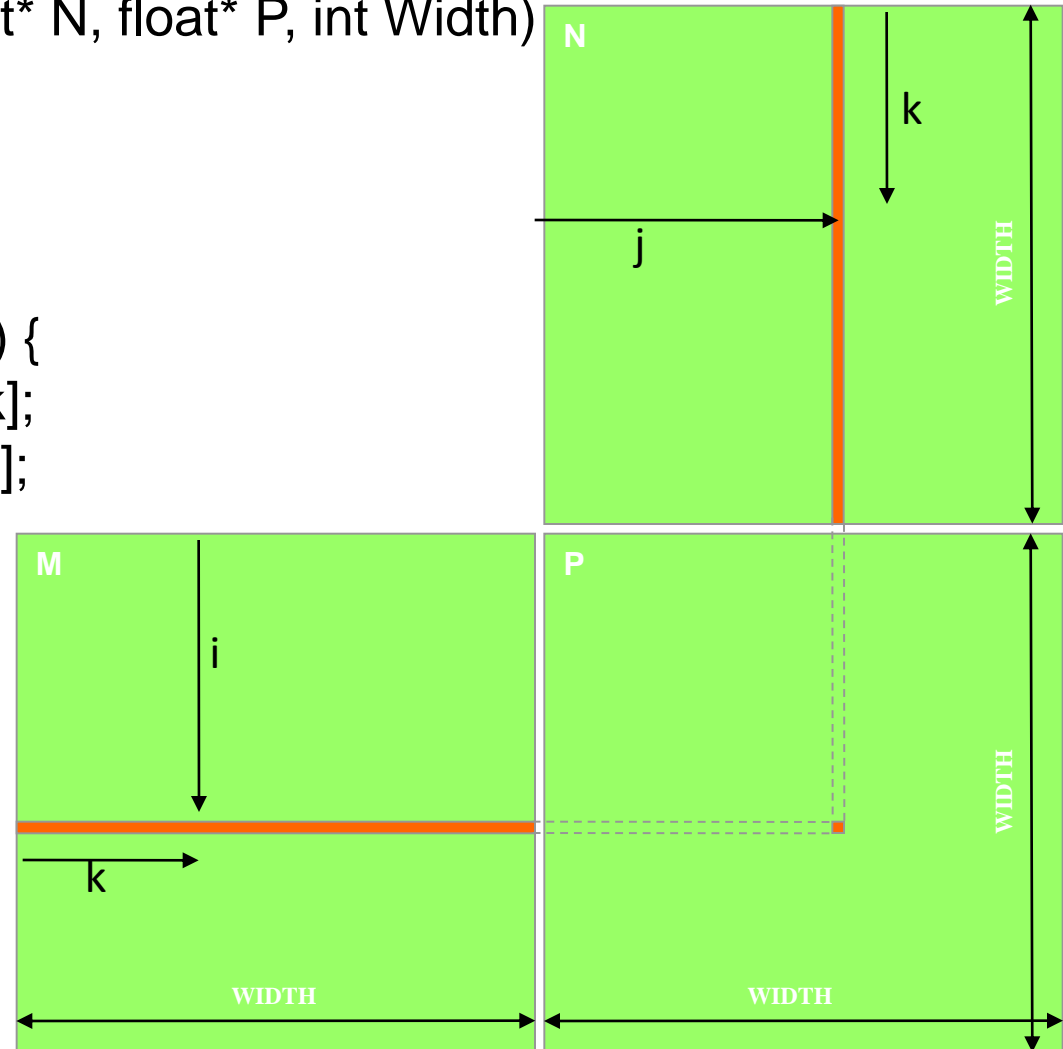
```
                sum += a * b;
```

```
            }
```

```
            P[i * Width + j] = sum;
```

```
        }
```

```
}
```



# The *Hello World* of Parallel Programming: Matrix Multiplication

```
void MatrixMultiplication(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);
    float* Md, Nd, Pd;

    1. // Transfer M and N to device memory
    cudaMalloc((void**) &Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**) &Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc((void**) &Pd, size);

    MatrixMulKernel(Md, Nd, Pd, Width);
    ...
    3. // Transfer P from device to host
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree (Pd);
}
```

# The *Hello World* of Parallel Programming: Matrix Multiplication

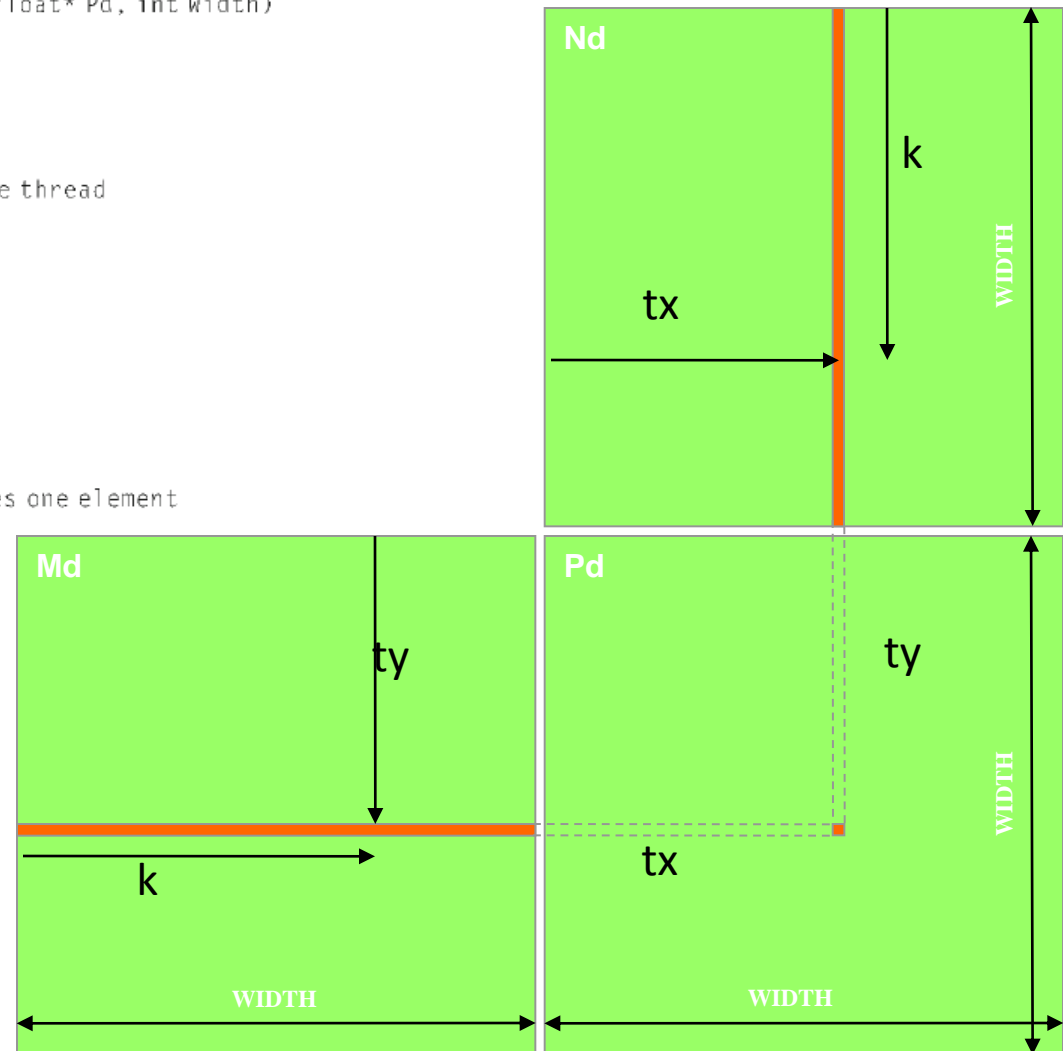
```
// Matrix multiplication kernel - thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Width + k];
        float Ndelement = Nd[k * Width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

## The Kernel Function



# More On Specifying Dimensions

```
// Setup the execution configuration
```

```
dim3 dimGrid(x, y);
```

```
dim3 dimBlock(x, y, z);
```

```
// Launch the device computation threads!
```

```
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

Important:

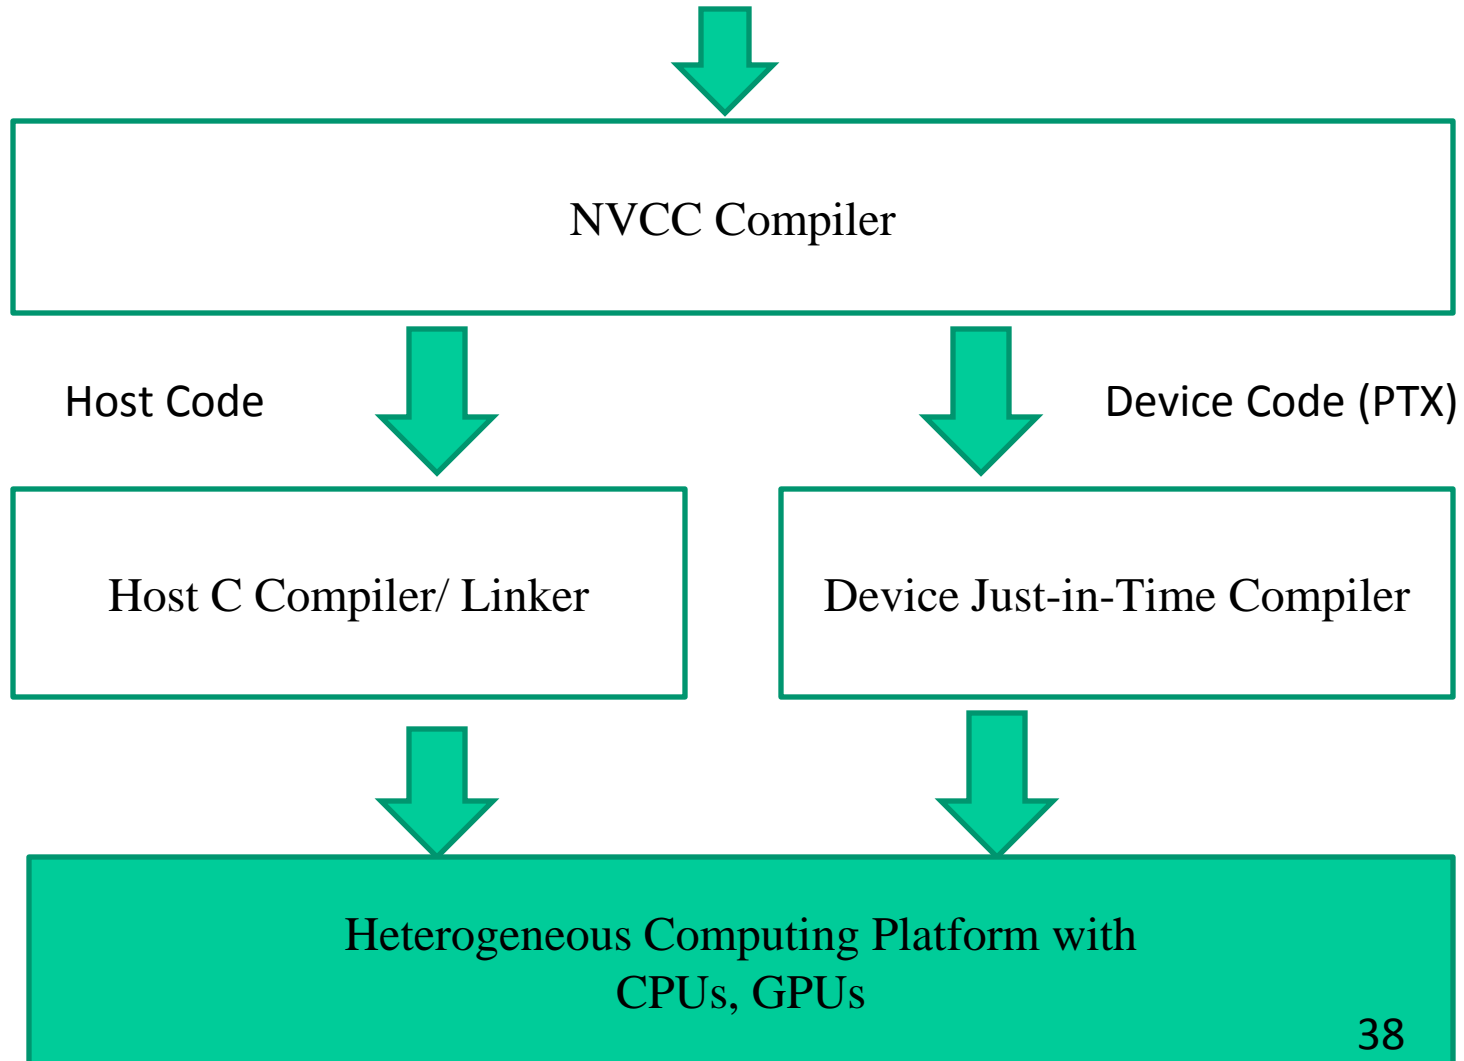
- dimGrid and dimBlock are user defined
- **gridDim** and **blockDim** are built-in predefined variable accessible in kernel functions

# Be Sure To Know:

- Maximum dimensions of a block
- Maximum number of threads per block
- Maximum dimensions of a grid
- Maximum number of blocks per thread

# Tools

Integrated C programs with CUDA extensions



# Conclusions

- Data parallelism is the main source of scalability for parallel programs
- Each CUDA source file can have a mixture of both host and device code.
- What we learned today about CUDA:
  - `KernelA<<< nBlk, nTid >>>(args)`
  - `cudaMalloc()`
  - `cudaFree()`
  - `cudaMemcpy()`
  - `gridDim` and `blockDim`
  - `threadIdx.x` and `threadIdx.y`
  - `dim3`