


# Fitting: The Hough transform

---



# Voting schemes

---

- Let each feature vote for all the models that are compatible with it
- Hopefully the noise features will not vote consistently for any single model
- Missing data doesn't matter as long as there are enough features remaining to agree on a good model

# Hough transform

---

- An early type of voting scheme
- General outline:
  - Discretize parameter space into bins
  - For each feature point in the image, put a vote in every bin in the parameter space that could have generated this point
  - Find bins that have the most votes






Image space




Hough parameter space

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

# Parameter space representation

---


- A line in the image corresponds to a point in Hough space



# Parameter space representation

---


- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?



# Parameter space representation

---


- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?
  - Answer: the solutions of  $b = -x_0m + y_0$
  - This is a line in Hough space



# Parameter space representation

---

- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?




# Parameter space representation


---

- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?
  - It is the intersection of the lines  $b = -x_0m + y_0$  and  $b = -x_1m + y_1$

Image space



Hough parameter space





# Parameter space representation


---

- Problems with the  $(m,b)$  space:
  - Unbounded parameter domains
  - Vertical lines require infinite  $m$

# Parameter space representation

---

- Problems with the (m,b) space:
  - Unbounded parameter domains
  - Vertical lines require infinite m
- Alternative: polar representation



Each point will add a sinusoid in the  $(\theta, \rho)$  parameter space

# Algorithm outline

---

- Initialize accumulator  $H$  to all zeros
- For each edge point  $(x,y)$  in the image

For  $\theta = 0$  to  $180$

$$\rho = x \cos \theta + y \sin \theta$$


$$H(\theta, \rho) = H(\theta, \rho) + 1$$

end

end


- Find the value(s) of  $(\theta, \rho)$  where  $H(\theta, \rho)$  is a local maximum
- The detected line in the image is given by  
$$\rho = x \cos \theta + y \sin \theta$$

$H$ : accumulator array (votes)




# Basic illustration

---



features




votes


# Other shapes

---

Square





Circle




# Several lines

---




# A more complicated image

---



# Effect of noise


---






# Effect of noise

---



features




votes

Peak gets fuzzy and hard to locate

# Effect of noise


---

- Number of votes for a line of 20 points with increasing noise:




# Random points

---



features




votes

Uniform noise can lead to spurious peaks in the array

# Random points

---

- As the level of uniform noise increases, the maximum number of votes increases too:



# Dealing with noise


---

- Choose a good grid / discretization
  - **Too coarse:** large votes obtained when too many different lines correspond to a single bucket
  - **Too fine:** miss lines because some points that are not exactly collinear cast votes for different buckets
- Increment neighboring bins (smoothing in accumulator array)
- Try to get rid of irrelevant features
  - Take only edge points with significant gradient magnitude

# Incorporating image gradients

---

- Recall: when we detect an edge point, we also know its gradient direction
- But this means that the line is uniquely determined!
- Modified Hough transform:



$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

For each edge point (x,y)

$\theta$  = gradient orientation at (x,y)

$\rho = x \cos \theta + y \sin \theta$

$H(\theta, \rho) = H(\theta, \rho) + 1$

end

# Hough transform for circles


---

- How many dimensions will the parameter space have?
- Given an oriented edge point, what are all possible bins that it can vote for?


# Hough transform for circles

---

image space



Hough parameter space






# Hough transform for circles

---

- Conceptually equivalent procedure: for each  $(x,y,r)$ , draw the corresponding circle in the image and compute its “support”




Is this more or less efficient than voting with features?

# Generalized Hough transform

---

- We want to find a template defined by its reference point (center) and several distinct types of landmark points in stable spatial configuration

## Template




# Generalized Hough transform


---

- Template representation:  
for each type of landmark  
point, store all possible  
displacement vectors  
towards the center

**Template**



**Model**




# Generalized Hough transform


---

- Detecting the template:
  - For each feature in a new image, look up that feature type in the model and vote for the possible center locations associated with that type in the model

**Test image**




**Model**




# Application in recognition

---

- Index displacements by “visual codeword”



training image



visual codeword with  
displacement vectors

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

# Application in recognition

---

- Index displacements by “visual codeword”




test image

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

# Implicit shape models: Training

---




1. Build codebook of patches around extracted interest points using clustering (more on this later in the course)



# Implicit shape models: Training

---

1. Build codebook of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry








# Implicit shape models: Training

---

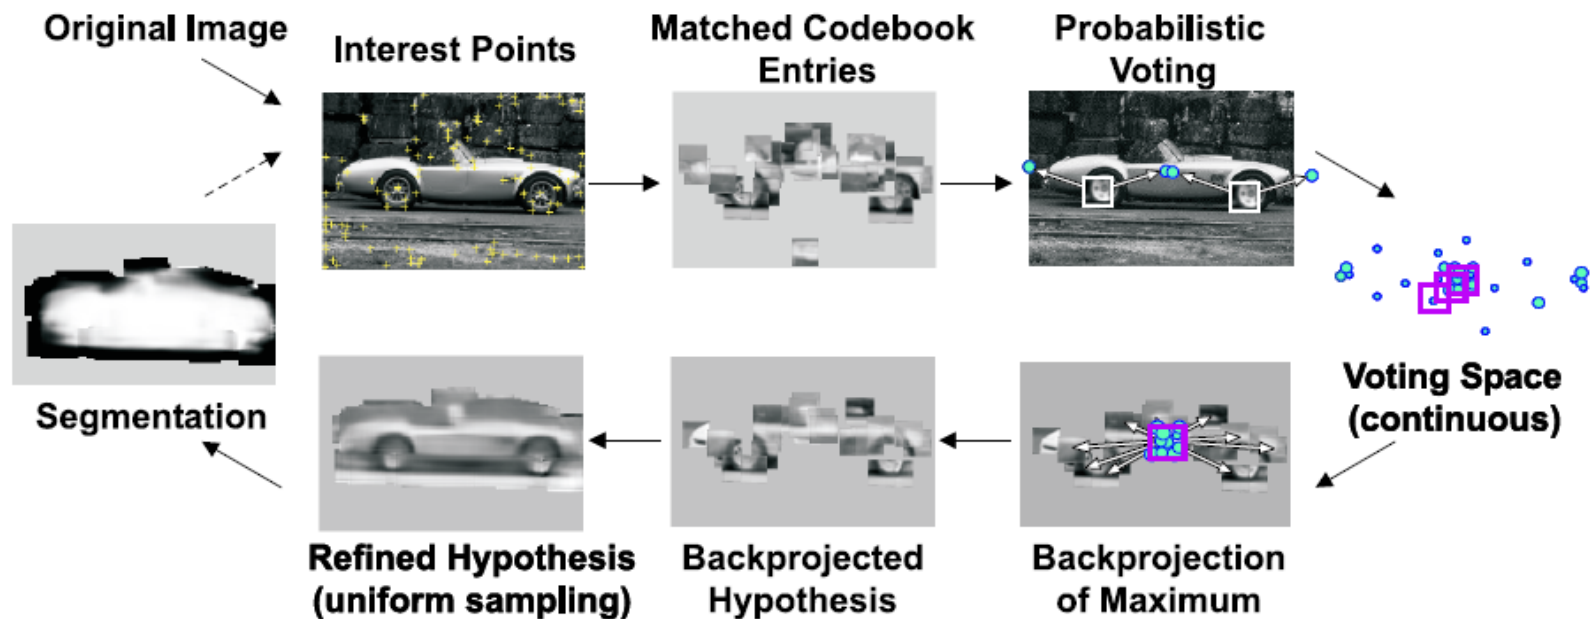
1. Build codebook of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry
3. For each codebook entry, store all positions it was found, relative to object center



# Implicit shape models: Testing






---

1. Given test image, extract patches, match to codebook entry
2. Cast votes for possible positions of object center
3. Search for maxima in voting space
4. Extract weighted segmentation mask based on stored masks for the codebook occurrences



# Additional examples

---



B. Leibe, A. Leonardis, and B. Schiele, [Robust Object Detection with Interleaved Categorization and Segmentation](#), IJCV 77 (1-3), pp. 259-289, 2008.

# Implicit shape models: Details

---

- Supervised training
  - Need reference location and segmentation mask for each training car
- Voting space is continuous, not discrete
  - Clustering algorithm needed to find maxima
- How about dealing with scale changes?
  - Option 1: search a range of scales, as in Hough transform for circles
  - Option 2: use interest points with characteristic scale
- Verification stage is very important
  - Once we have a location hypothesis, we can overlay a more detailed template over the image and compare pixel-by-pixel, transfer segmentation masks, etc.

# Hough transform: Discussion

---

- Pros
  - Can deal with non-locality and occlusion
  - Can detect multiple instances of a model
  - Some robustness to noise: noise points unlikely to contribute consistently to any single bin
- Cons
  - Complexity of search time increases exponentially with the number of model parameters
  - Non-target shapes can produce spurious peaks in parameter space
  - It's hard to pick a good grid size