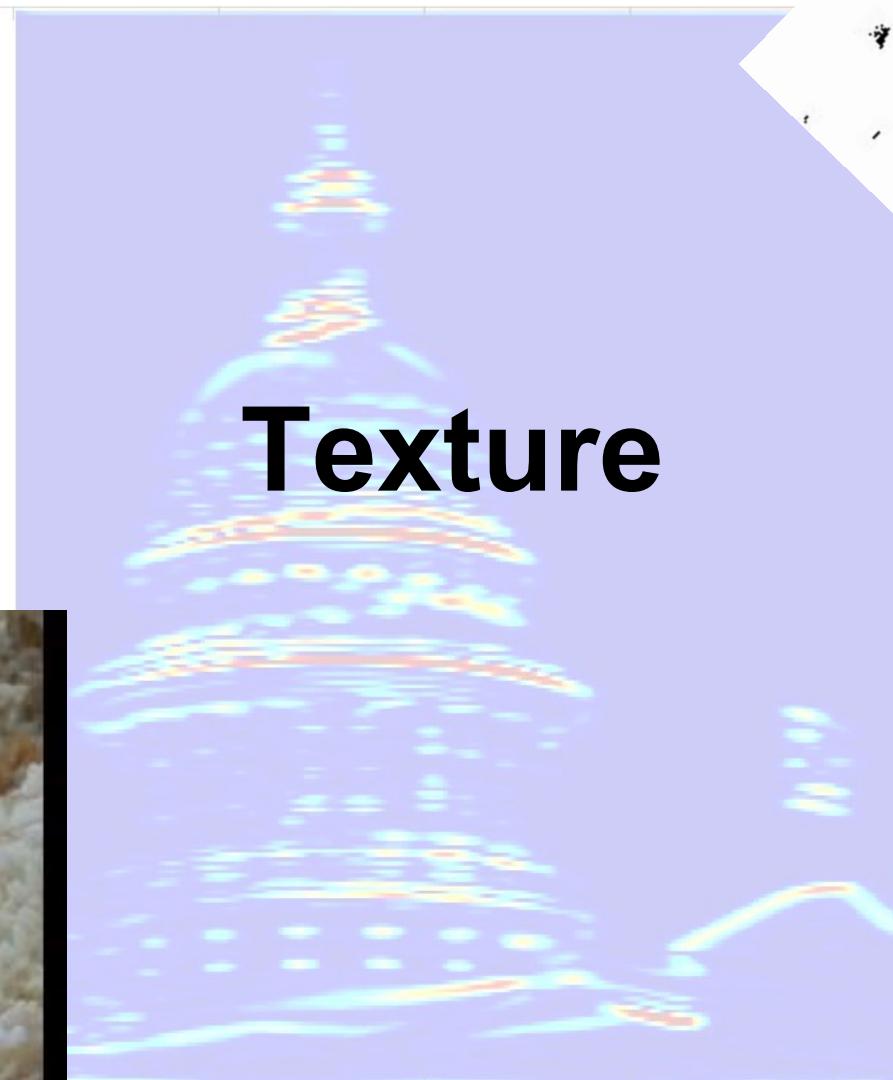


# Texture

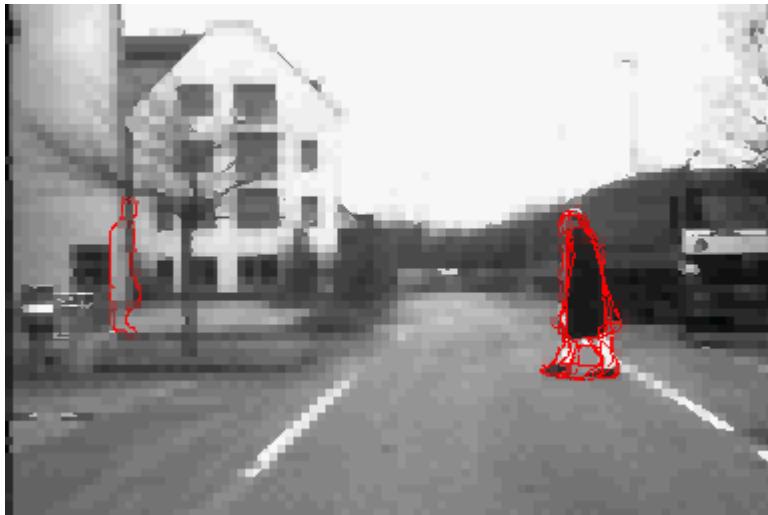


# Chamfer matching system



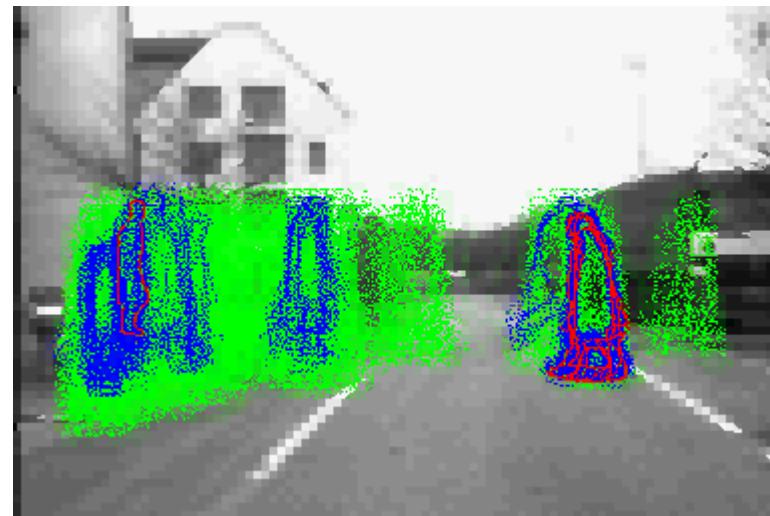
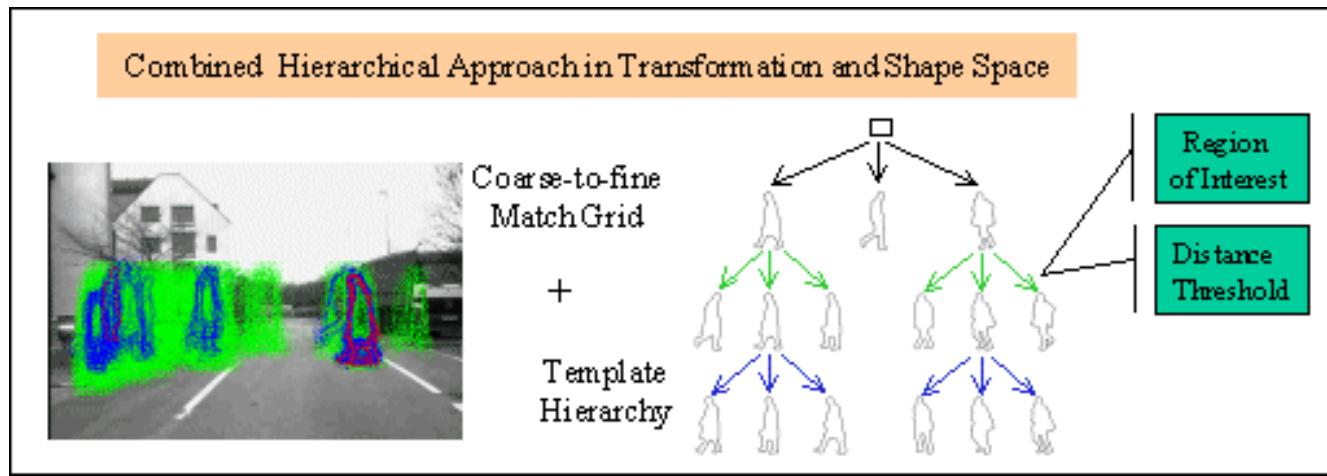
- Gavrila et al.  
[http://gavrila.net/Research/Chamfer\\_System/chamfer\\_system.html](http://gavrila.net/Research/Chamfer_System/chamfer_system.html)

# Chamfer matching system



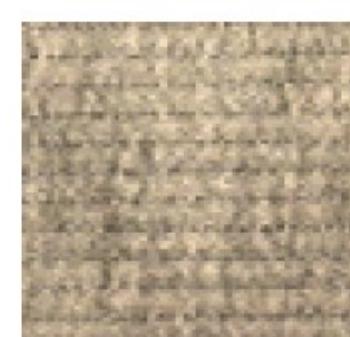
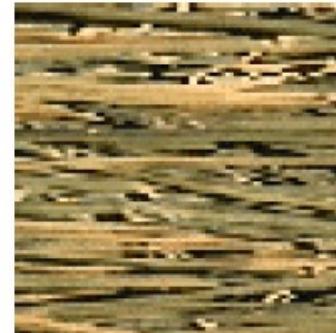
- Gavrila et al.  
[http://gavrila.net/Research/Chamfer\\_System/chamfer\\_system.html](http://gavrila.net/Research/Chamfer_System/chamfer_system.html)

# Chamfer matching system



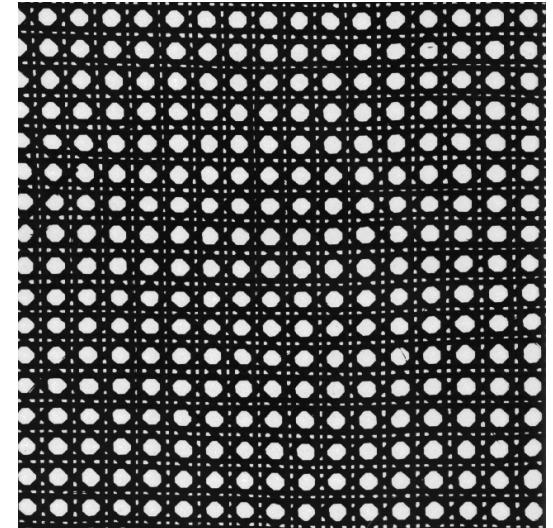
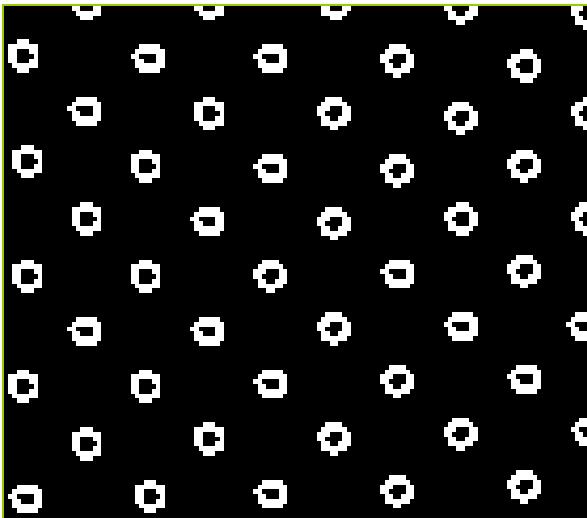
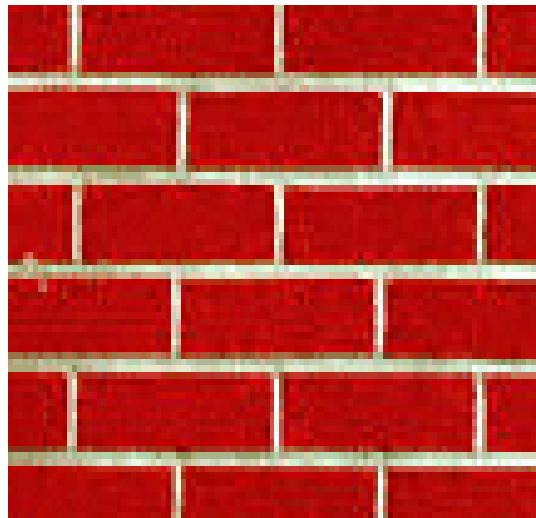
- Gavrila et al.  
[http://gavrila.net/Research/Chamfer\\_System/chamfer\\_system.html](http://gavrila.net/Research/Chamfer_System/chamfer_system.html)

# Today: Texture

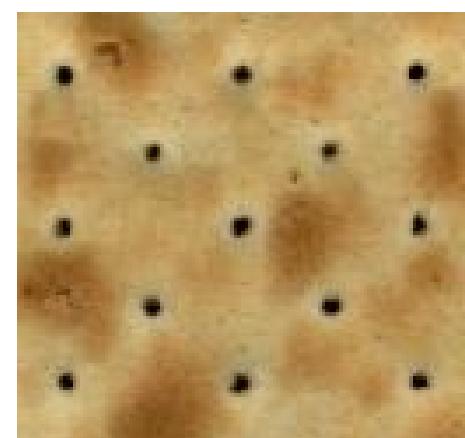
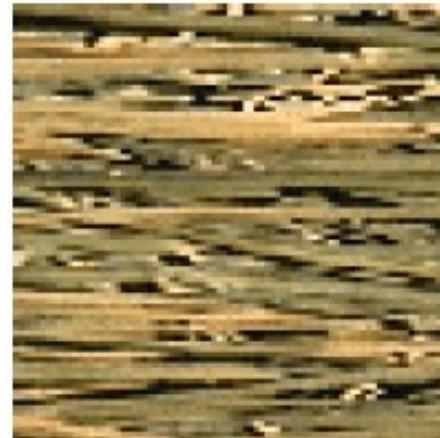
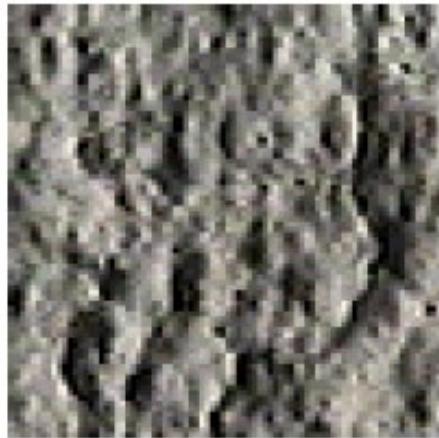


What defines a texture?

# Includes: more regular patterns



# Includes: more random patterns

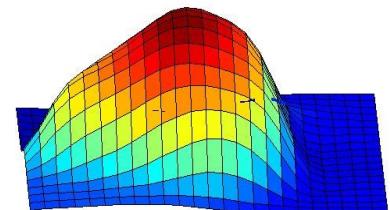
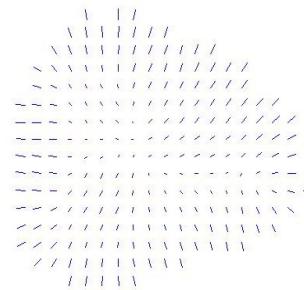
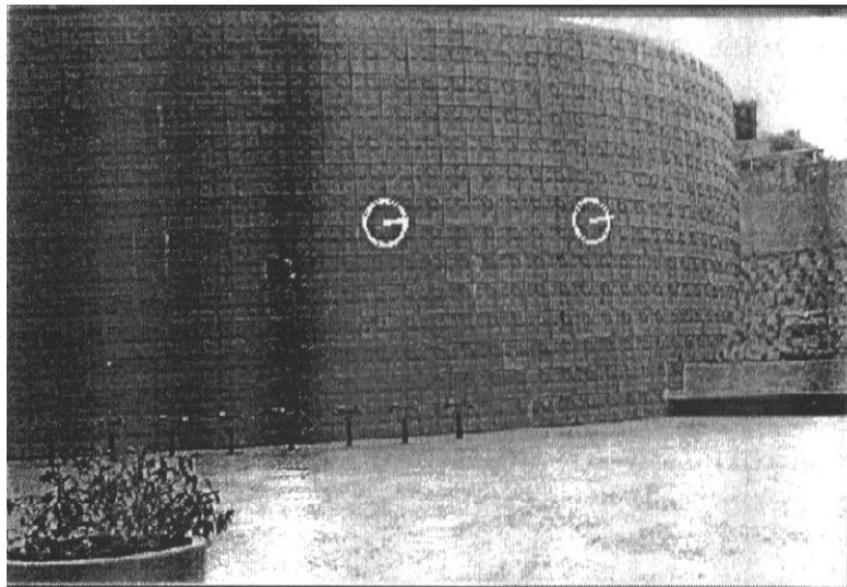


# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture

# Shape from texture

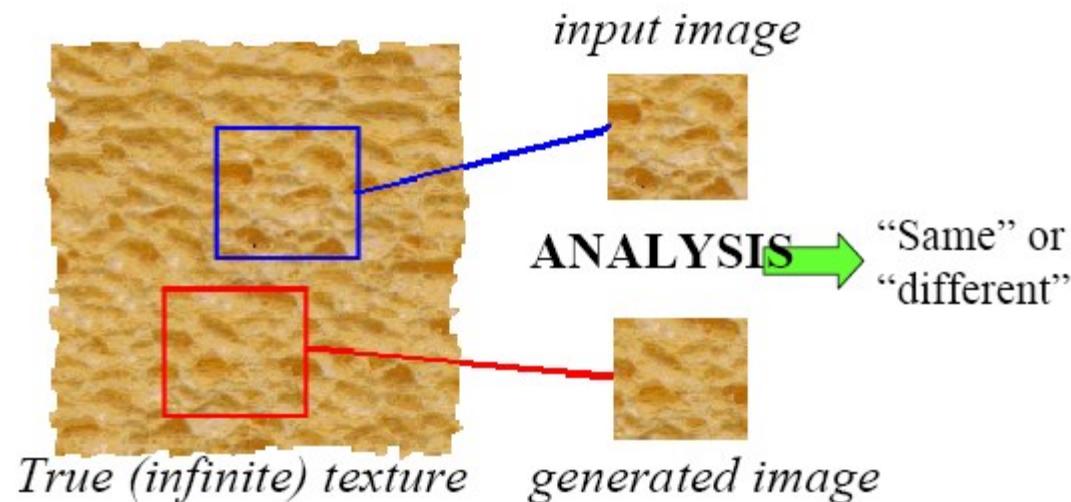
- Use deformation of texture from point to point to estimate surface shape



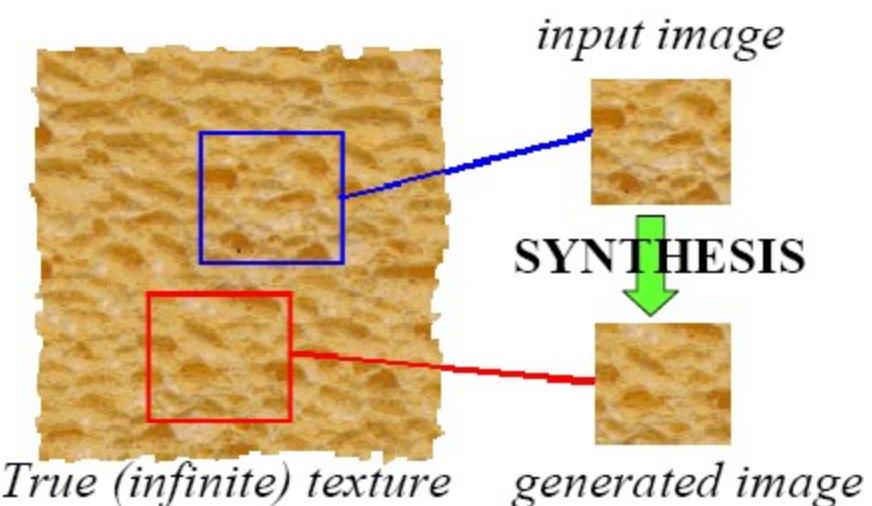
# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
  - Analyze, represent texture
  - Group image regions with consistent texture
- **Synthesis**
  - Generate new texture patches/images given some examples

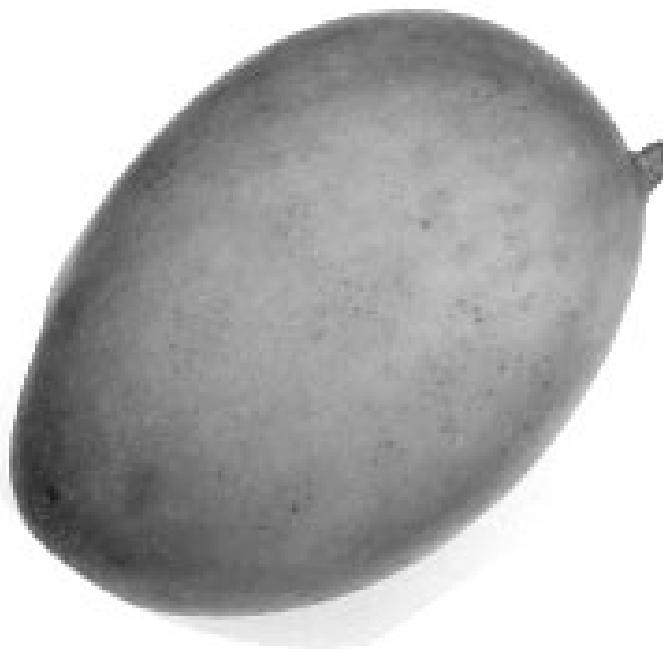
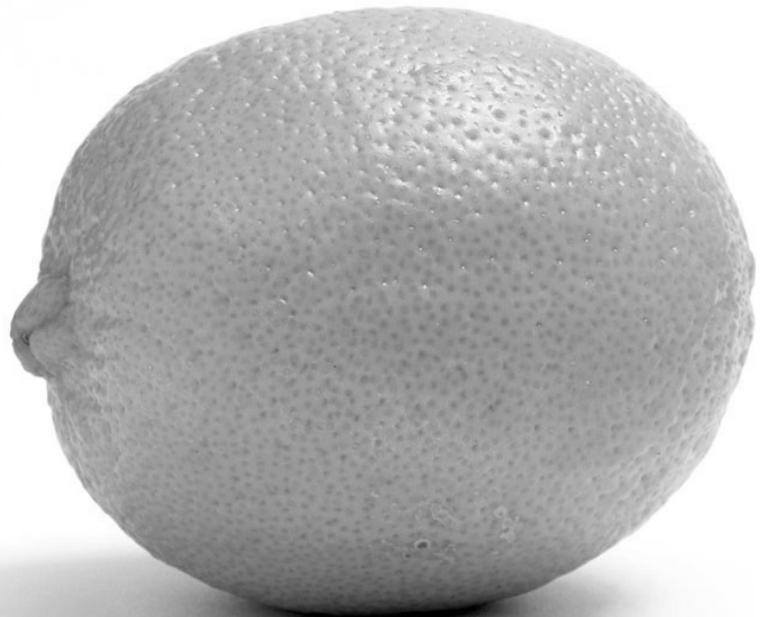
# Analysis vs. Synthesis

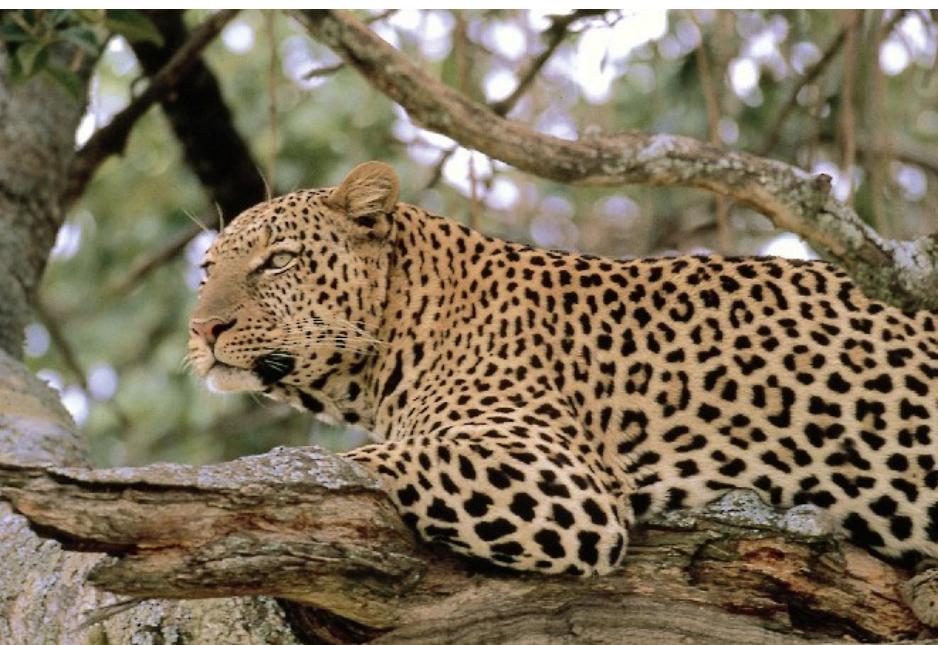


Why analyze  
texture?











What kind of response will we get with an edge detector for these images?



...and for this image?

# Why analyze texture?

Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.

A 7x7 grid of black L-shaped blocks on a white background. Each block is composed of three squares: one vertical column of two squares and one horizontal square extending from the middle of the left side of the column. The blocks are arranged in a staggered pattern: the first row has blocks at columns 2 through 8; the second row has blocks at columns 1 through 7; the third row has blocks at columns 2 through 8; and so on, creating a repeating pattern of 5x5 subgrids.

The image shows a large, dense grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, creating a complex and intricate design. The grid consists of approximately 10 columns and 10 rows of these L-shaped blocks.

A 10x10 grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, where each block is positioned such that its vertical stem is aligned with the horizontal bar of the block directly above it. This creates a continuous, winding path across the entire grid.

The image shows a large, dense grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, creating a complex and intricate design. The grid consists of approximately 10 columns and 10 rows of these L-shaped blocks, resulting in a total of about 100 individual blocks. The blocks are rendered in a solid black color and have a thick, bold outline. The overall effect is a high-contrast, geometric artwork.

# Capturing the local patterns with image measurements



[Bergen &  
Adelson,  
*Nature* 1988]

Scale of  
patterns  
influences  
discriminability

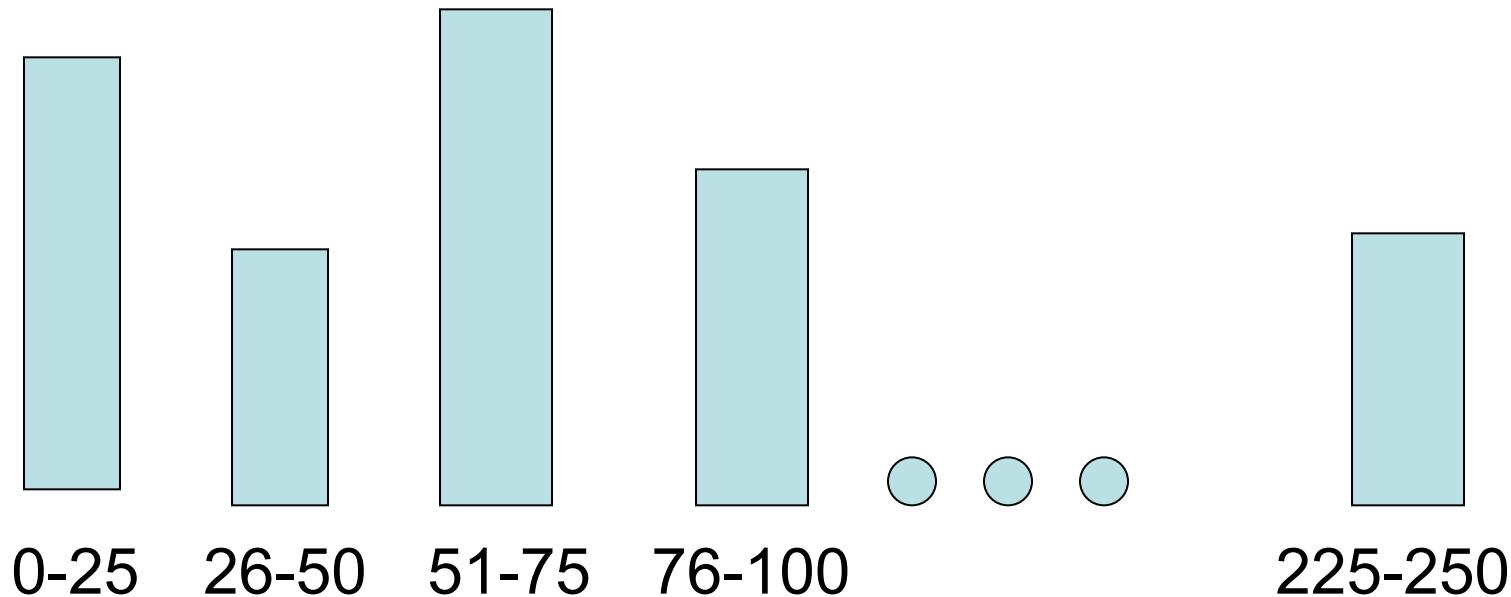
Size-tuned  
linear filters

# Texture representation

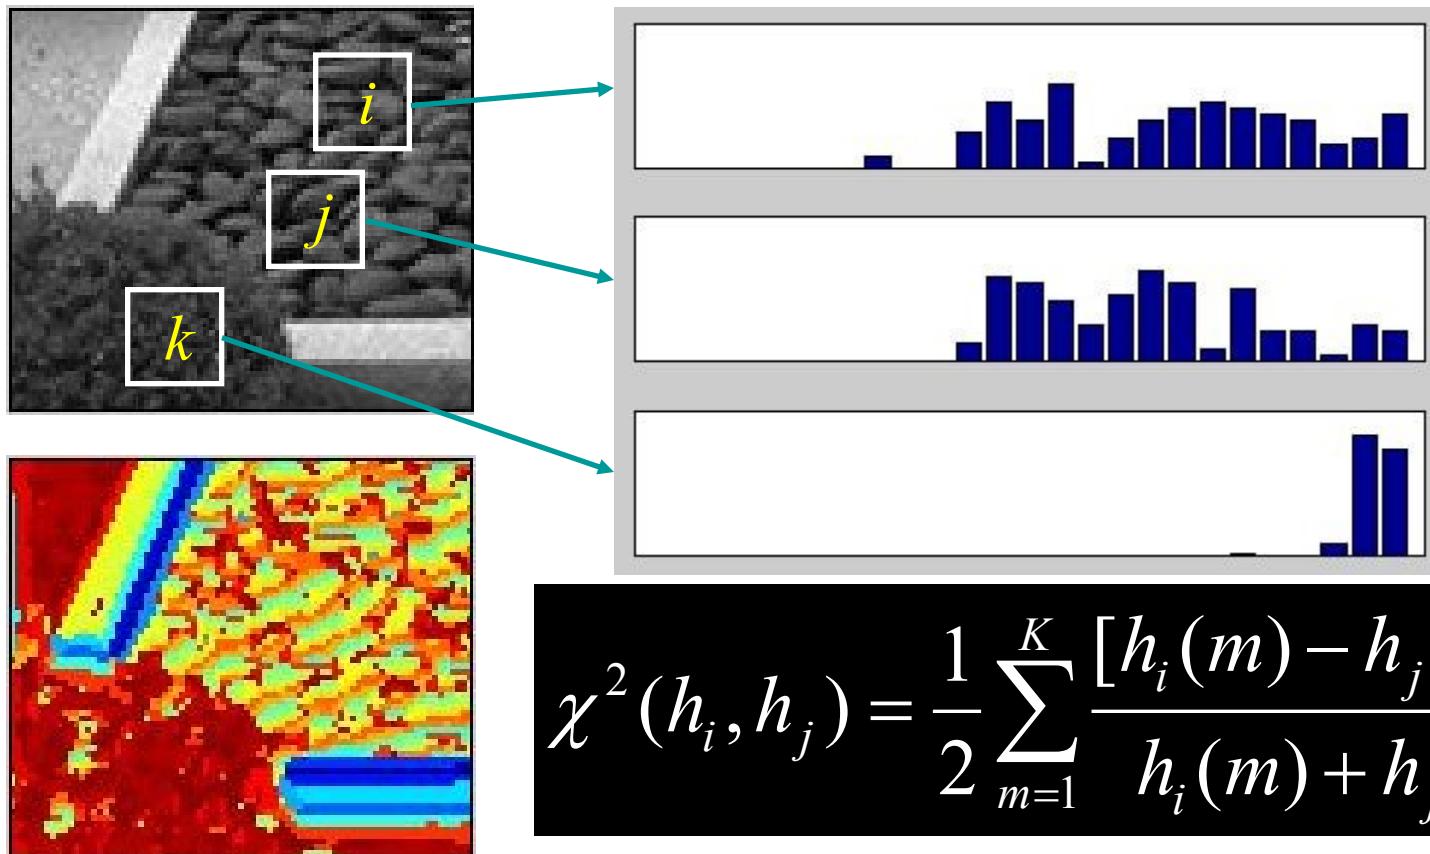
- Textures are made up of repeated local patterns, so:
  - Find the patterns
    - Use filters that look like patterns (spots, bars, raw patches...)
    - Consider magnitude of response
  - Describe their statistics within each local window
    - Mean, standard deviation
    - Histogram
    - Histogram of “prototypical” feature occurrences

# Simplest Texture Discrimination

- Compare histograms.
  - Divide intensities into discrete ranges.
  - Count how many pixels in each range.



# Chi square distance between texton histograms



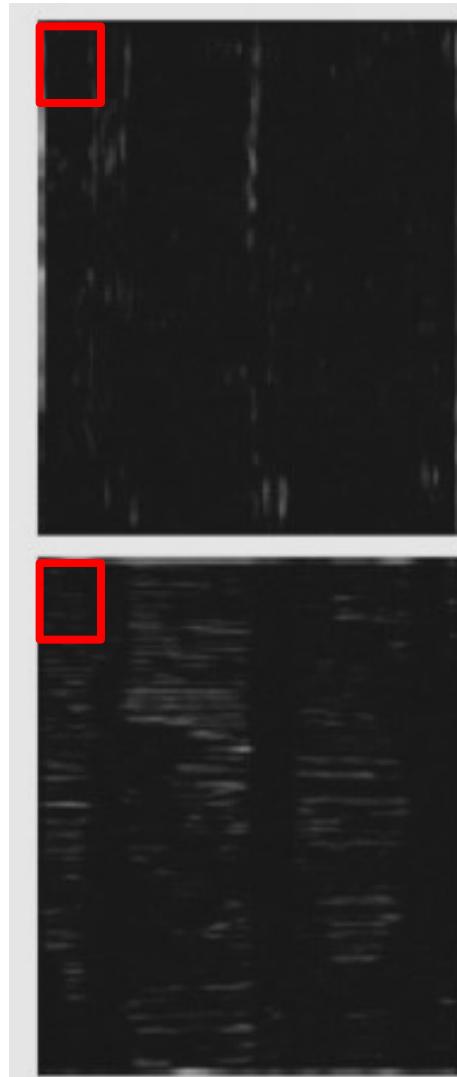
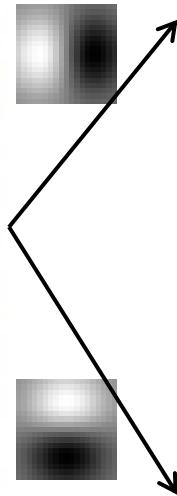
$$\chi^2(h_i, h_j) = \frac{1}{2} \sum_{m=1}^K \frac{[h_i(m) - h_j(m)]^2}{h_i(m) + h_j(m)}$$

(Malik)

# Texture representation: example



original image



derivative filter  
responses, squared

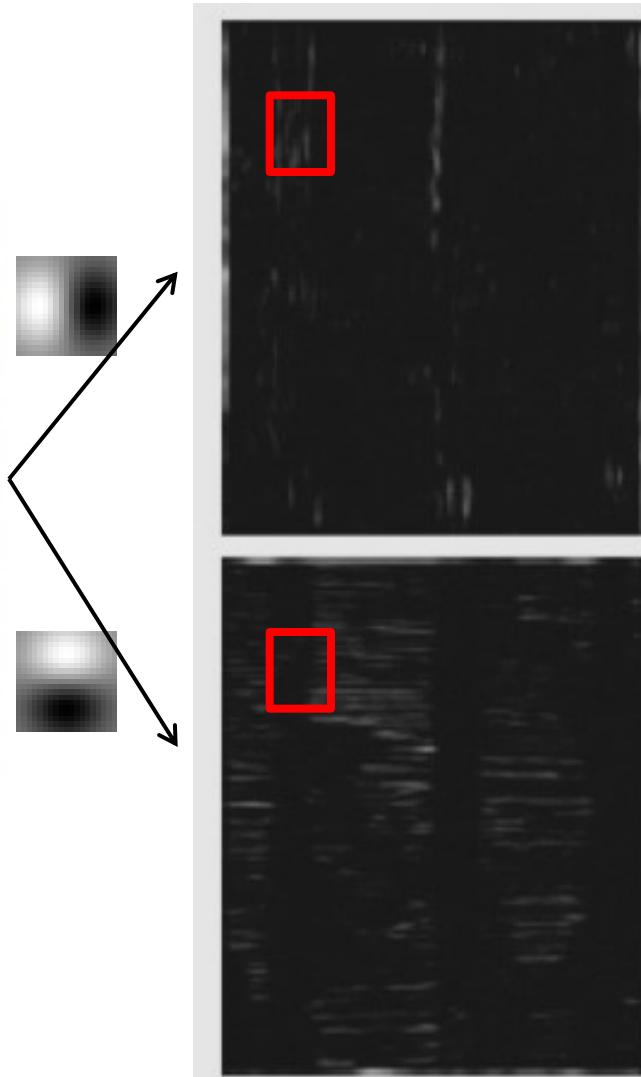
	<u>mean <math>d/dx</math> value</u>	<u>mean <math>d/dy</math> value</u>
Win. #1	4	10
⋮	⋮	⋮
⋮	⋮	⋮
⋮	⋮	⋮

statistics to  
summarize patterns  
in small windows

# Texture representation: example



# original image



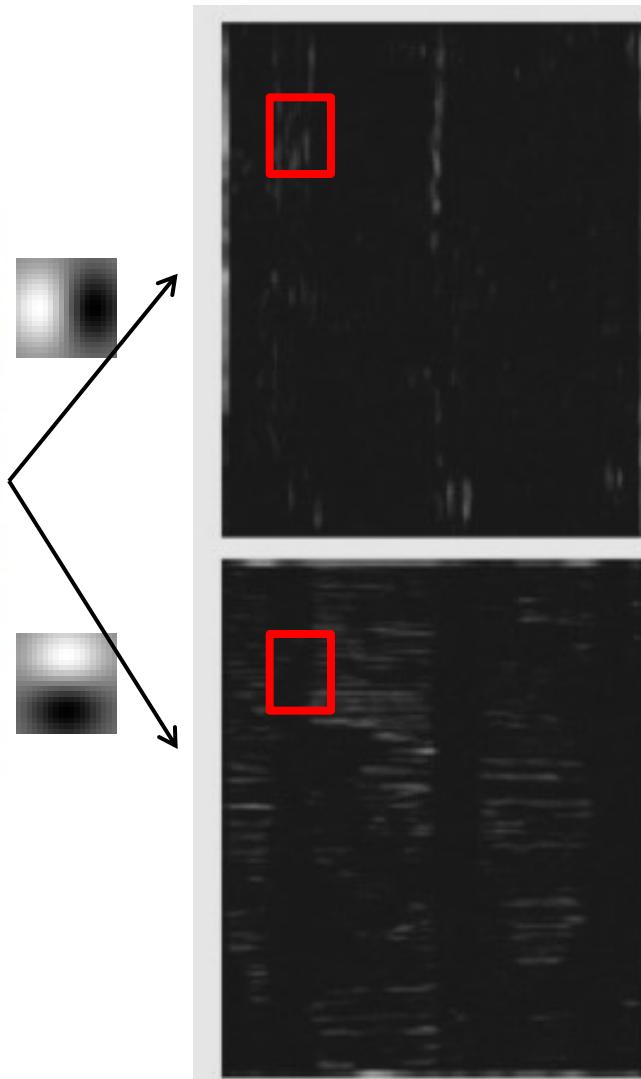
## derivative filter responses, squared

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



# original image



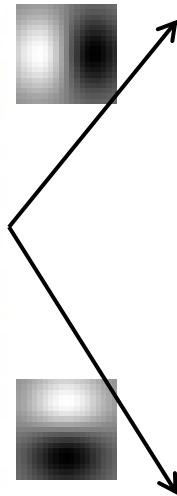
## derivative filter responses, squared

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



original image



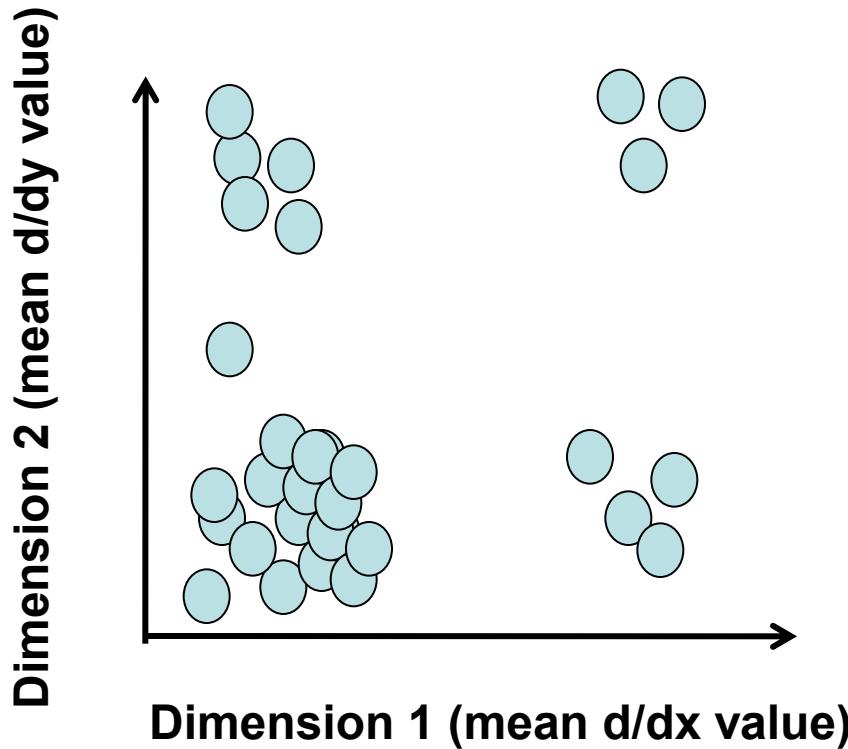
derivative filter  
responses, squared

	<u>mean <math>d/dx</math> value</u>	<u>mean <math>d/dy</math> value</u>
Win. #1	4	10
Win.#2	18	7
:		
Win.#9	20	20

⋮

statistics to  
summarize patterns  
in small windows

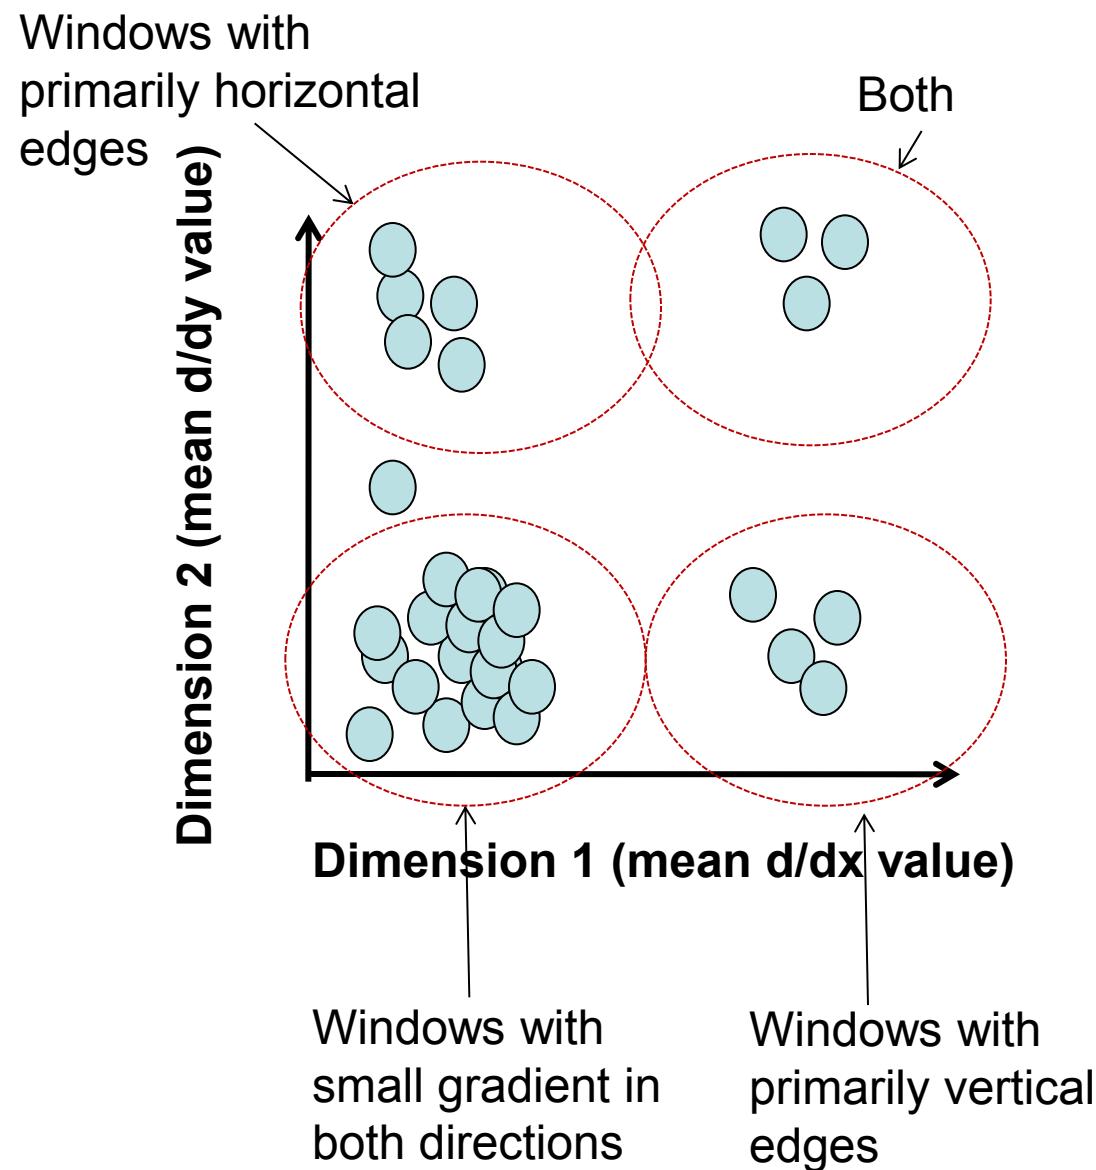
# Texture representation: example



	<u>mean <math>d/dx</math> value</u>	<u>mean <math>d/dy</math> value</u>
Win. #1	4	10
Win.#2	18	7
⋮		
Win.#9	20	20
⋮		

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



	<u>mean <math>d/dx</math> value</u>	<u>mean <math>d/dy</math> value</u>
Win. #1	4	10
Win.#2	18	7
:		
Win.#9	20	20

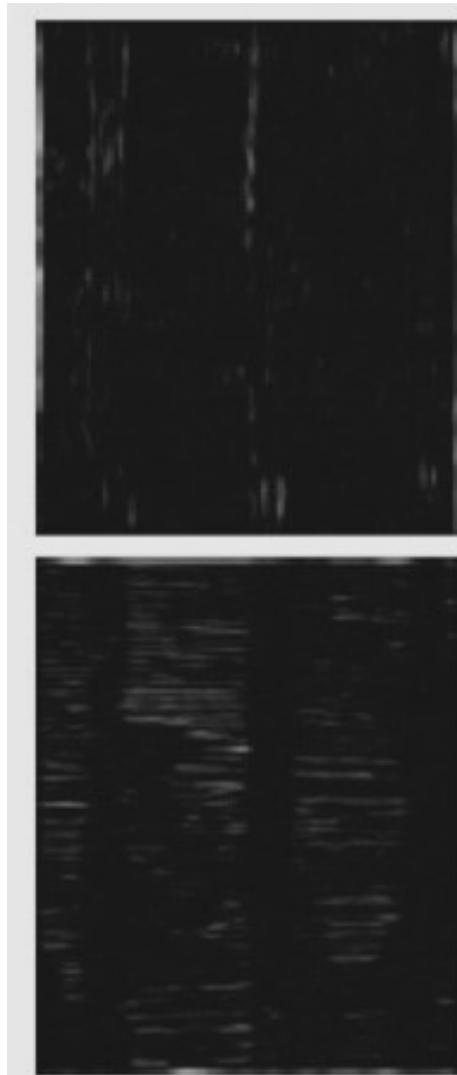
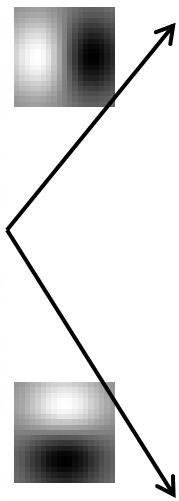
⋮

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



original image

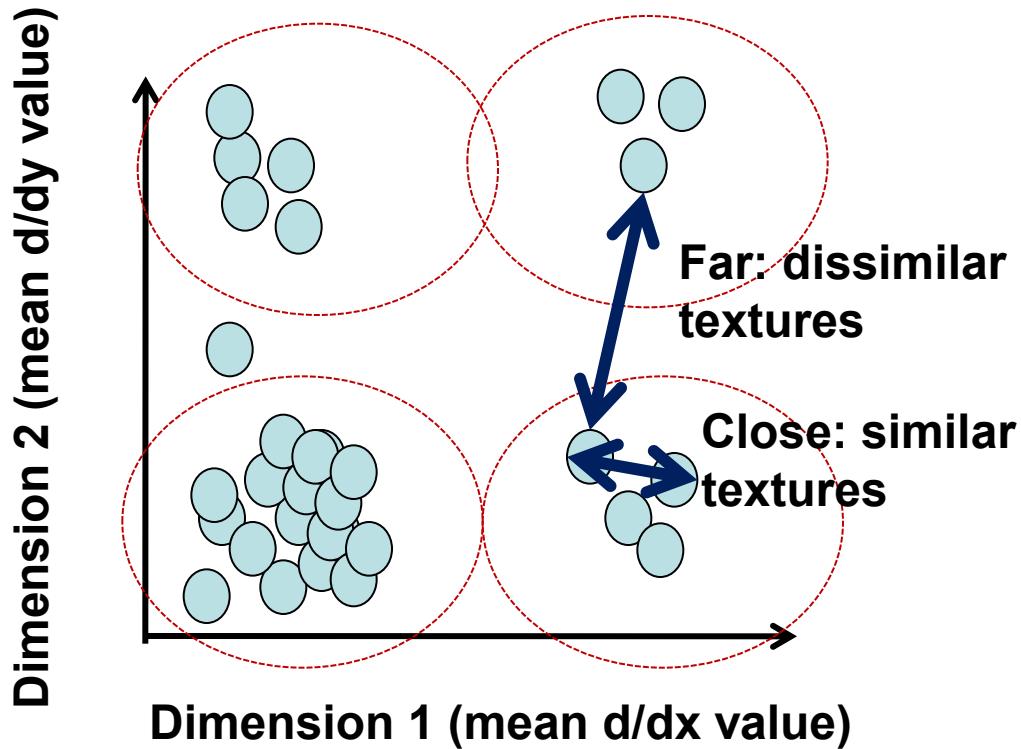


derivative filter  
responses, squared



visualization of the  
assignment to  
texture “types”

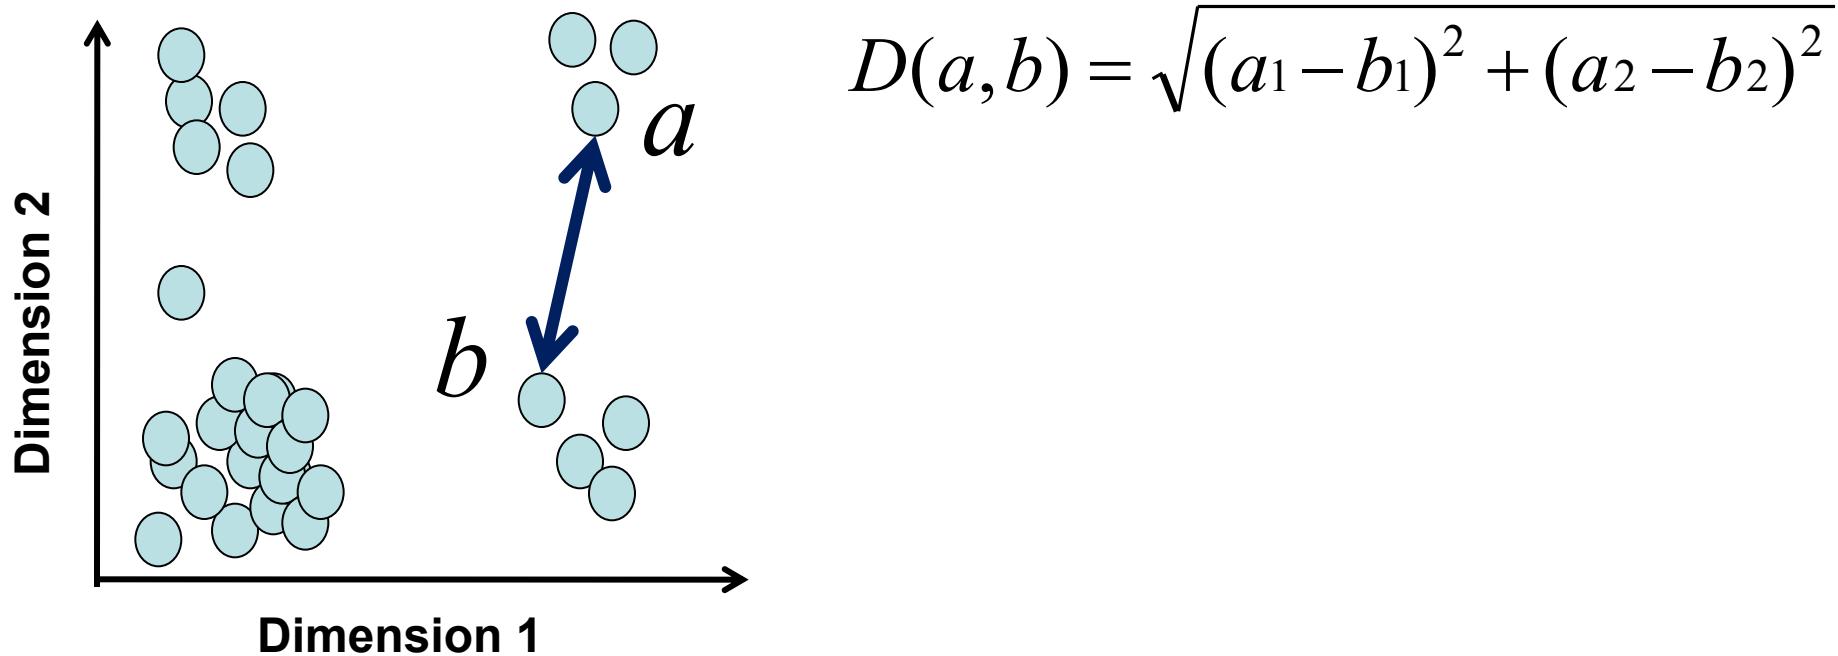
# Texture representation: example



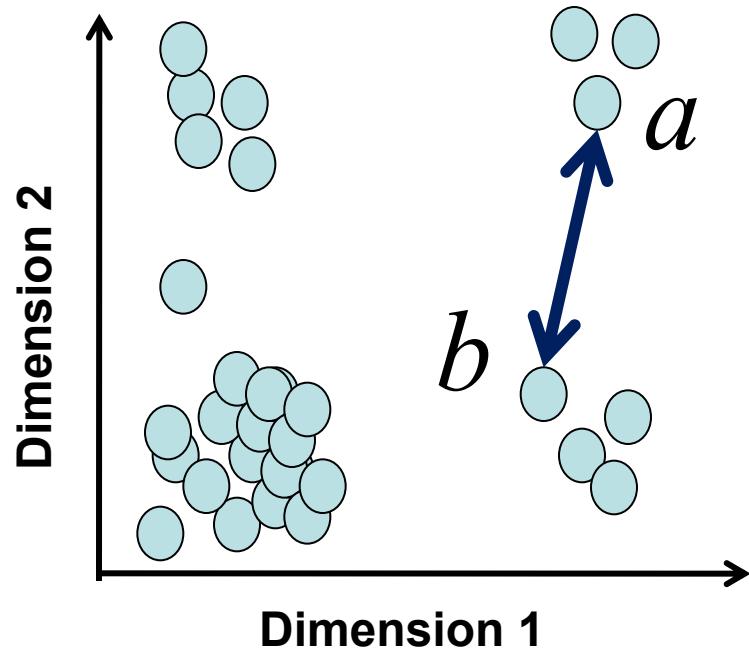
	<u>mean <math>d/dx</math> value</u>	<u>mean <math>d/dy</math> value</u>
Win. #1	4	10
Win. #2	18	7
:		
Win. #9	20	20
⋮		

statistics to  
summarize patterns  
in small windows

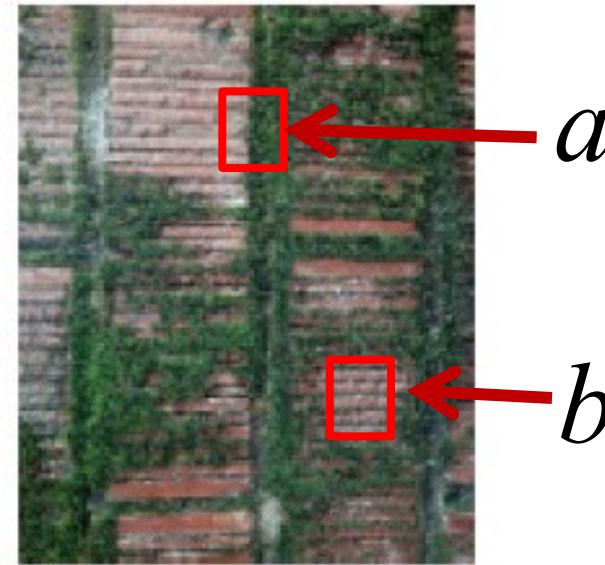
# Texture representation: example



# Texture representation: example



Distance reveals how dissimilar texture from window *a* is from texture in window *b*.



# Texture Analysis Using Edges

## ■ Edge Density and Directions

Consider a region of  $N$  Pixels.

$$\text{方向: } G = i \frac{\partial I}{\partial x} + j \frac{\partial I}{\partial y} \quad \text{大小: } |G| = \sqrt{\left[ \frac{\partial I}{\partial x} \right]^2 + \left[ \frac{\partial I}{\partial y} \right]^2}$$

$$Edgeness = \frac{|\{p \mid Mag(p) \geq T\}|}{N}$$

# Textures

- Edge Density and Directions

Consider a region of N Pixels.

Direction: Divide the angle to different bins,  
90, 45, or 30 degrees

Magnitude: sampling with a fixed size

Edge Histogram:  $H_{dir}(i), H_{mag}(i)$

$$L(H_1, H_2) = \sum_{n=1}^N |H_1[i] - H_2[i]|$$

# Second-order gray-level statistics

- Co-occurrence Matrices and Features

Let  $d=(dx,dy)$ .

<b>1</b>	<b>1</b>		
<b>1</b>	<b>1</b>		
		<b>2</b>	<b>2</b>
		<b>2</b>	<b>2</b>

The gray-tone cooccurrence matrix  $C_d$  is:

$$C_d(m,n) = |\{(x,y) \mid I(x,y)=m \text{ and } I(x+dx,y+dy)=n\}|$$

$(dx,dy)=(1,0)$

	0	1	2
0	4	0	2
1	2	2	0
2	0	0	2

$(dx,dy)=(0,1)$

	0	1	2
0	4	0	2
1	2	2	0
2	0	0	2

$(dx,dy)=(1,1)$

	0	1	2
0	2	0	1
1	2	1	1
2	0	0	1

# Textures

- Co-occurrence Matrix

Normalized co-occurrence matrix:

$$N_d(m,n) = \frac{C_d(m,n)}{\sum_{i,j} C_d(i,j)}$$

$$Energy = \sum_{i,j} N_d^2(i,j)$$

$$Entropy = \sum_{i,j} N_d(i,j) \log N_d(i,j)$$

$$Contrast = \sum_{i,j} (i-j)^2 N_d(i,j)$$

$$Homogeneity = \sum_{i,j} \frac{N_d(i,j)}{1 + |i-j|}$$

$$Correlation = \sum_{i,j} \frac{(i - u_i)(j - u_j) N_d(i,j)}{\sigma_i \sigma_j}$$

# Textures

## ■ Transform-based Statistical Methods

- Power Spectrum Method:

Autocorrelation function:

$$\rho(x, y) = \frac{\sum_{m,n} I(m, n)I(m + x, n + y)}{\sum_{m,n} |I(m, n)|^2}$$

Texture:

coarse: drop off slows;

flat: drop off very quick.

regular: have peaks and valleys.

Let  $F(u, v)$  = Fourier Transform of  $I(m, n)$ .

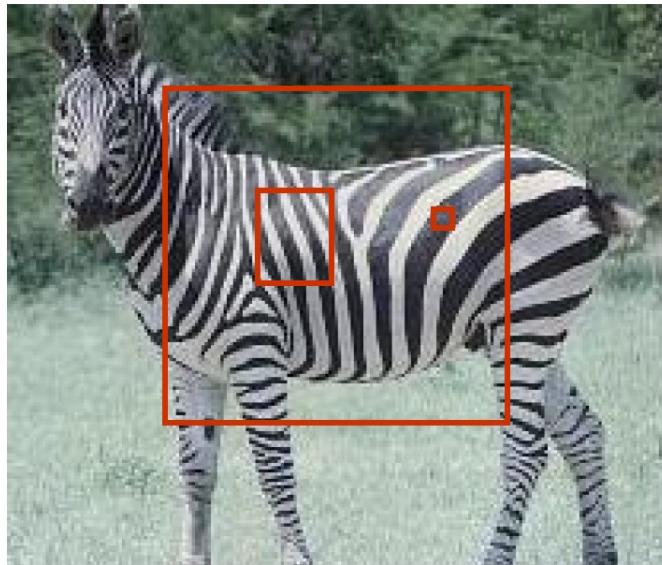
Power Spectrum =  $|F(u, v)|^2$

# More Complex Discrimination

- Histogram comparison is very limiting
  - Every pixel is independent.
  - Everything happens at a tiny scale.
- Use output of filters of different scales.

# Texture representation: window scale

- We're assuming we know the relevant window size for which we collect these statistics.

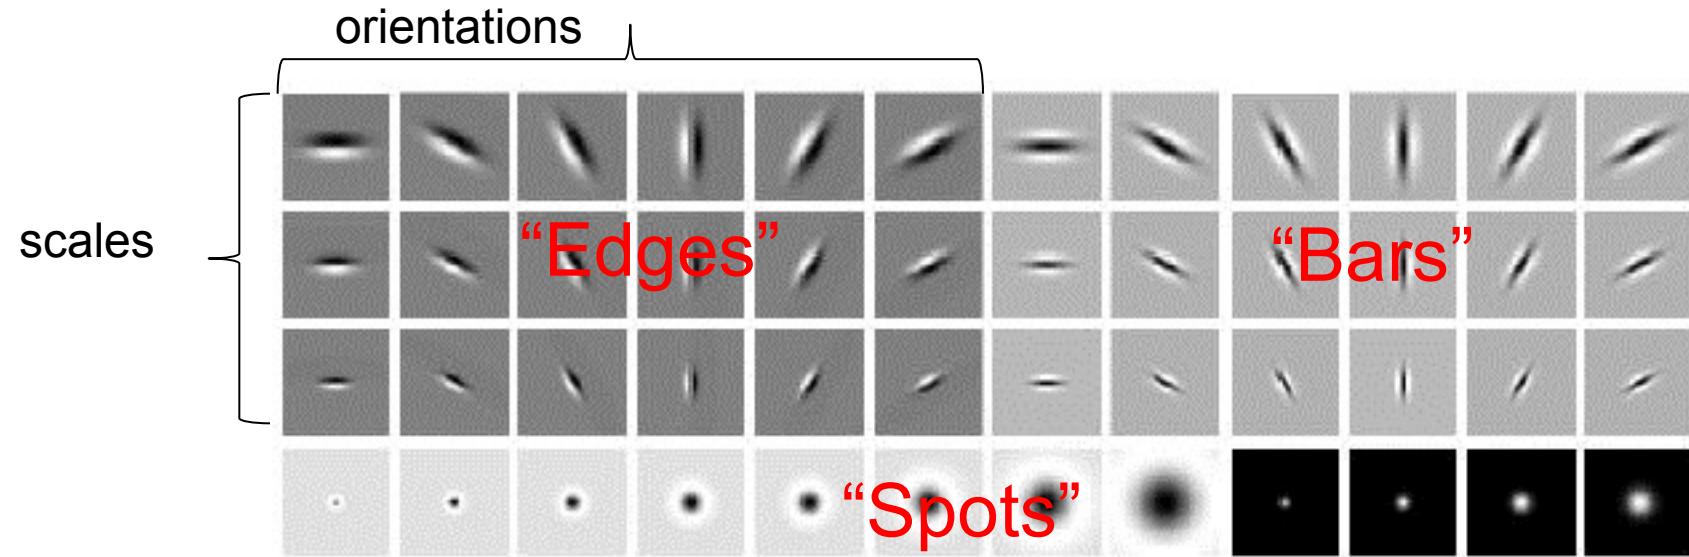


Possible to perform scale selection by looking for window scale where texture description not changing.

# Filter banks

- Our previous example used two filters, and resulted in a 2-dimensional feature vector to describe texture in a window.
  - x and y derivatives revealed something about local structure.
- We can generalize to apply a collection of multiple ( $d$ ) filters: a “filter bank”
- Then our feature vectors will be  $d$ -dimensional.
  - still can think of nearness, farness in feature space

# Filter banks



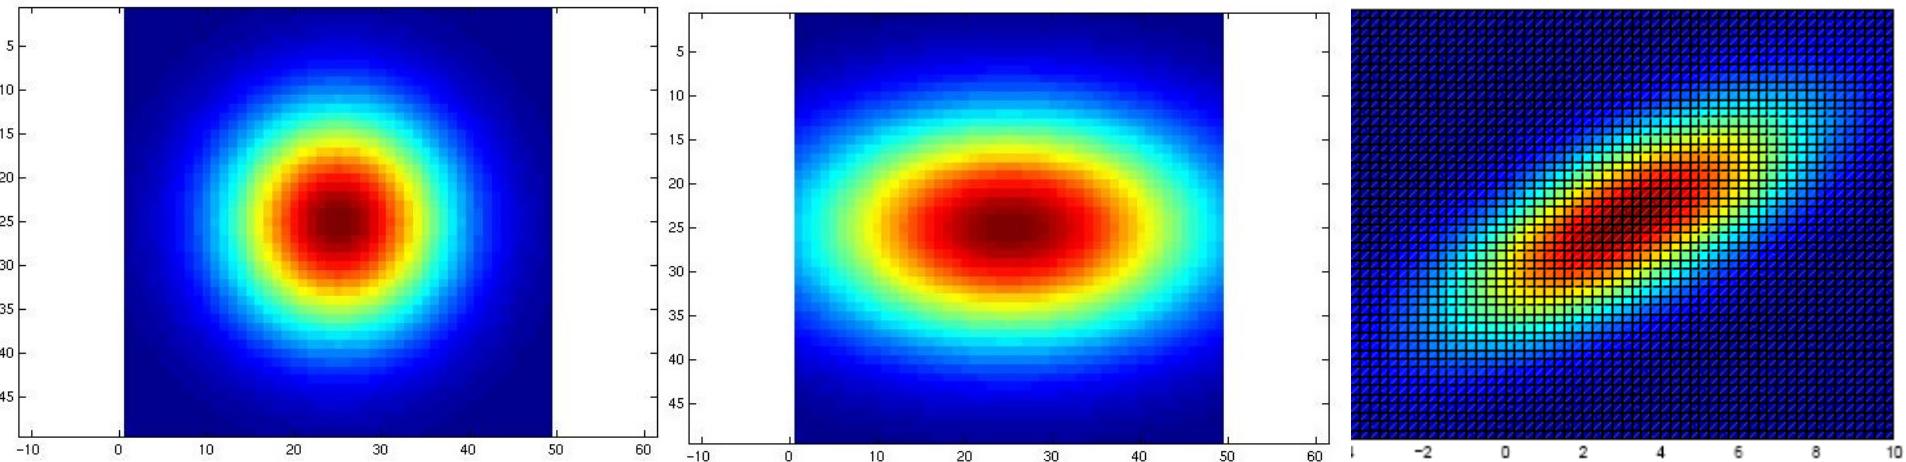
- What filters to put in the bank?
  - Typically we want a combination of scales and orientations, different types of patterns.

Matlab code available for these examples:

<http://www.robots.ox.ac.uk/~vgg/research/texclass/filters.html>

# Multivariate Gaussian

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

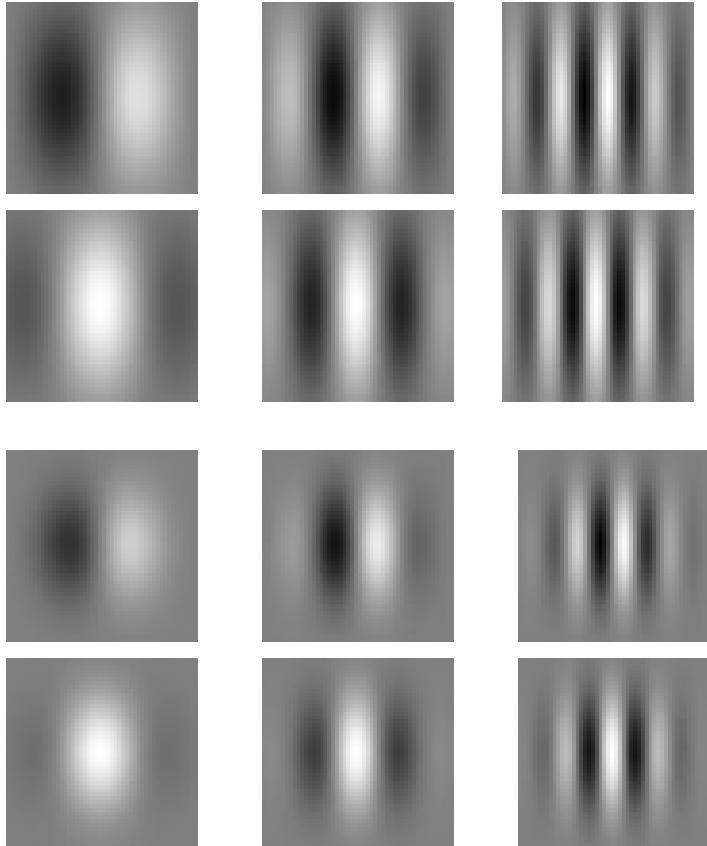


$$\Sigma = \begin{bmatrix} 9 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 16 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 10 & 5 \\ 5 & 5 \end{bmatrix}$$

# Gabor Filters



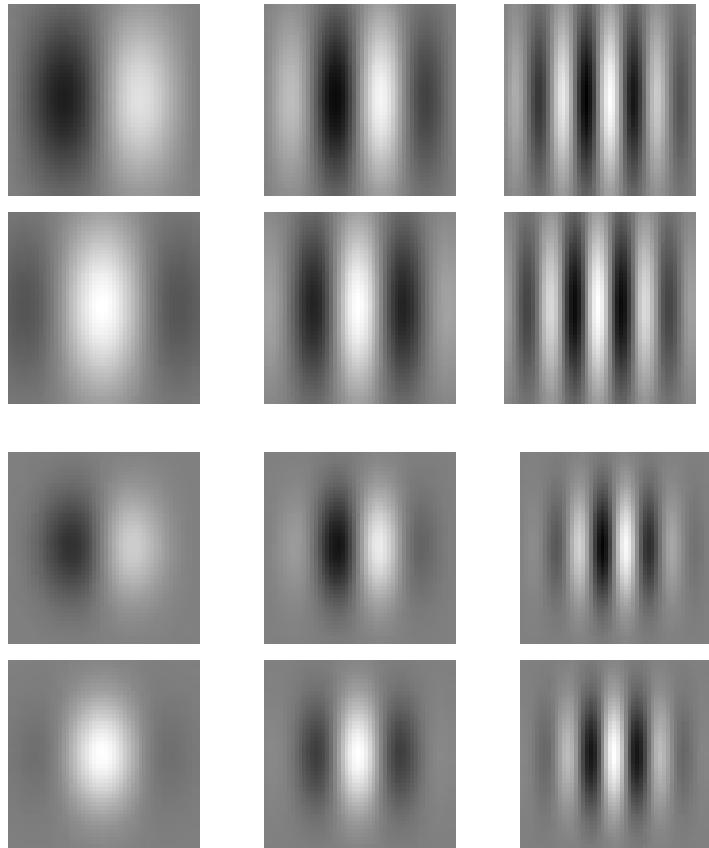
Gabor filters at different scales and spatial frequencies

top row shows anti-symmetric (or odd) filters, bottom row the symmetric (or even) filters.

$$\text{symmetric: } \cos(k_x x + k_y y) \exp - \left\{ \frac{x^2 + y^2}{2\sigma^2} \right\}$$

$$\text{antisymmetric : } \sin(k_x x + k_y y) \exp - \left\{ \frac{x^2 + y^2}{2\sigma^2} \right\}$$

# Gabor Filters



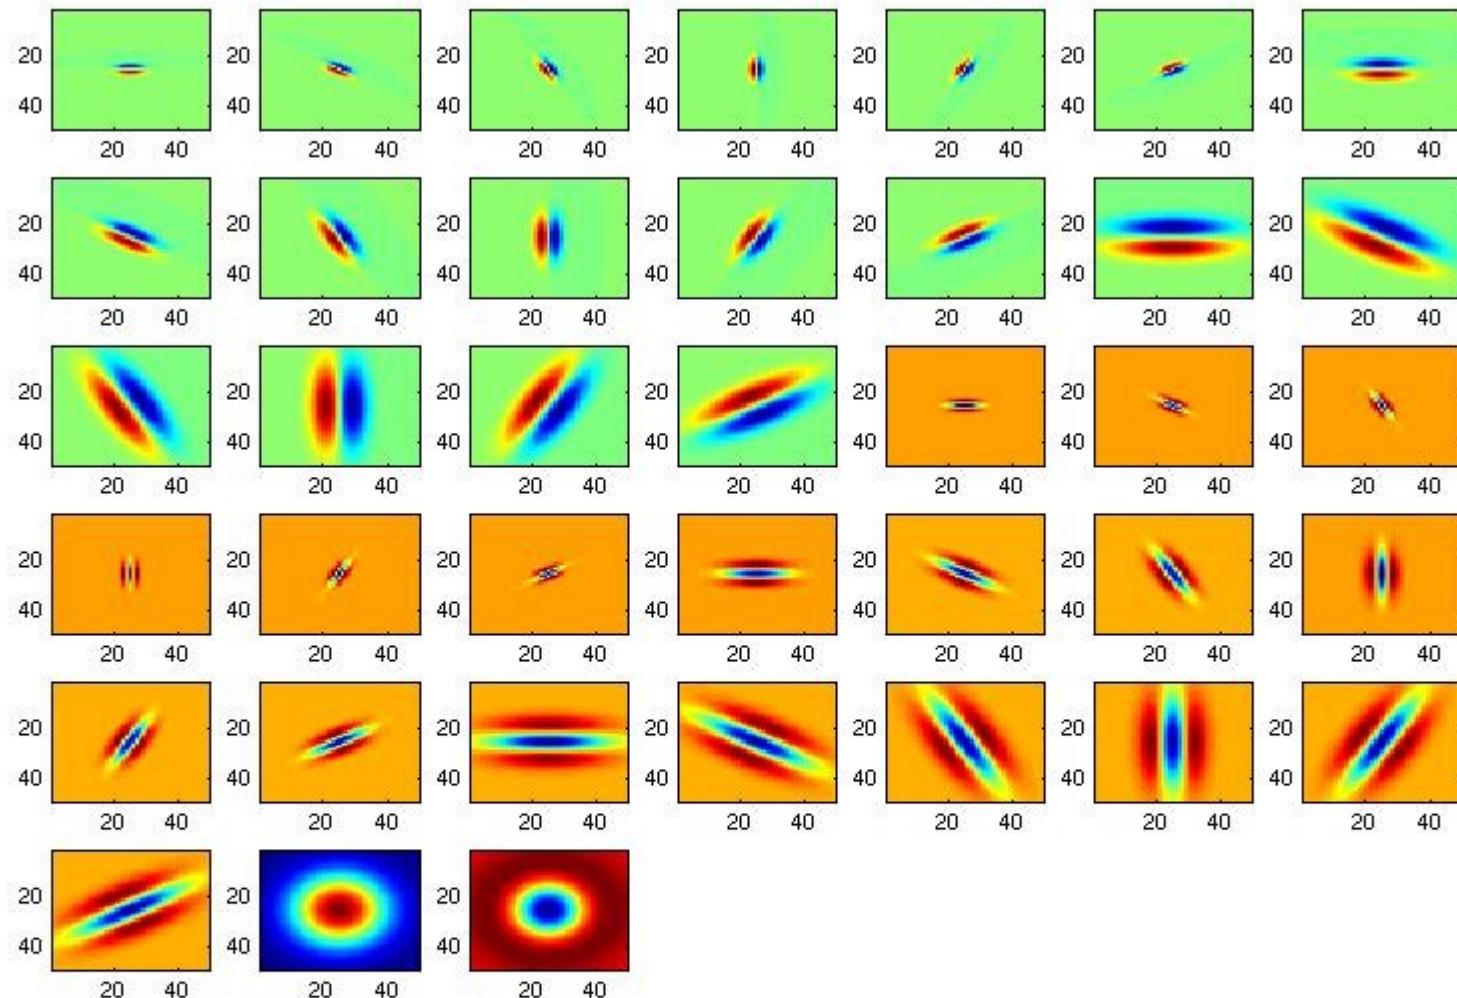
Gabor filters at different scales and spatial frequencies

top row shows anti-symmetric (or odd) filters, bottom row the symmetric (or even) filters.

$$\text{symmetric: } \cos(k_x x + k_y y) \exp - \left\{ \frac{x^2 + y^2}{2\sigma^2} \right\}$$

$$\text{antisymmetric : } \sin(k_x x + k_y y) \exp - \left\{ \frac{x^2 + y^2}{2\sigma^2} \right\}$$

# Filter bank



# Gabor Filters

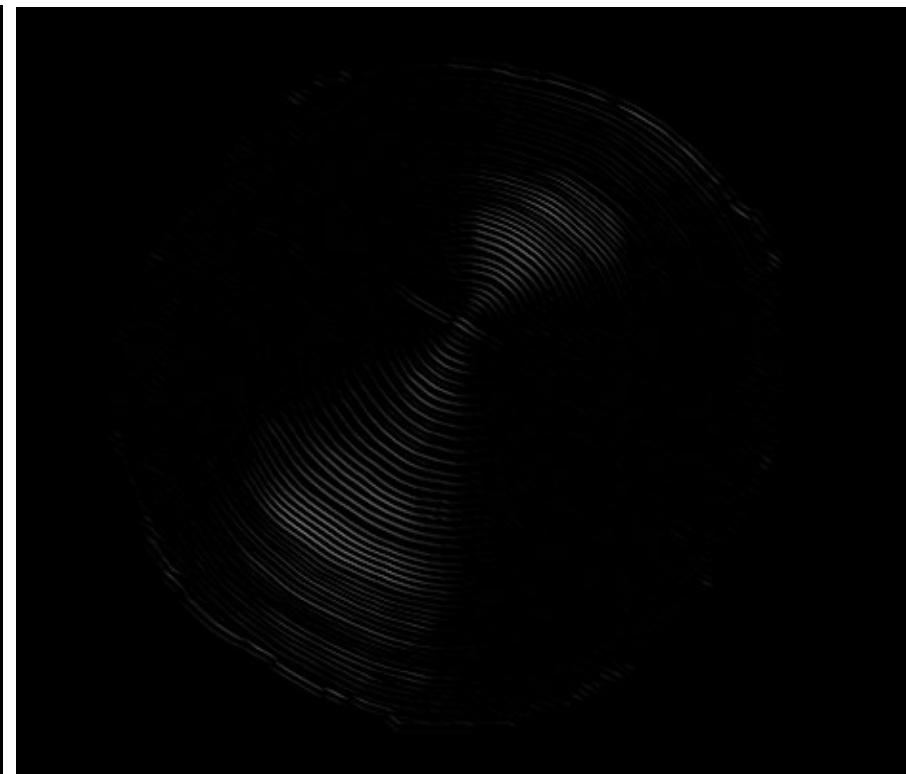
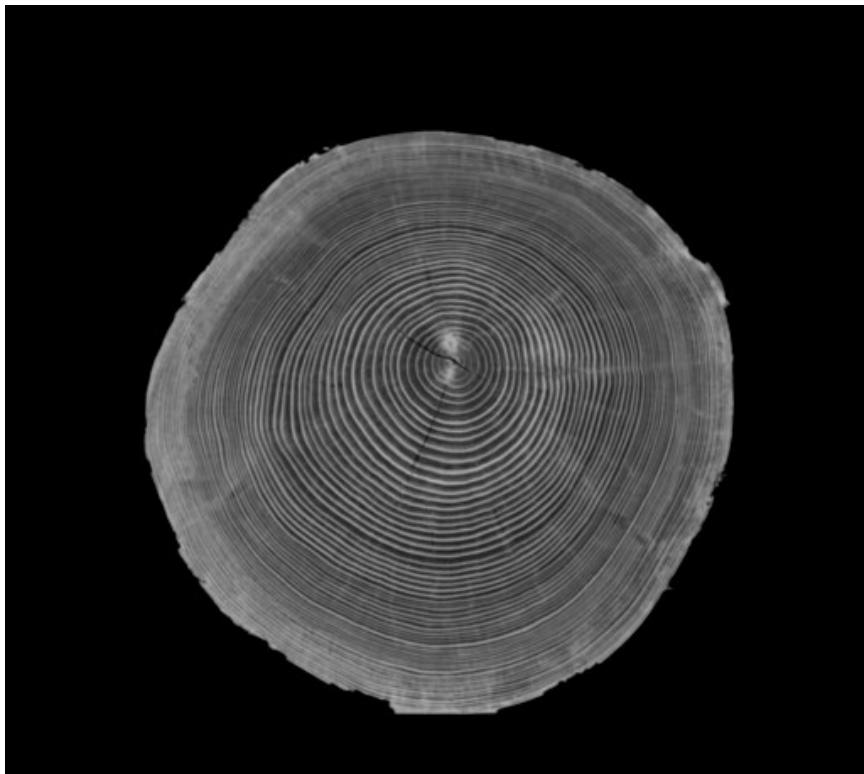
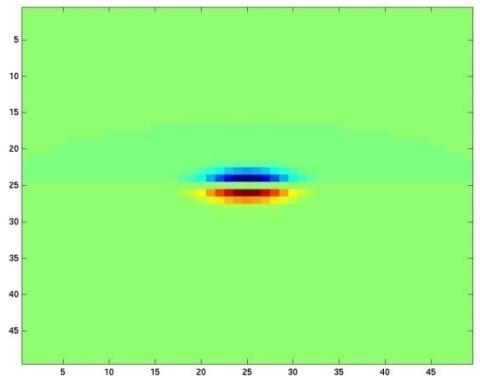
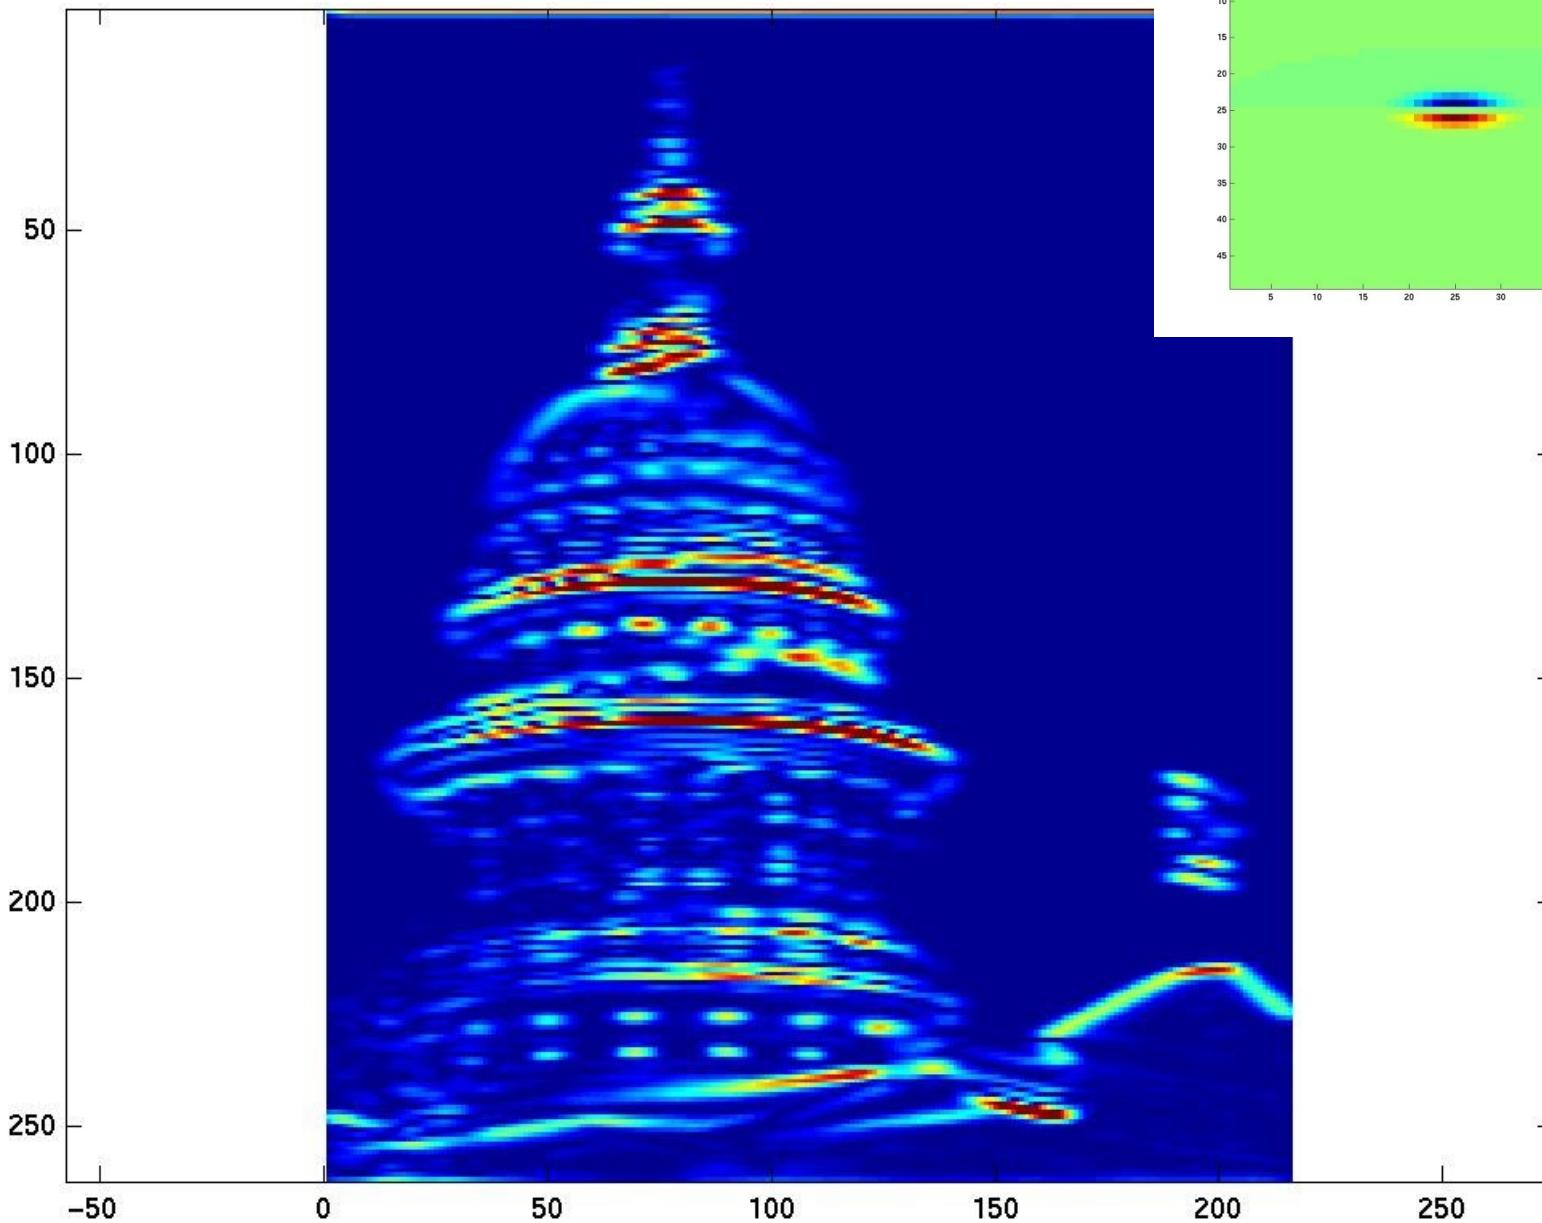
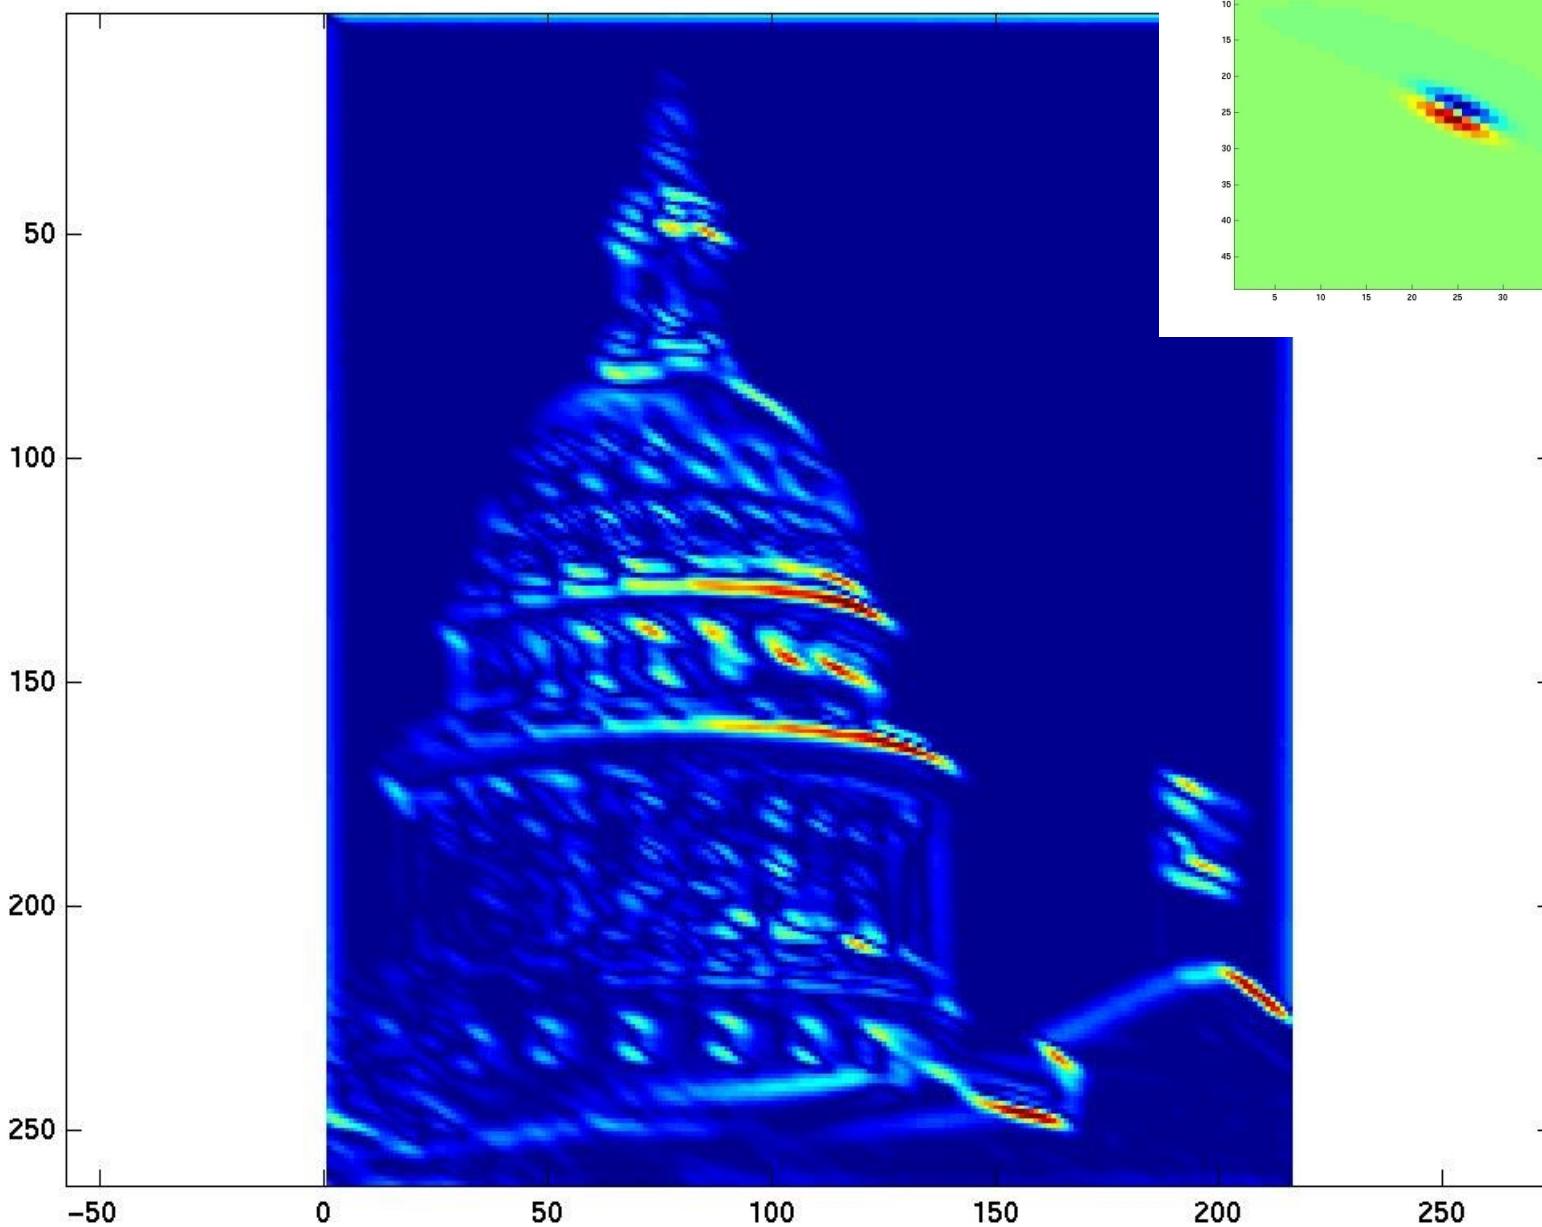


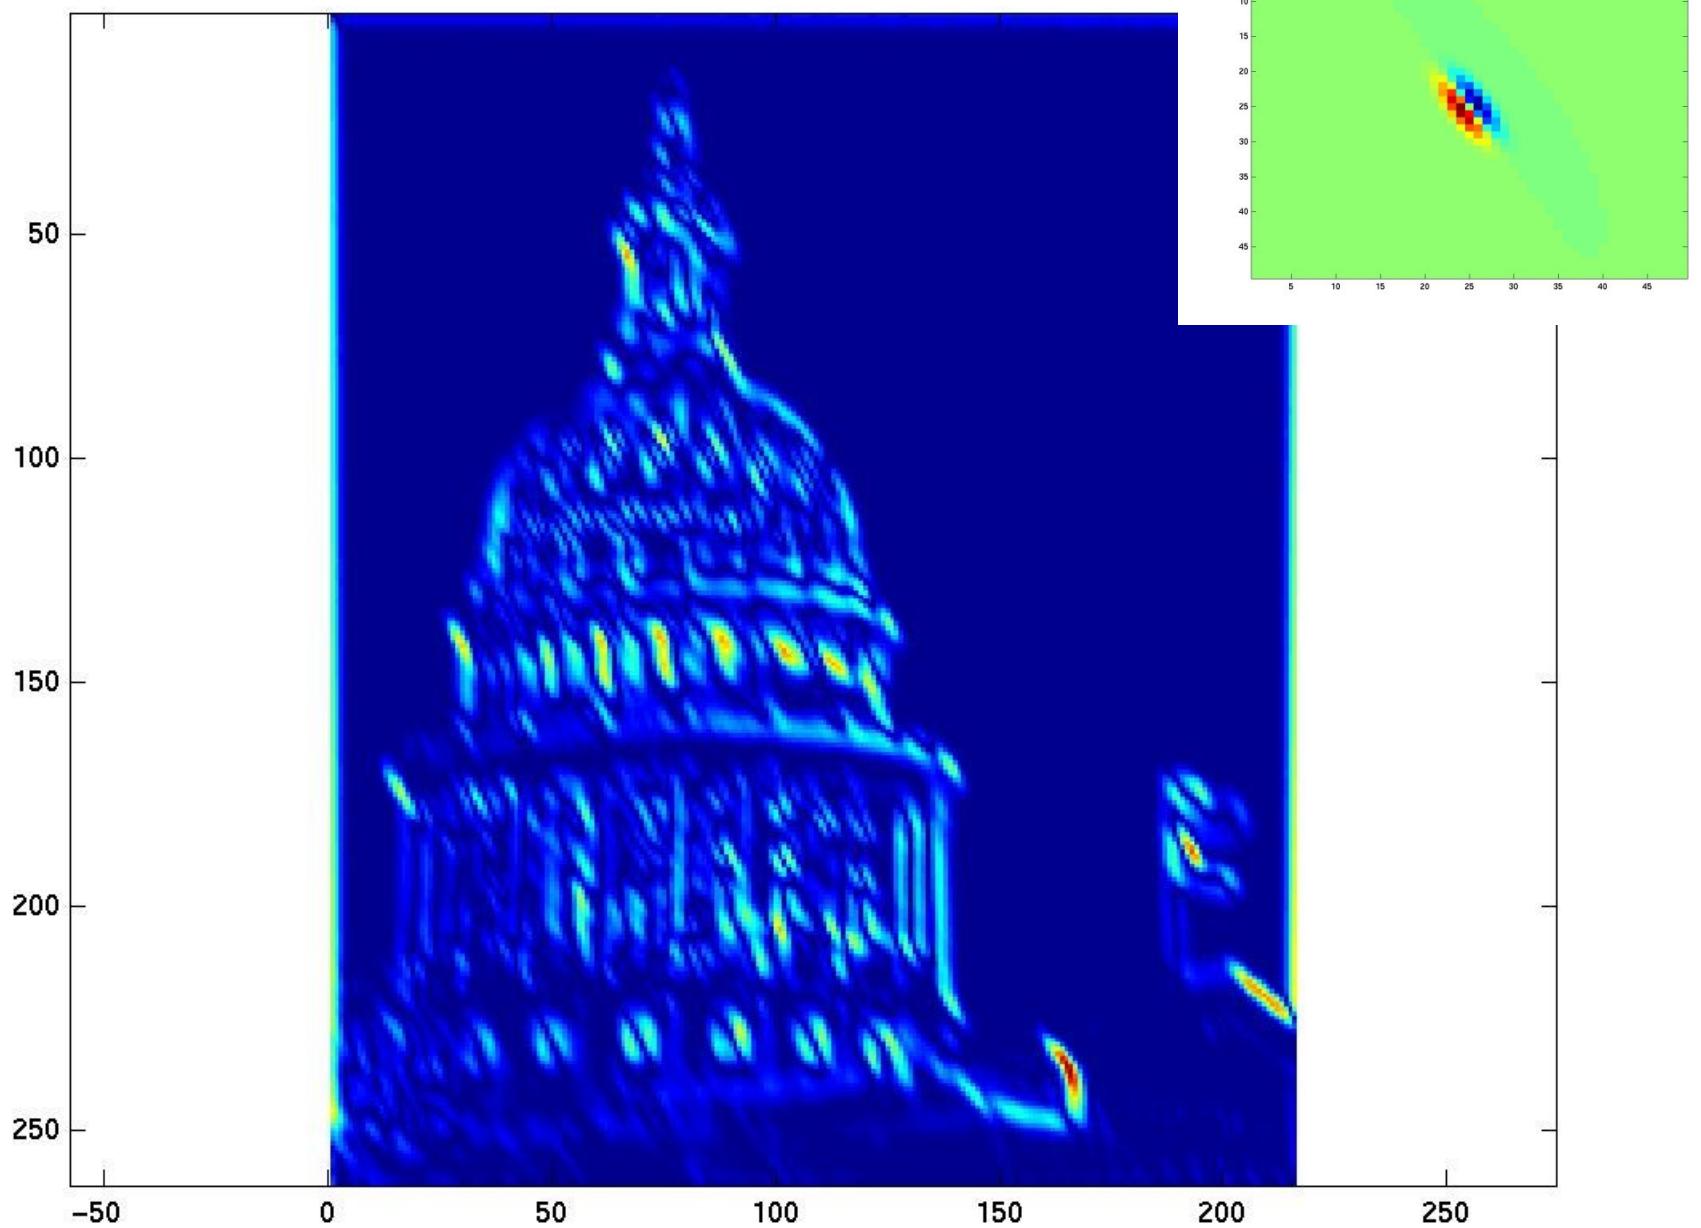
 Image from <http://www.texasexplorer.com/austincap2.jpg>

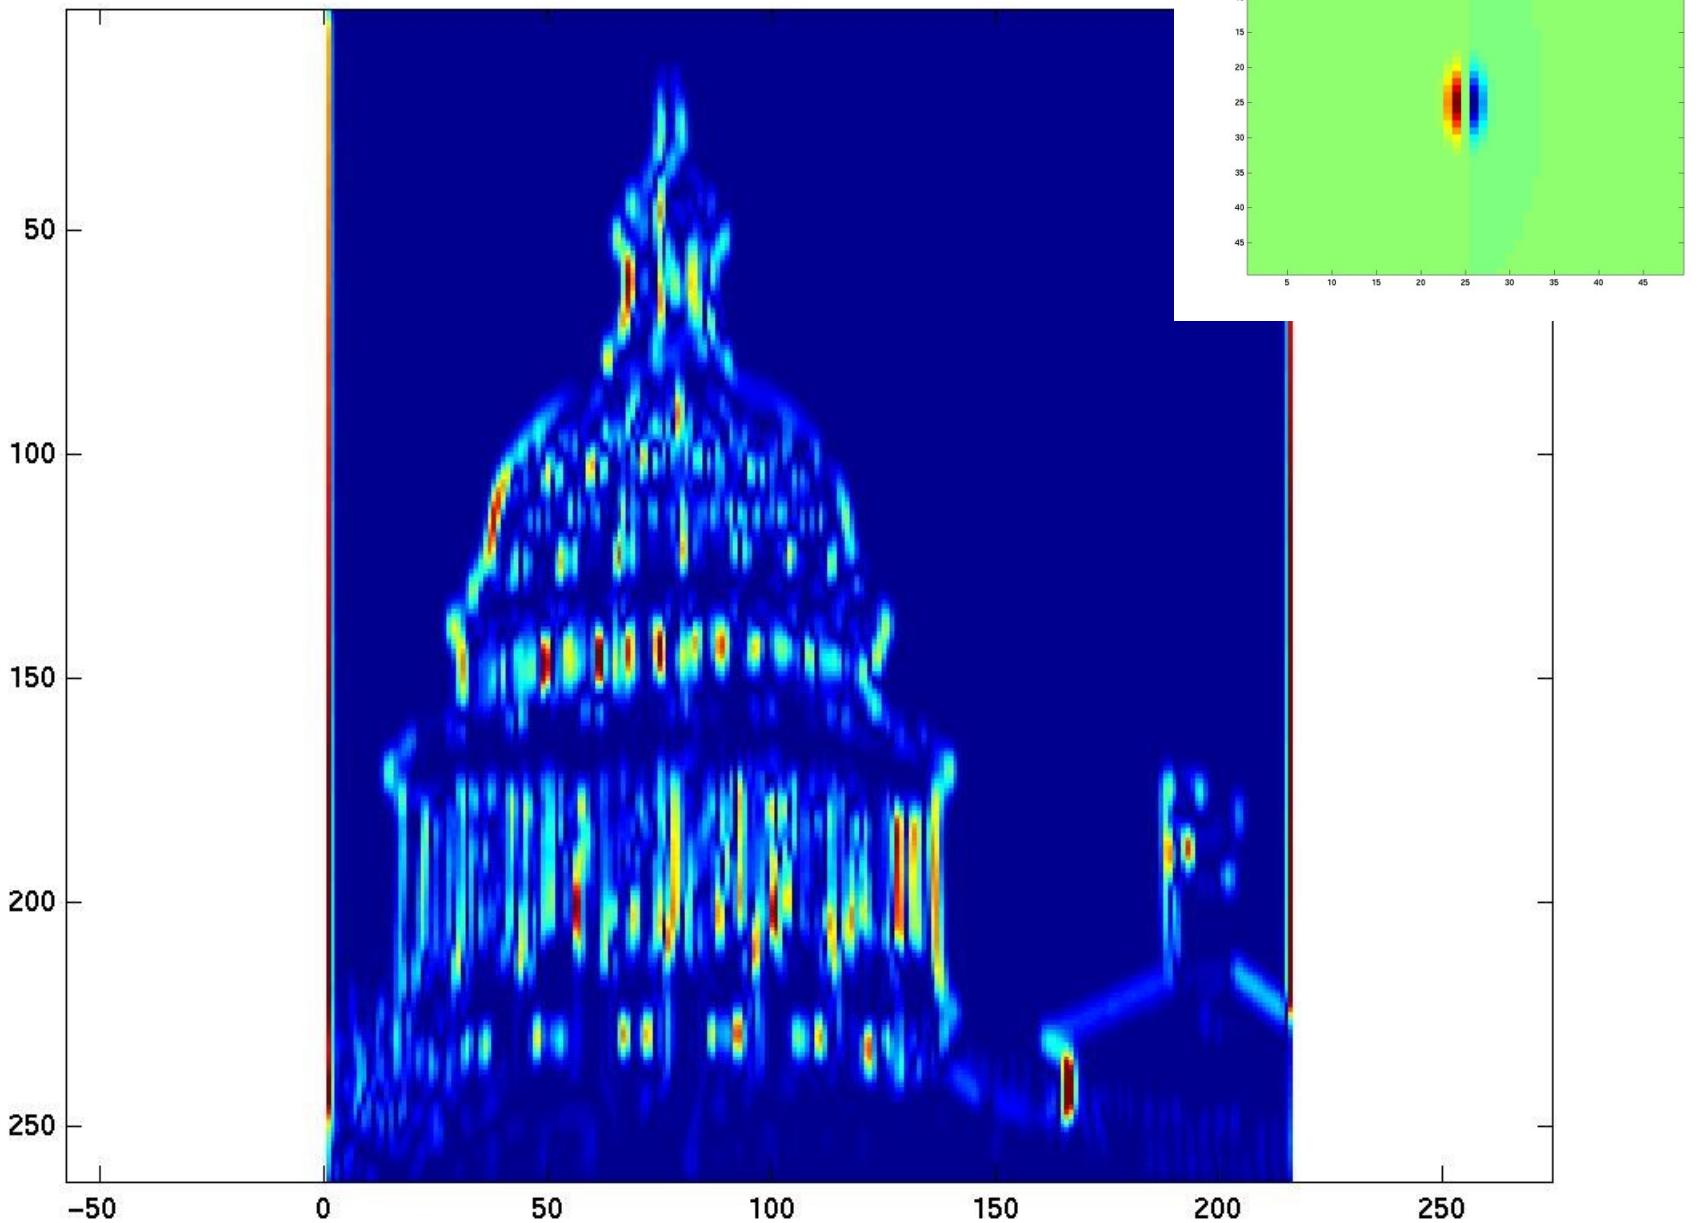


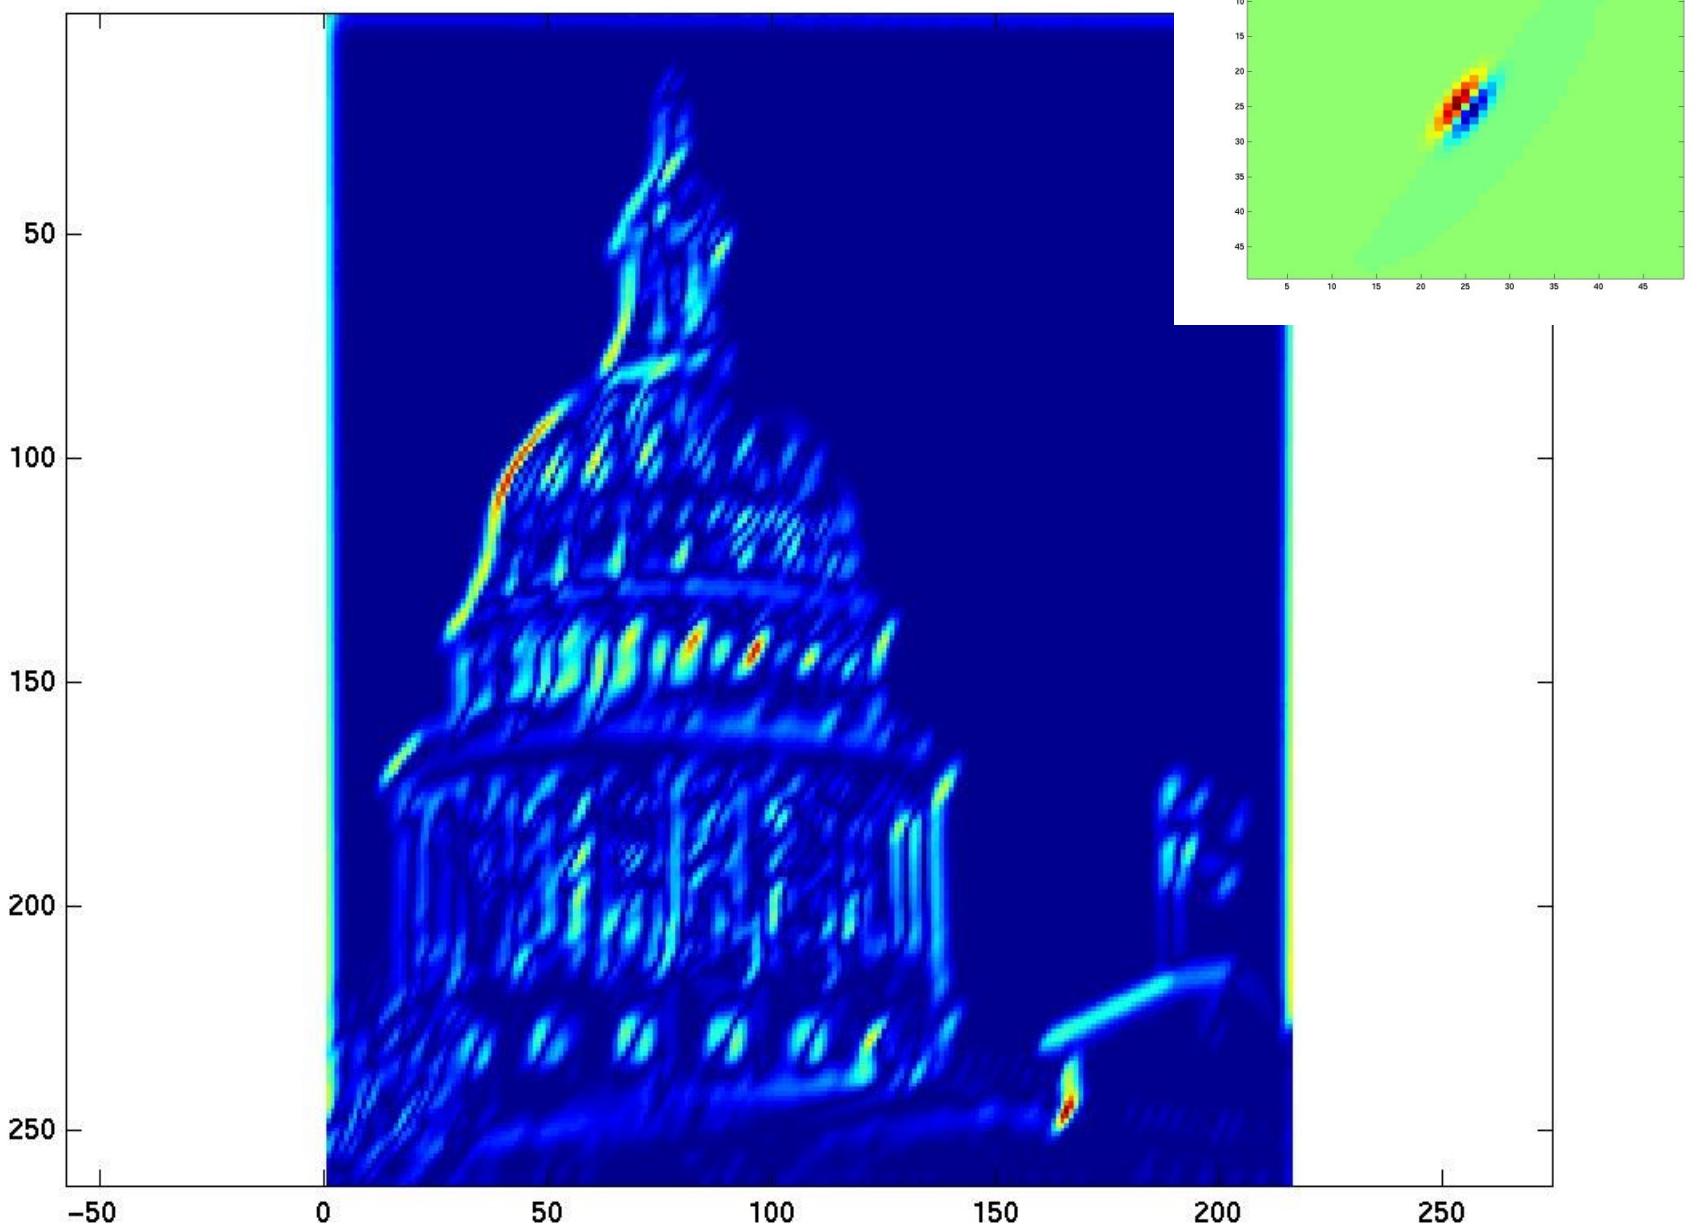
 Kristen Grauman

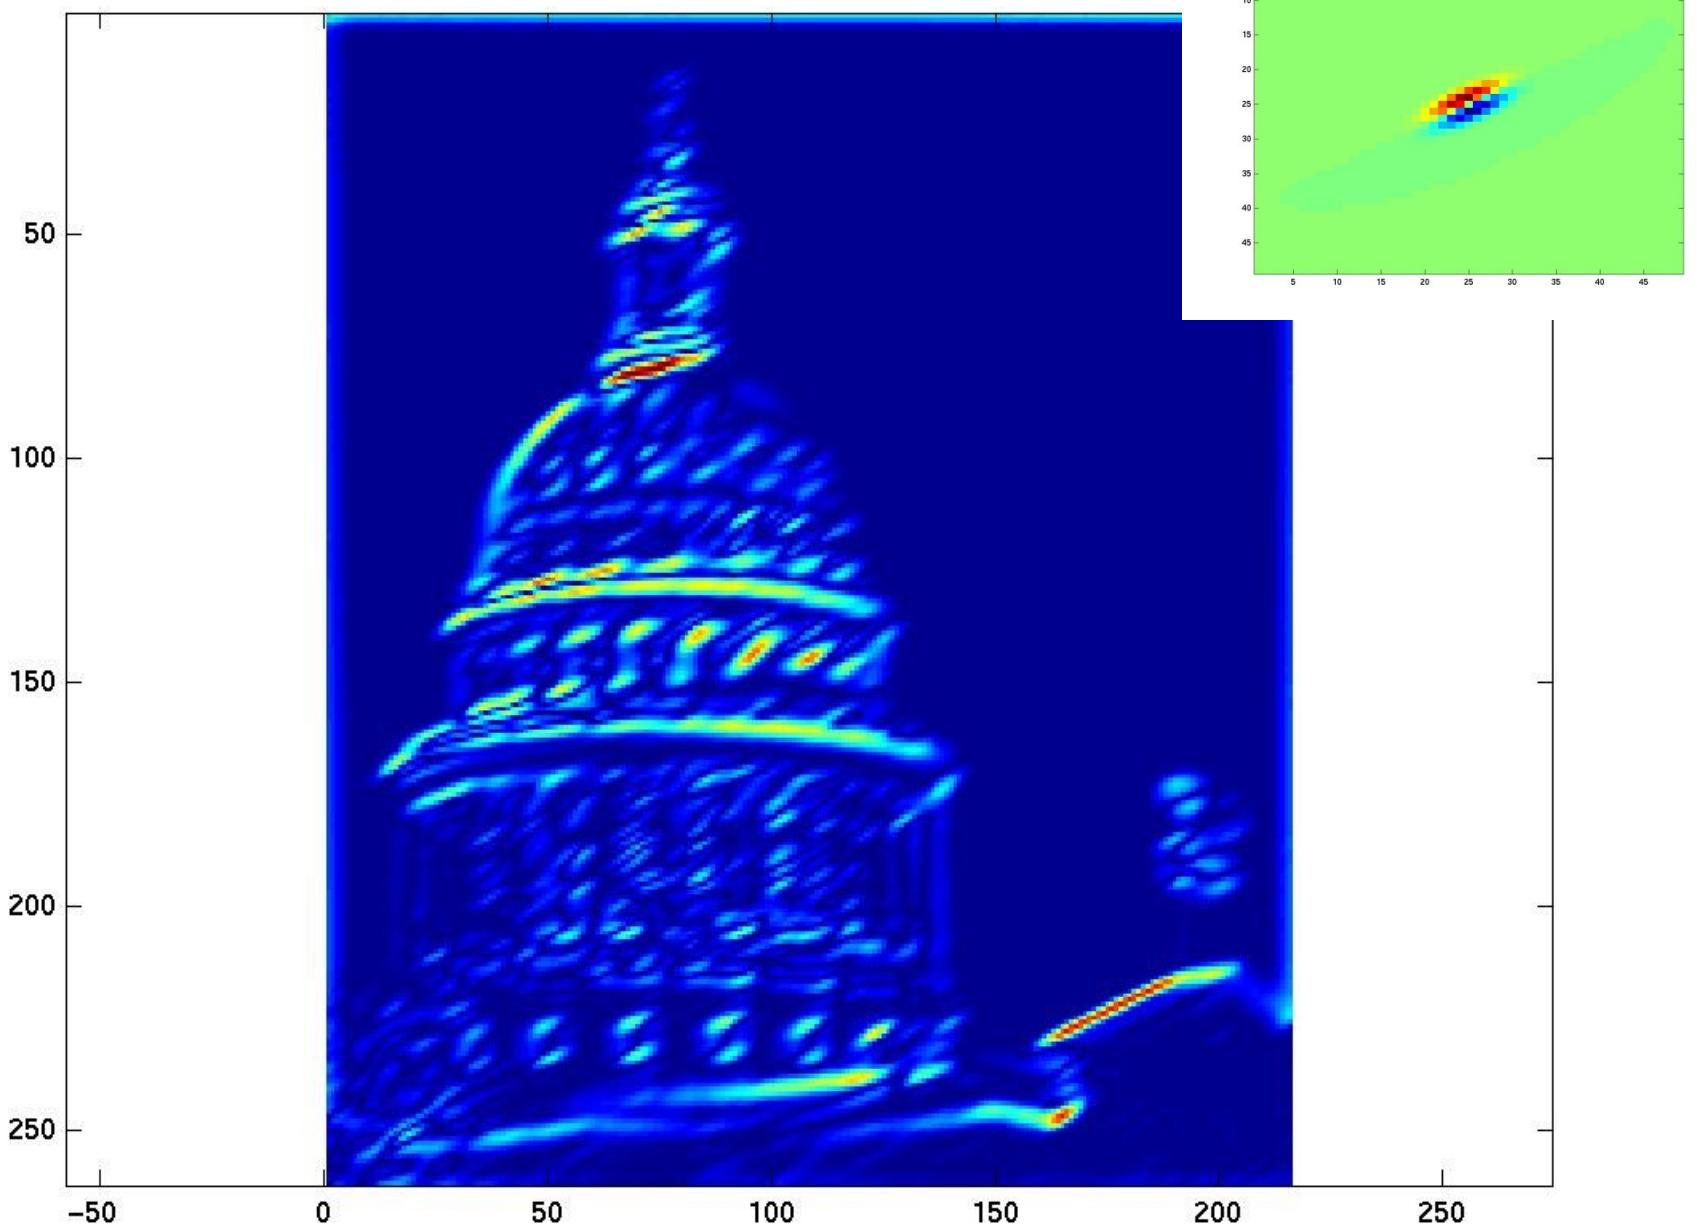


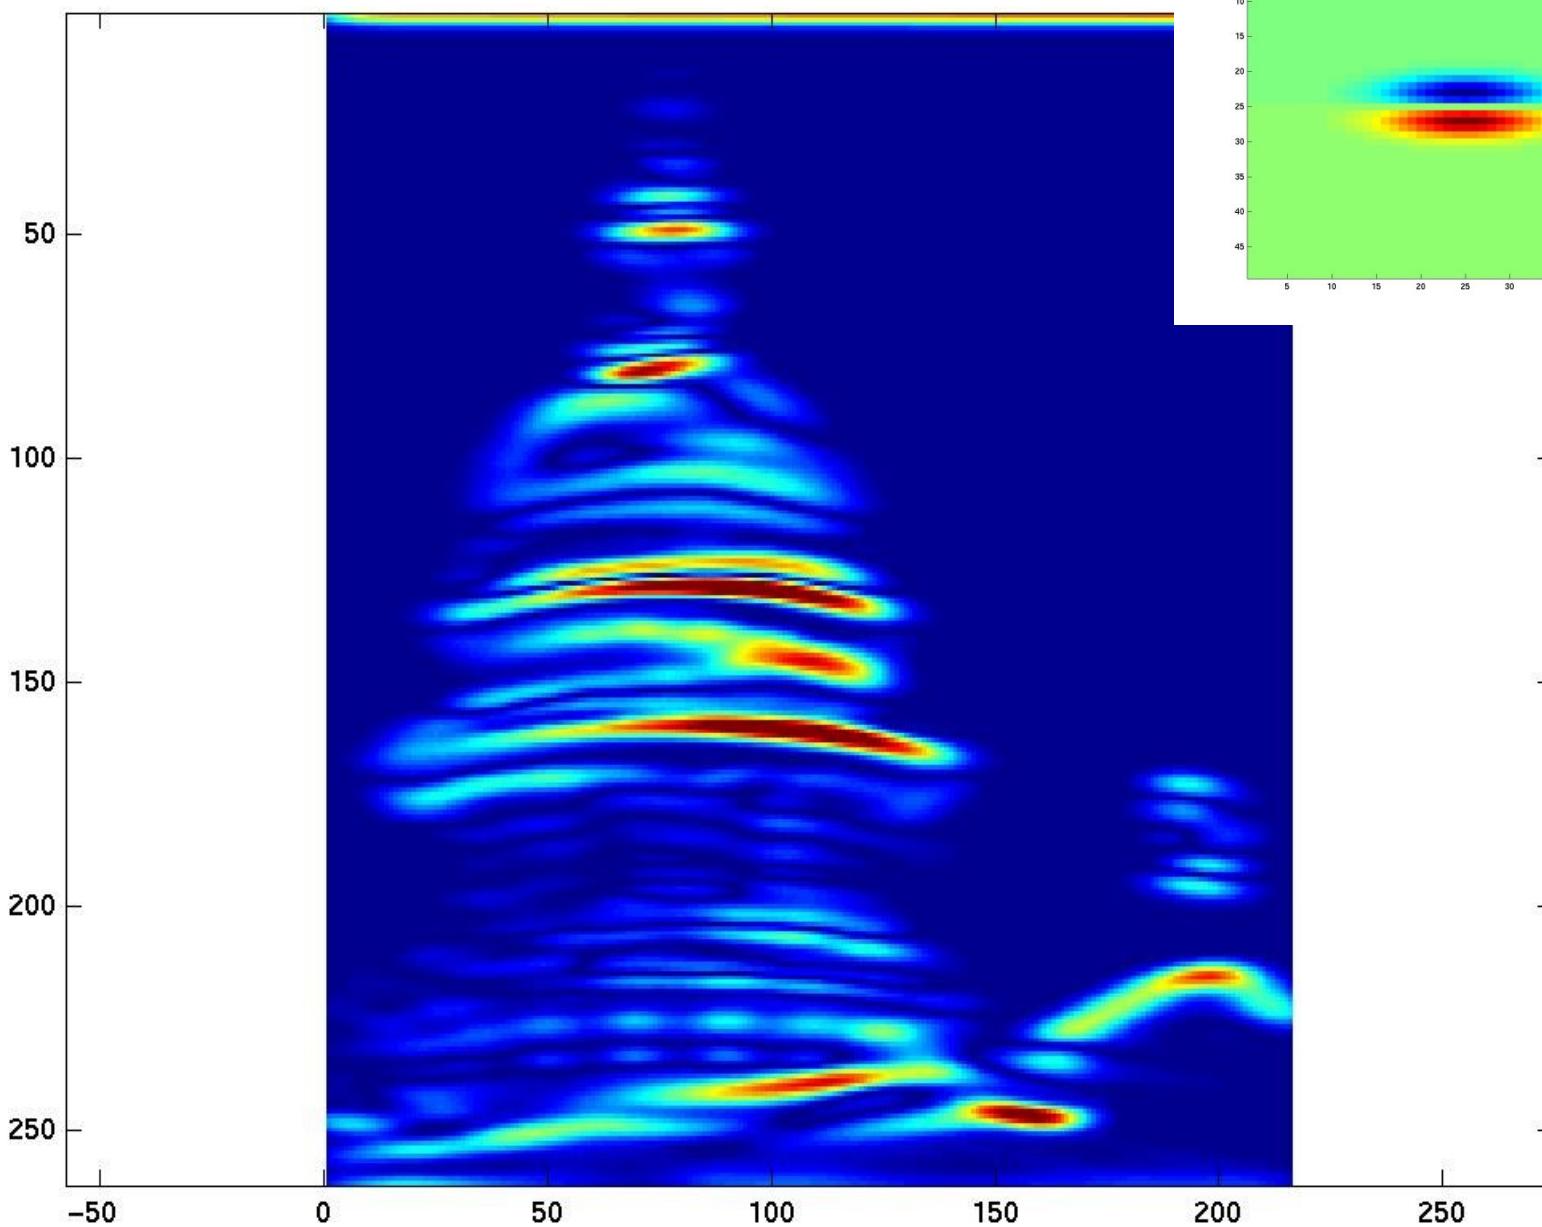


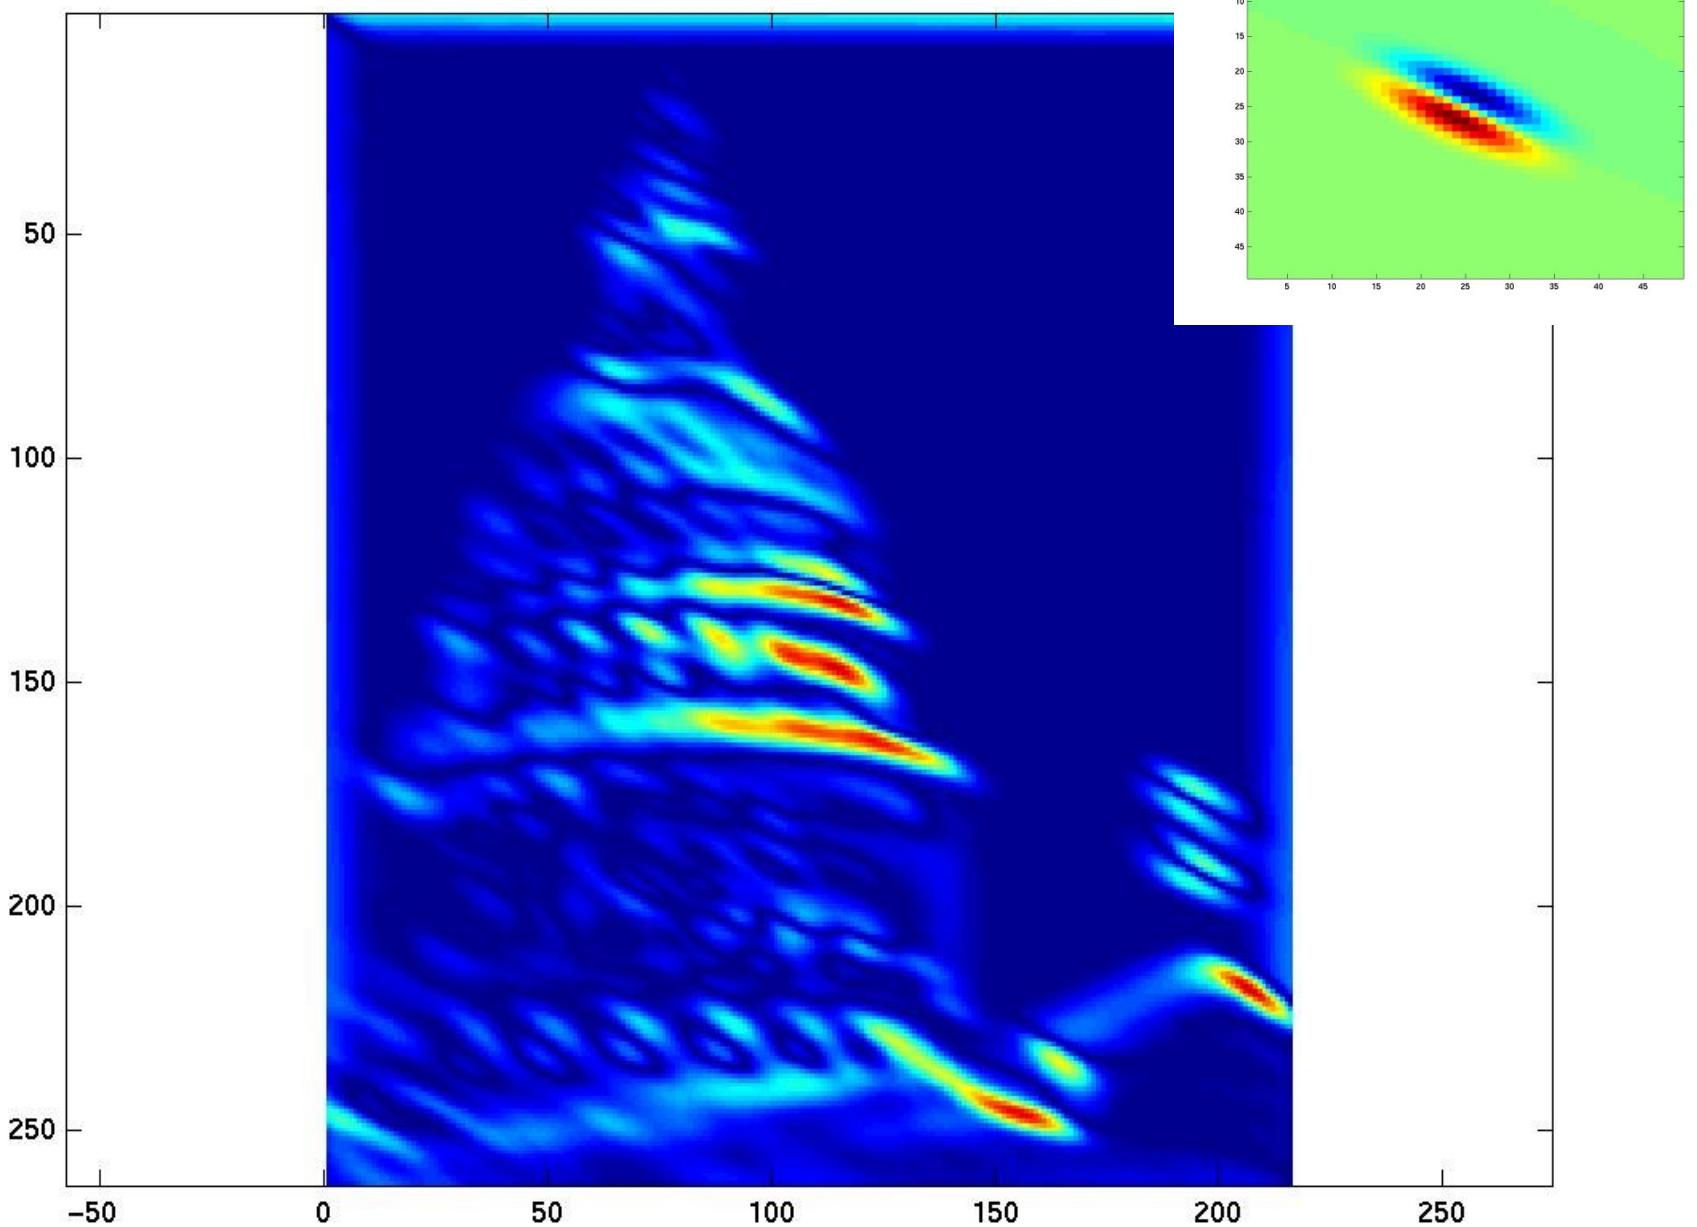


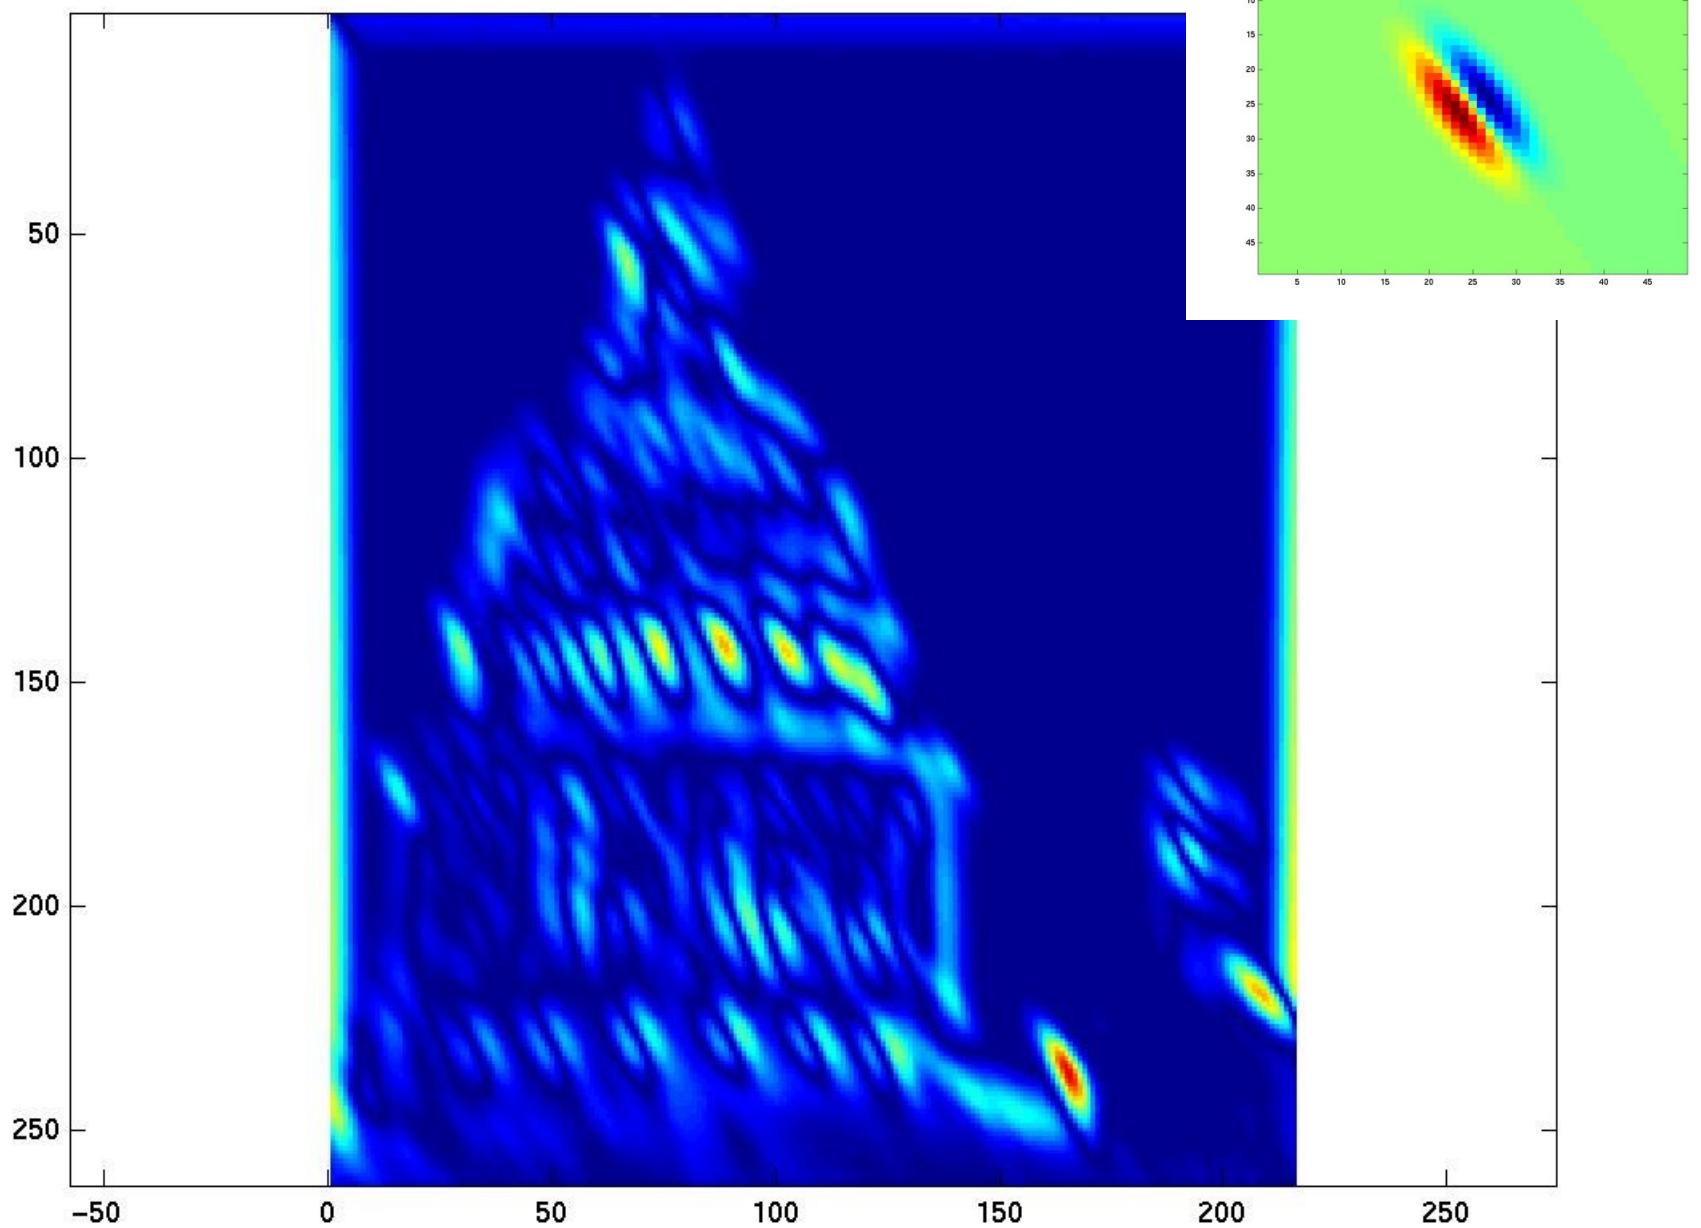


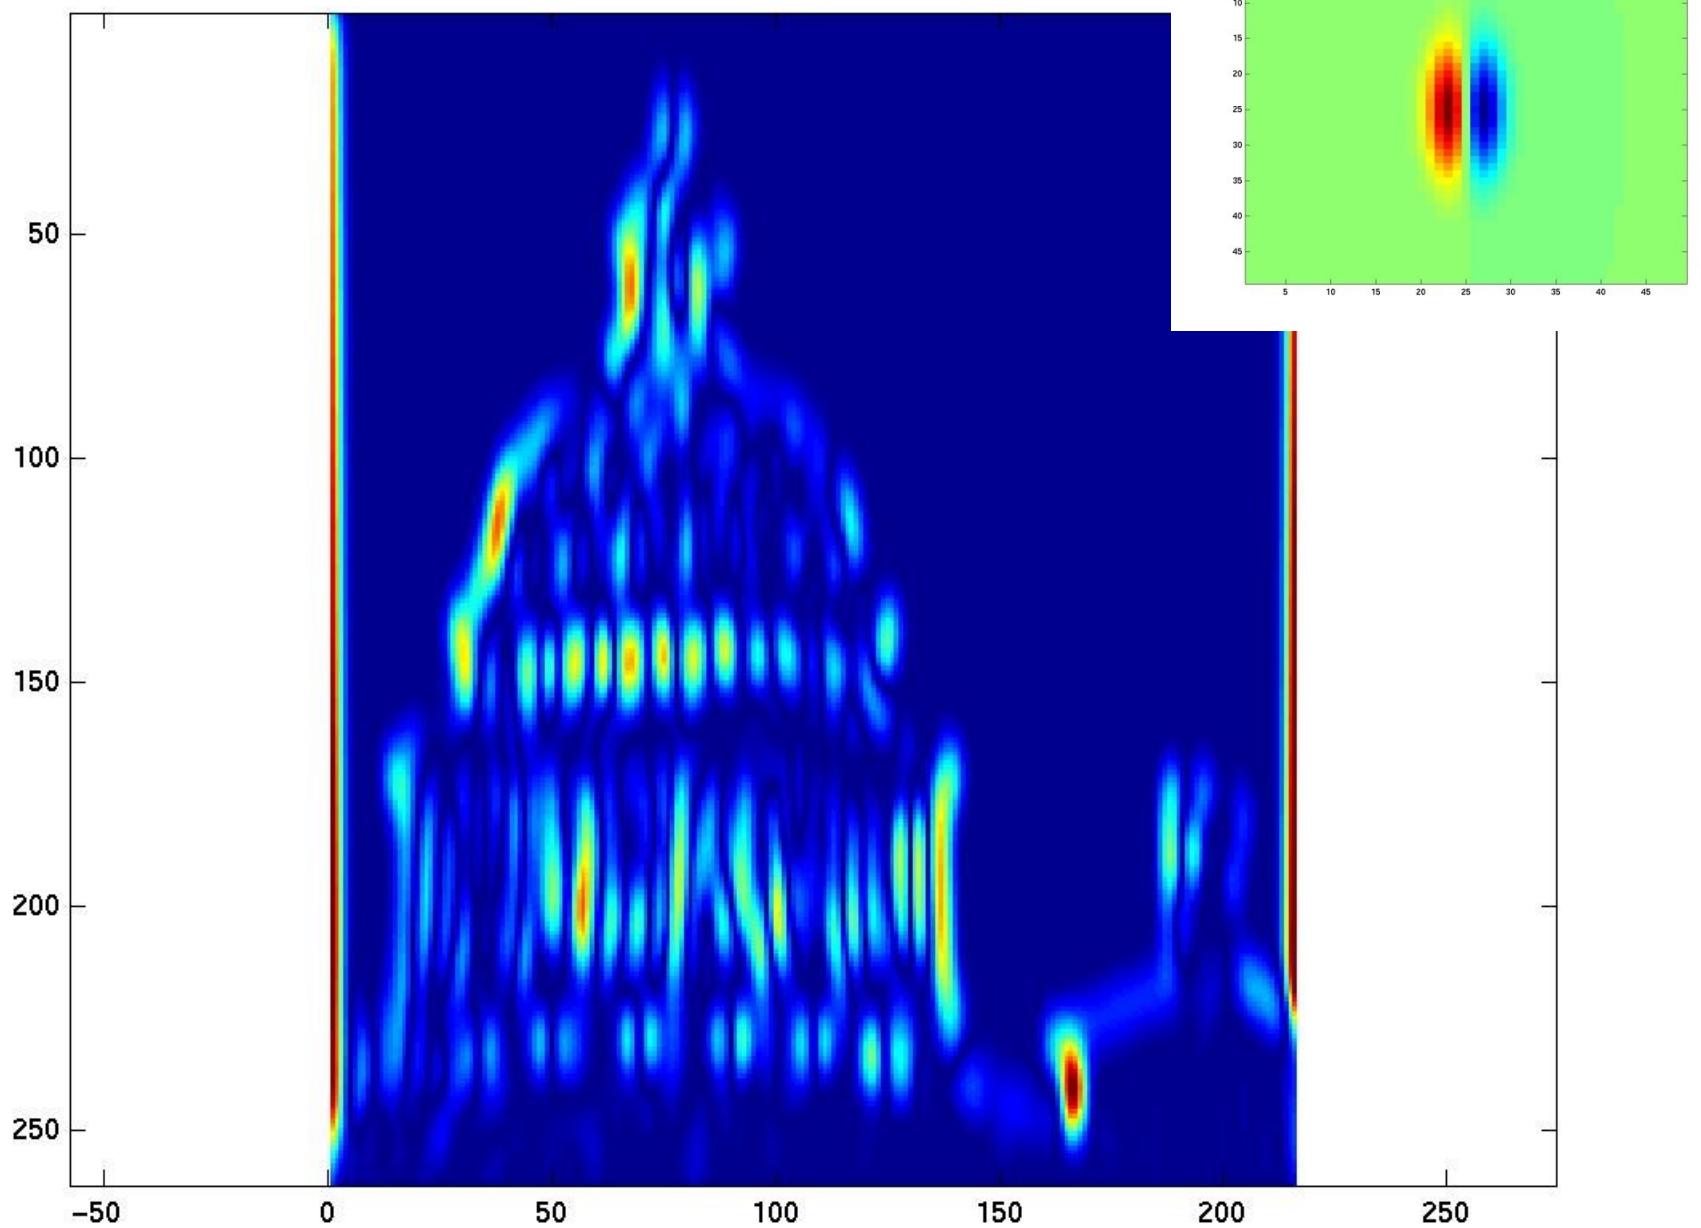


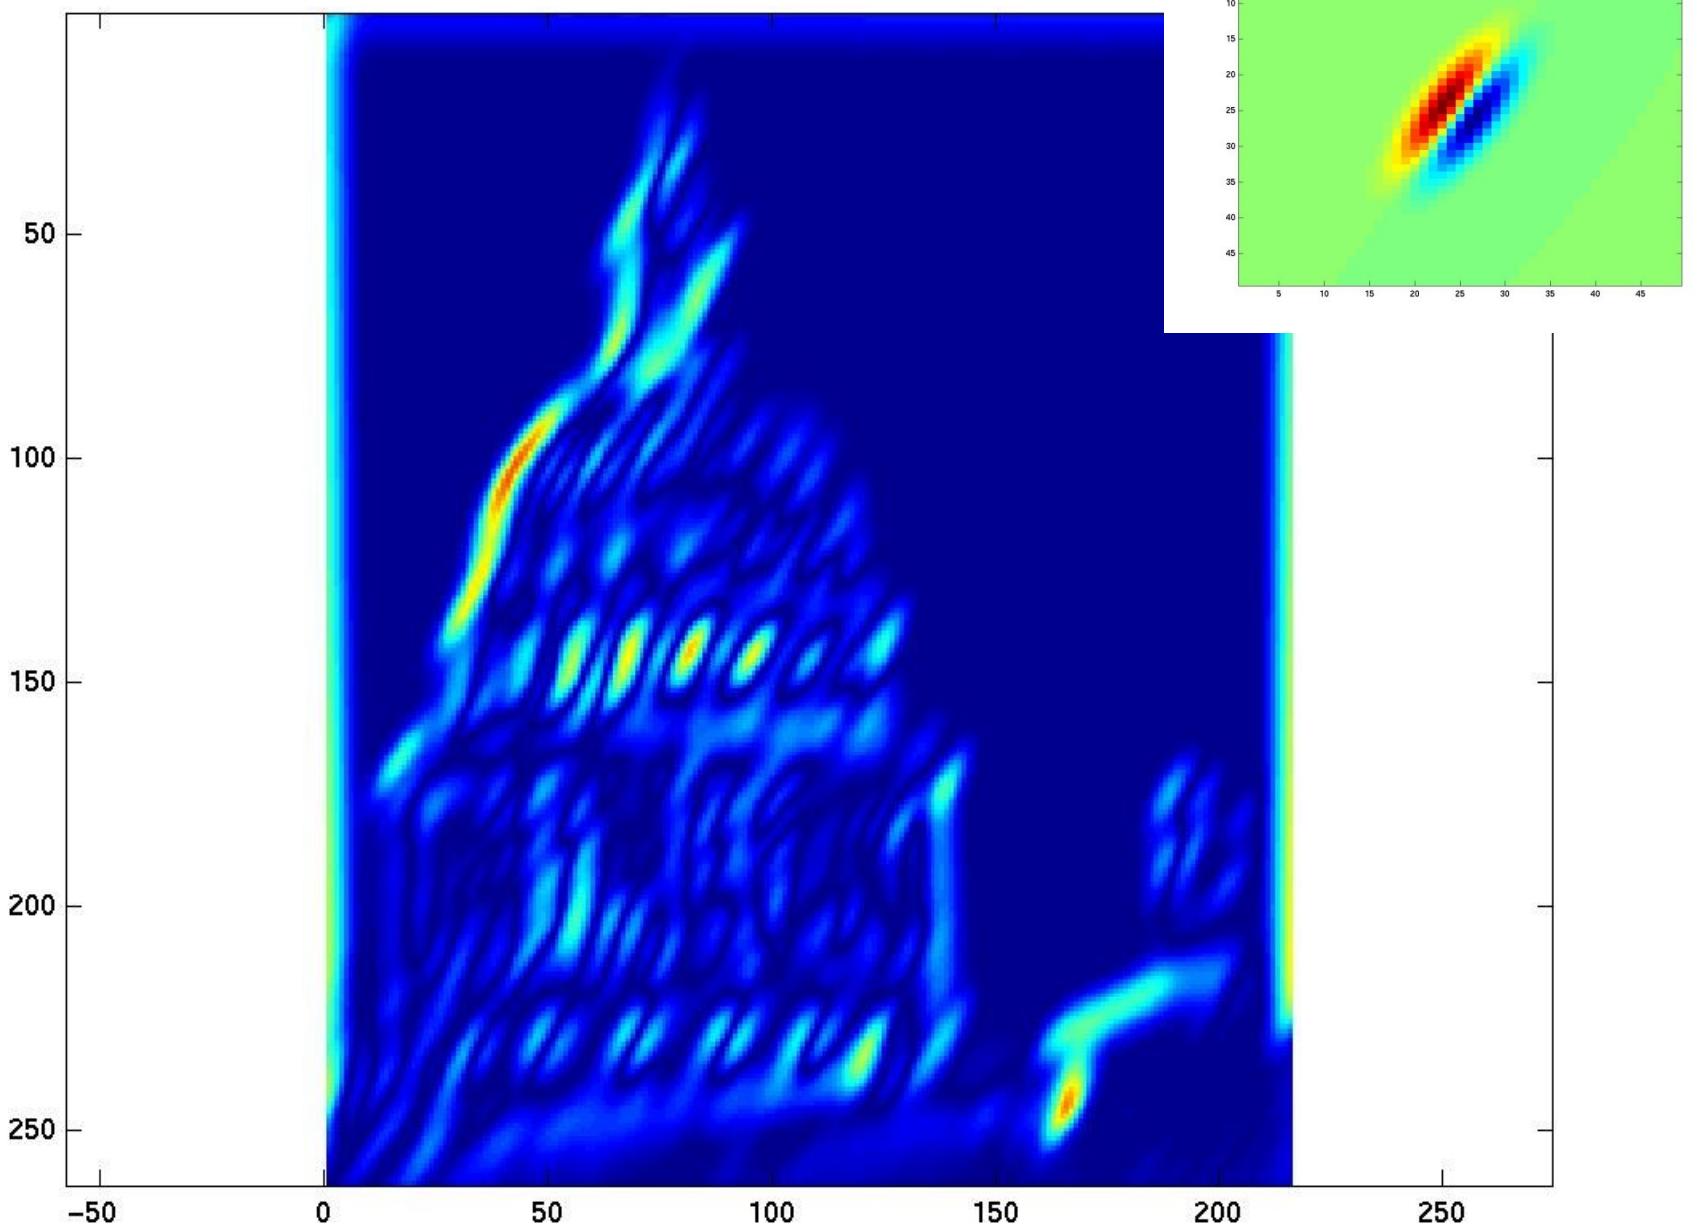


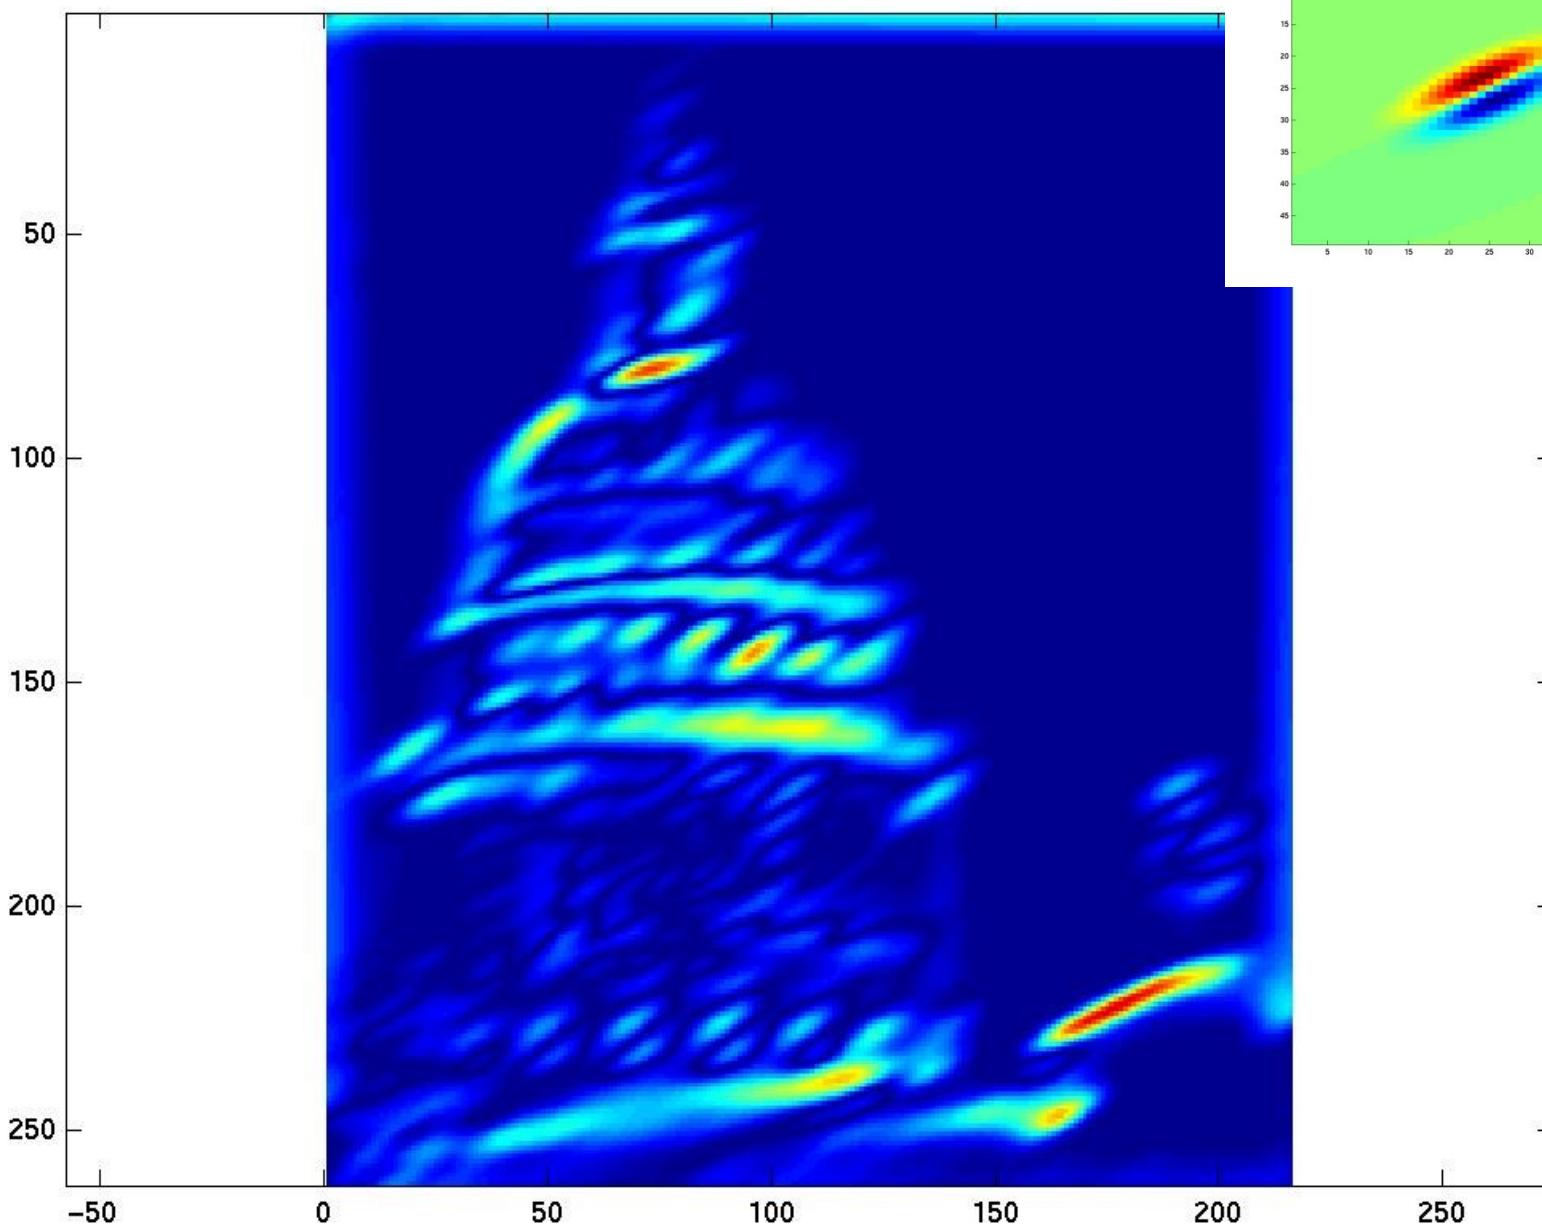


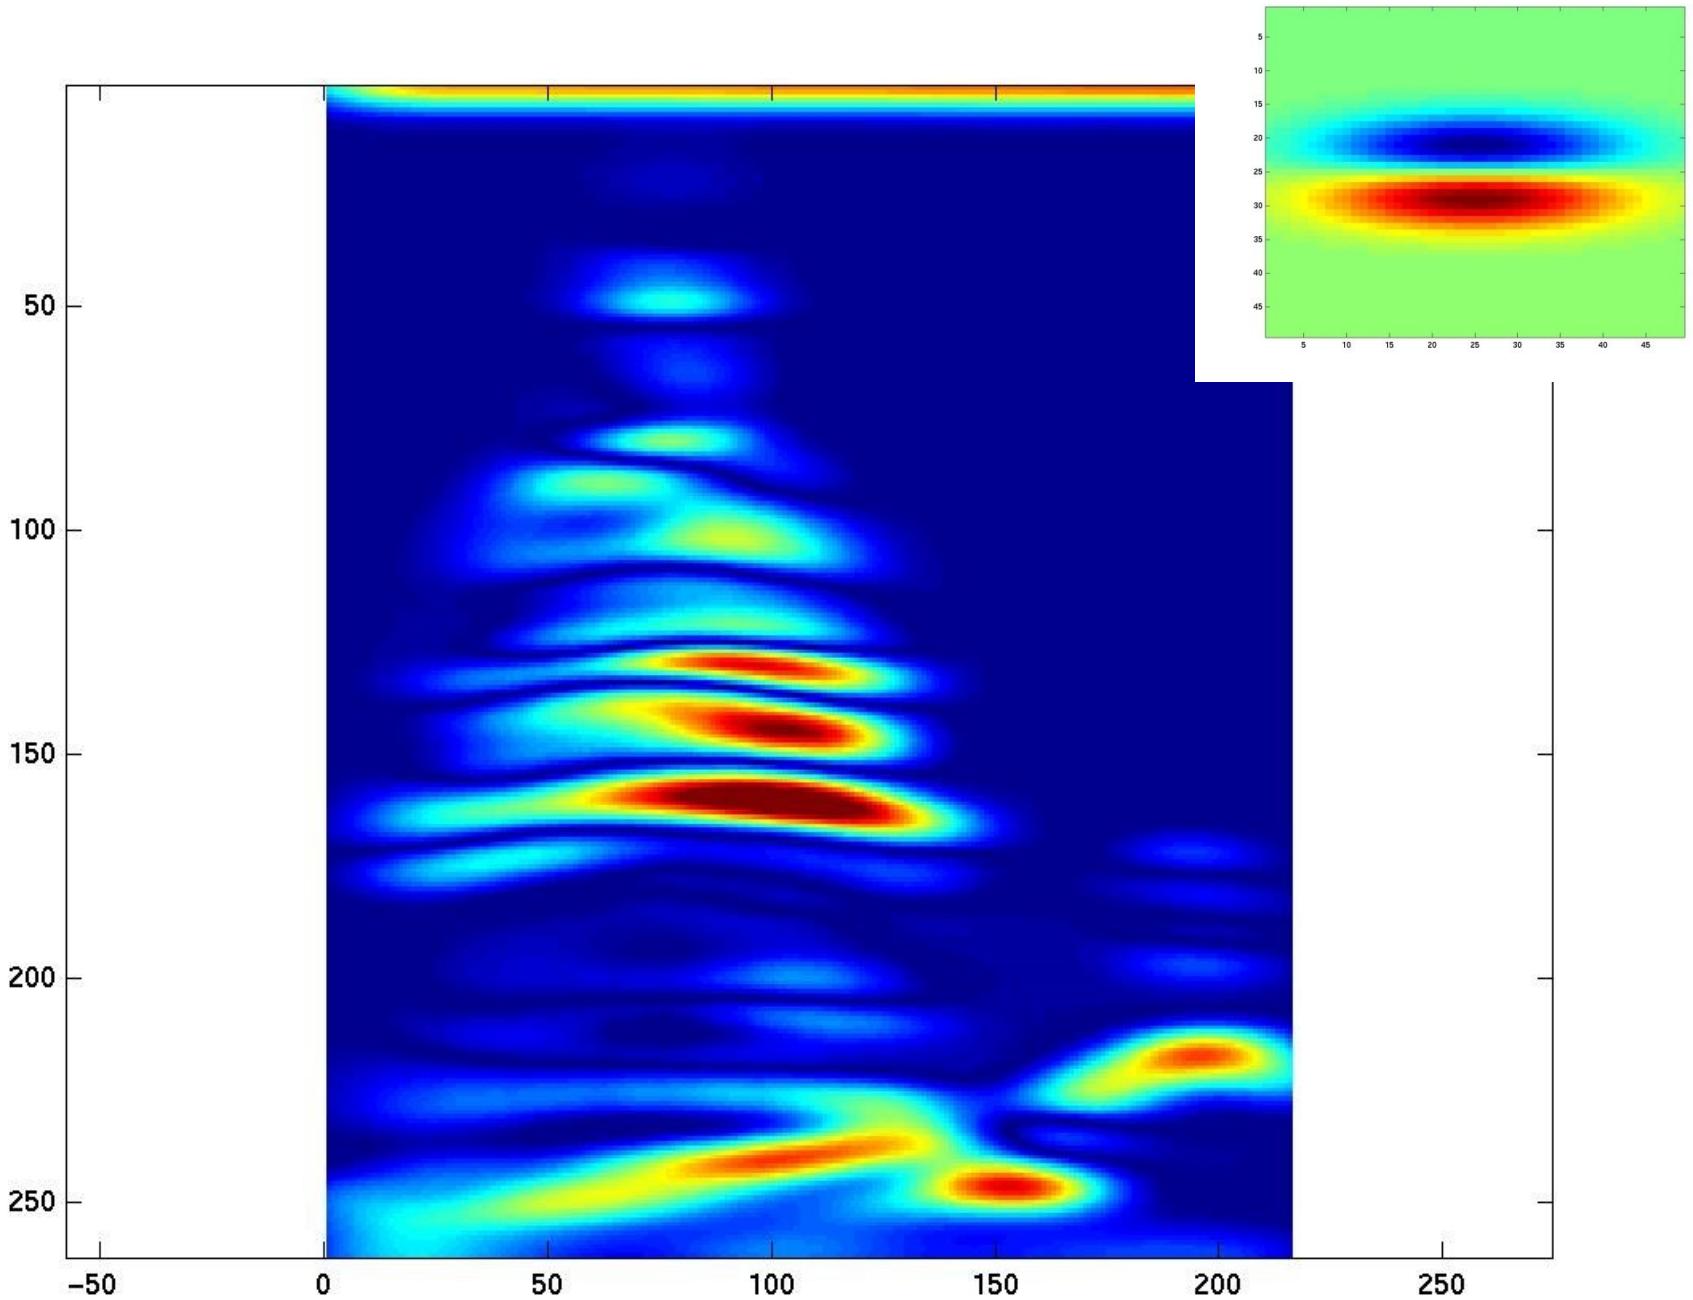


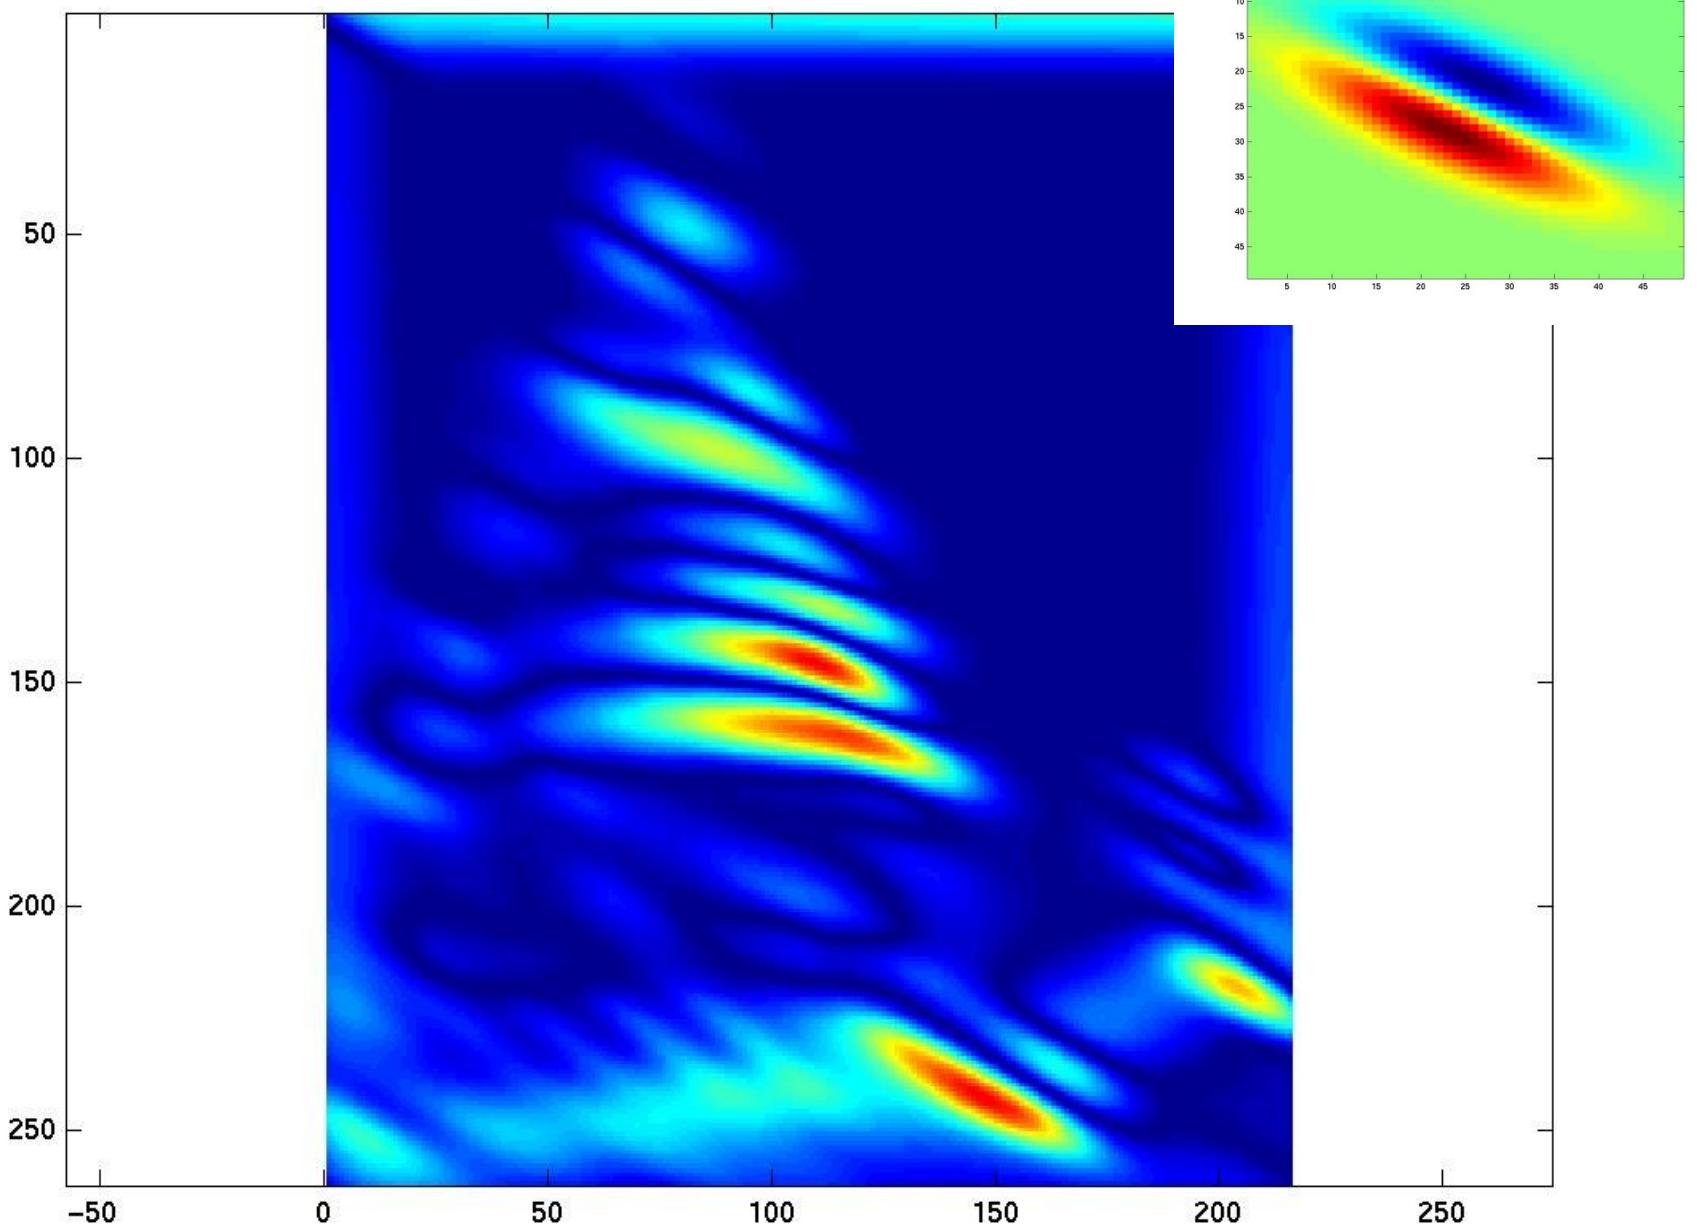


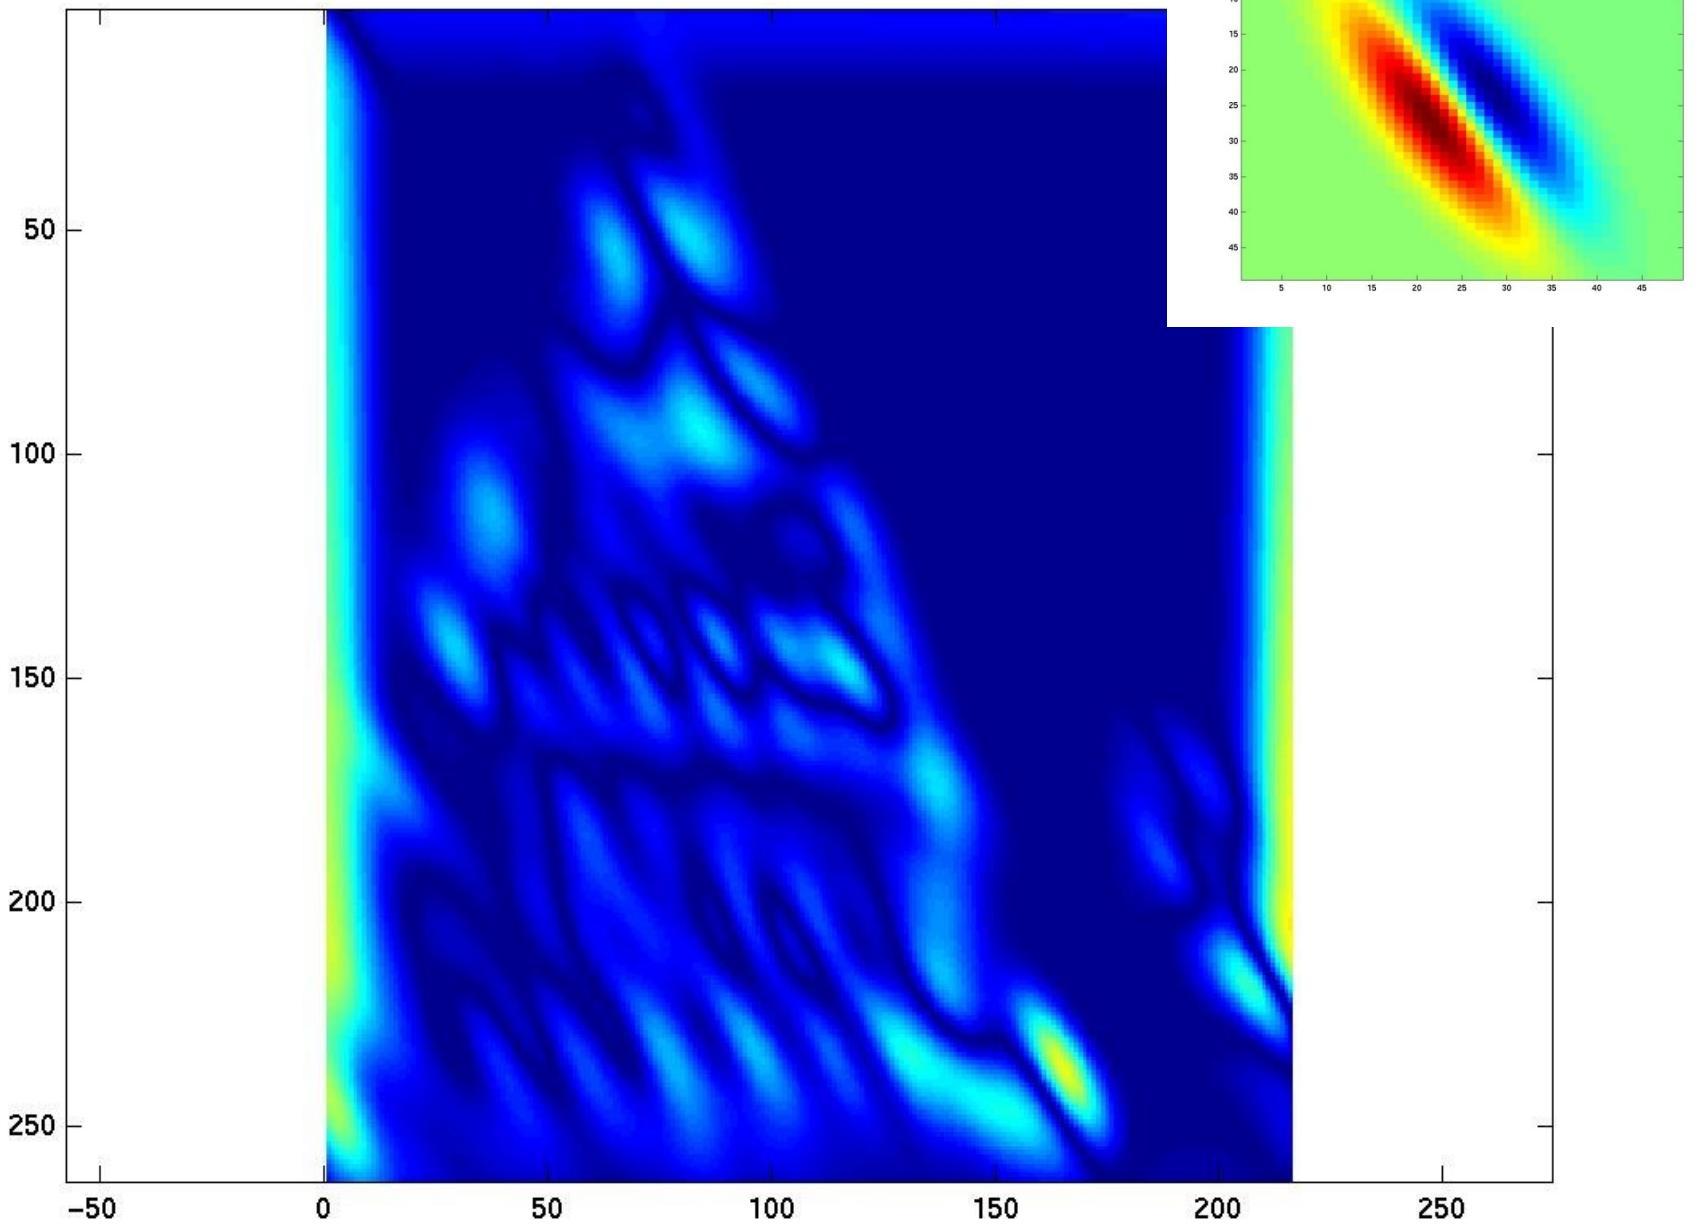


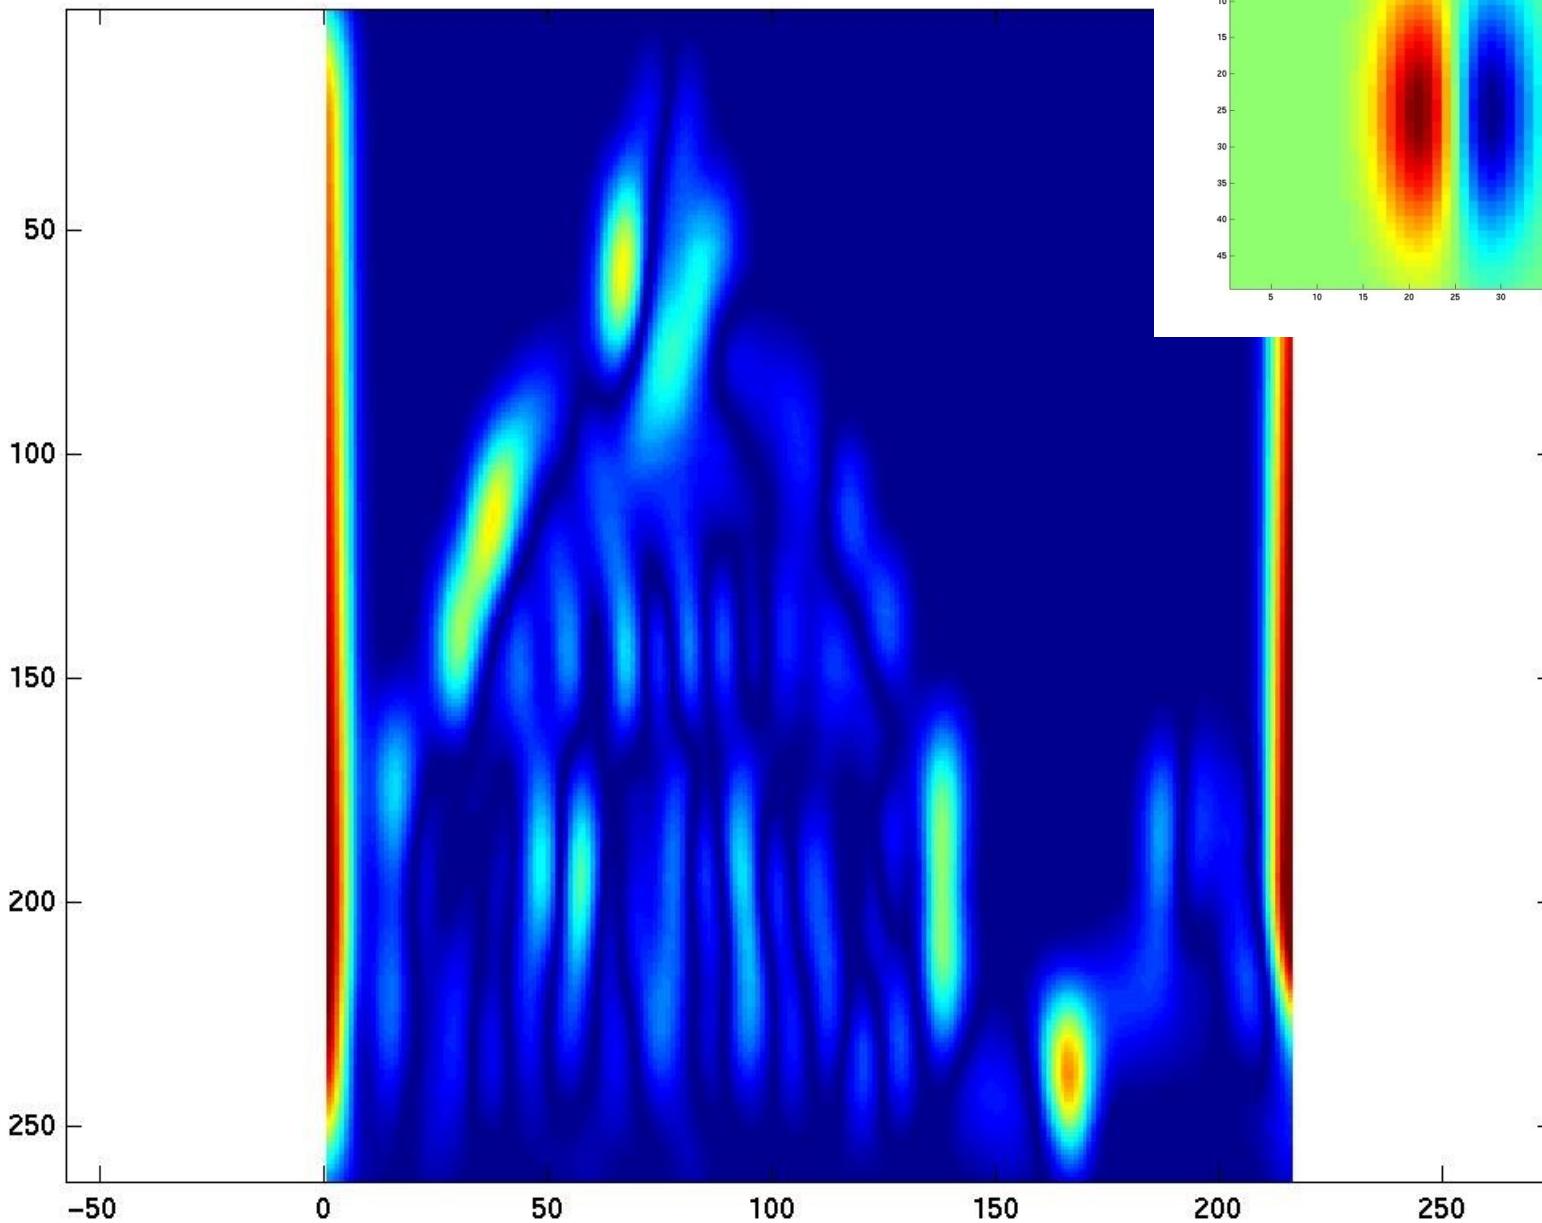


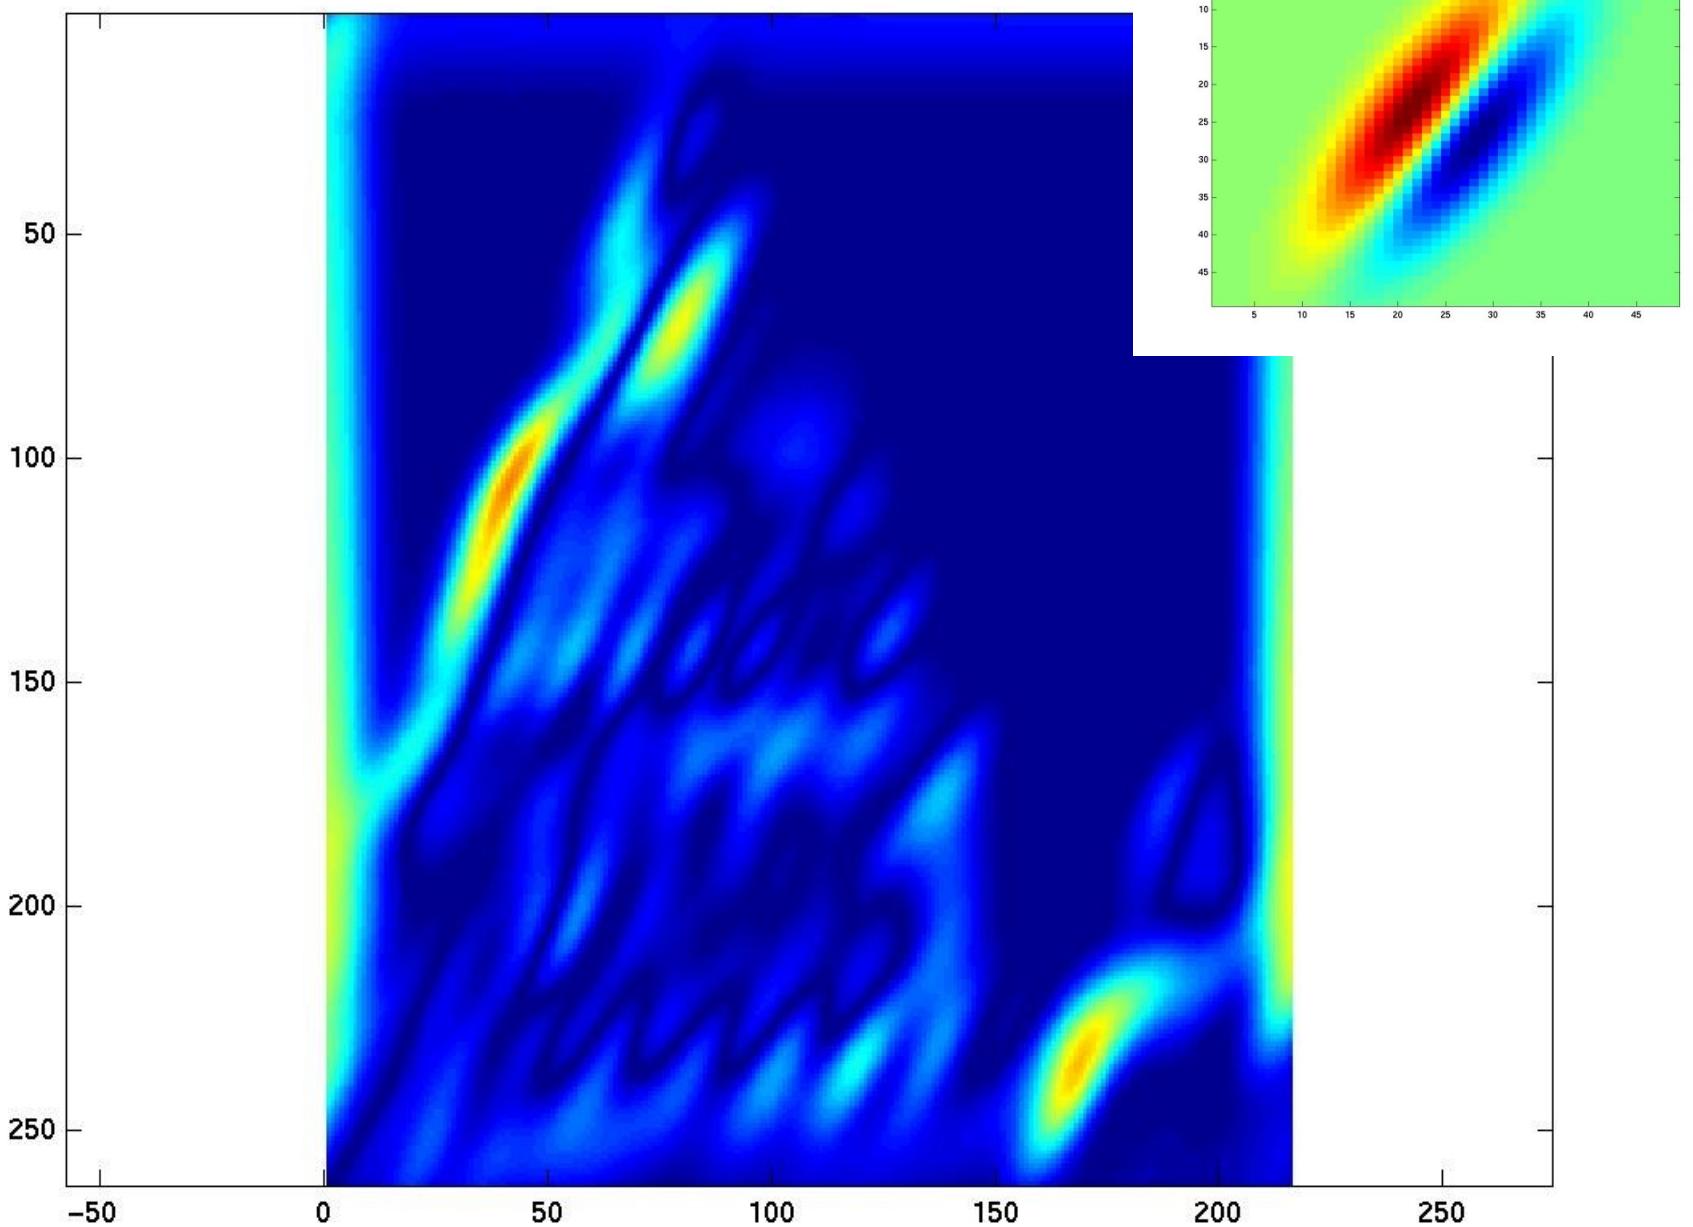


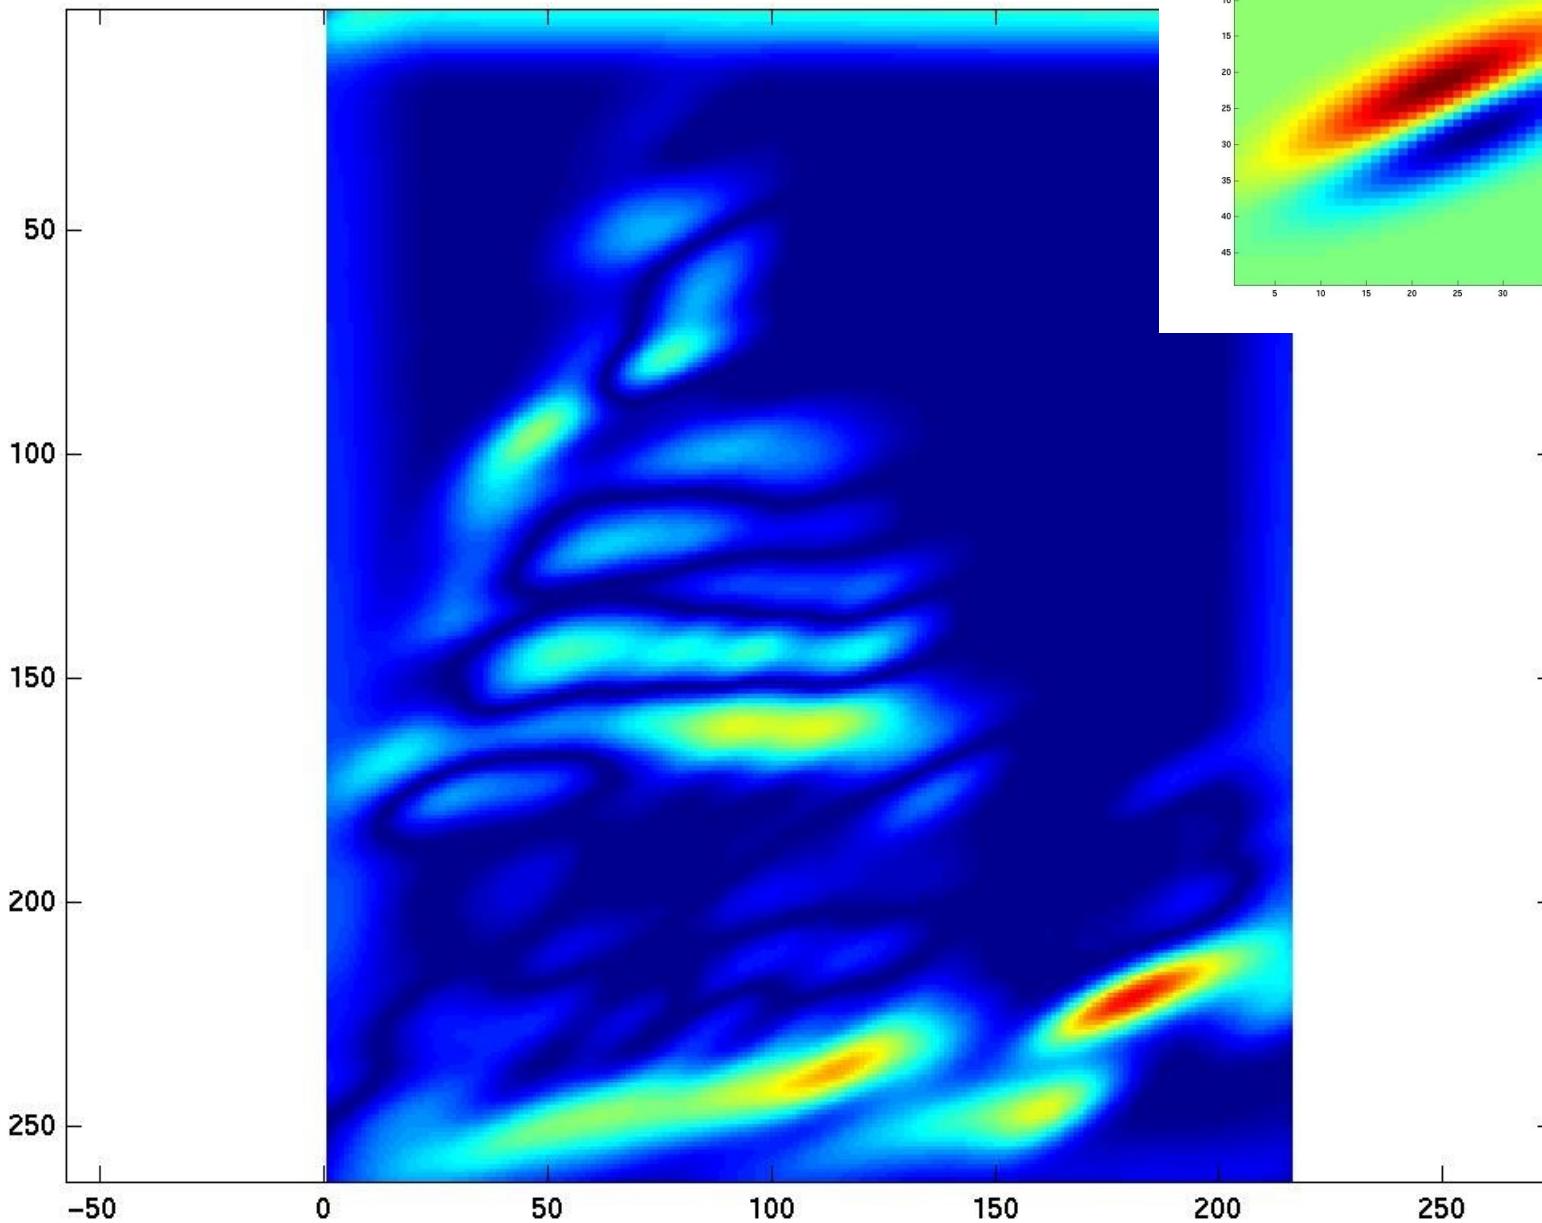


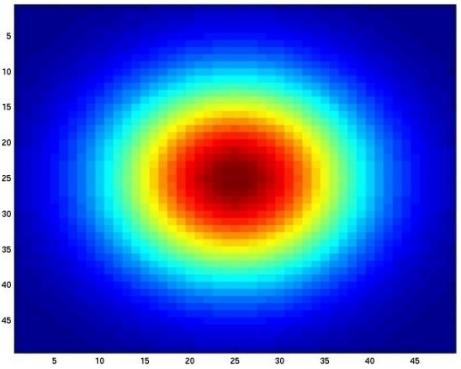
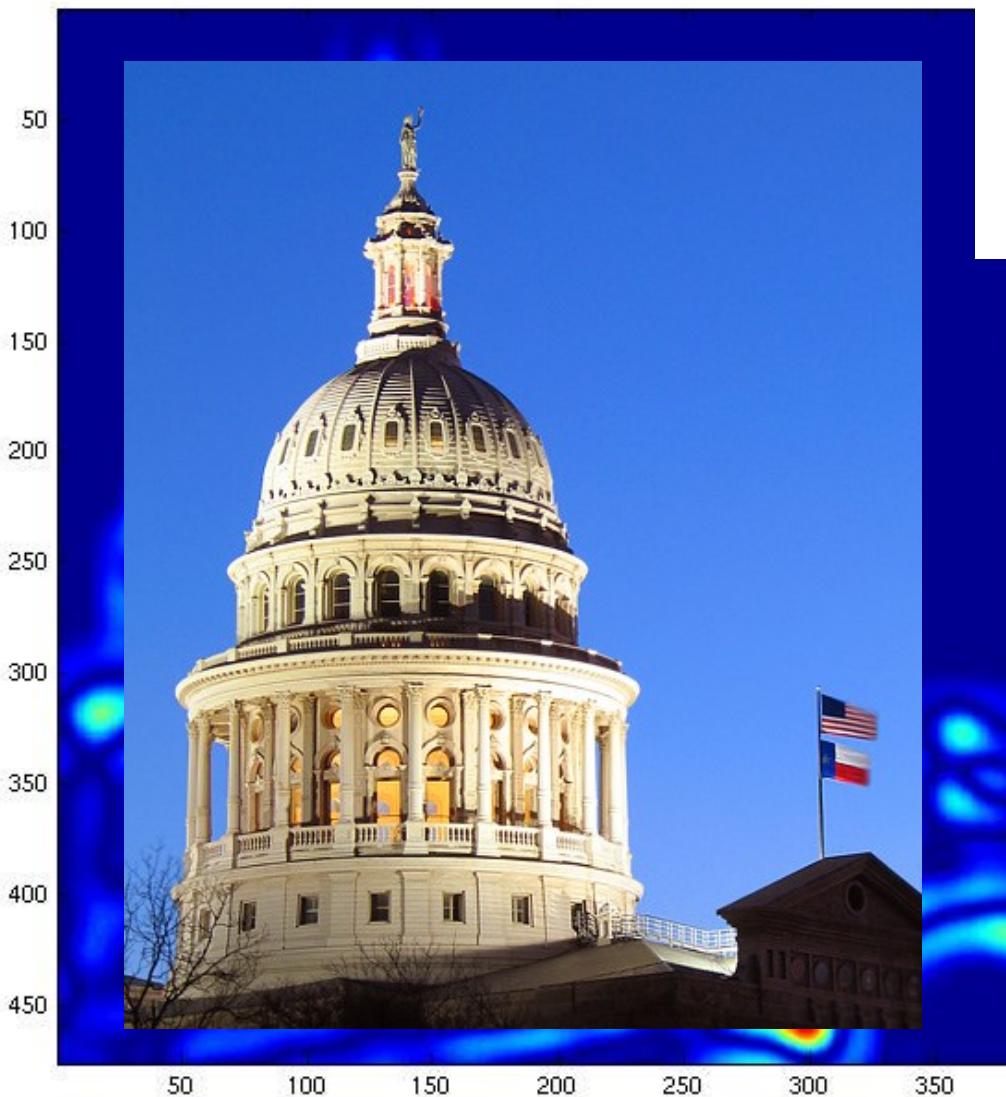




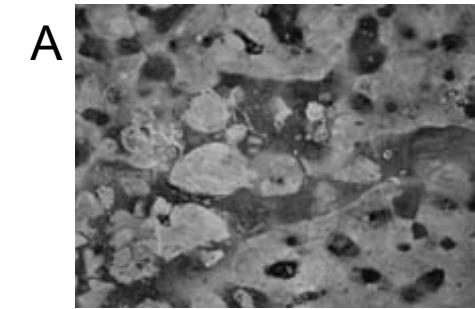
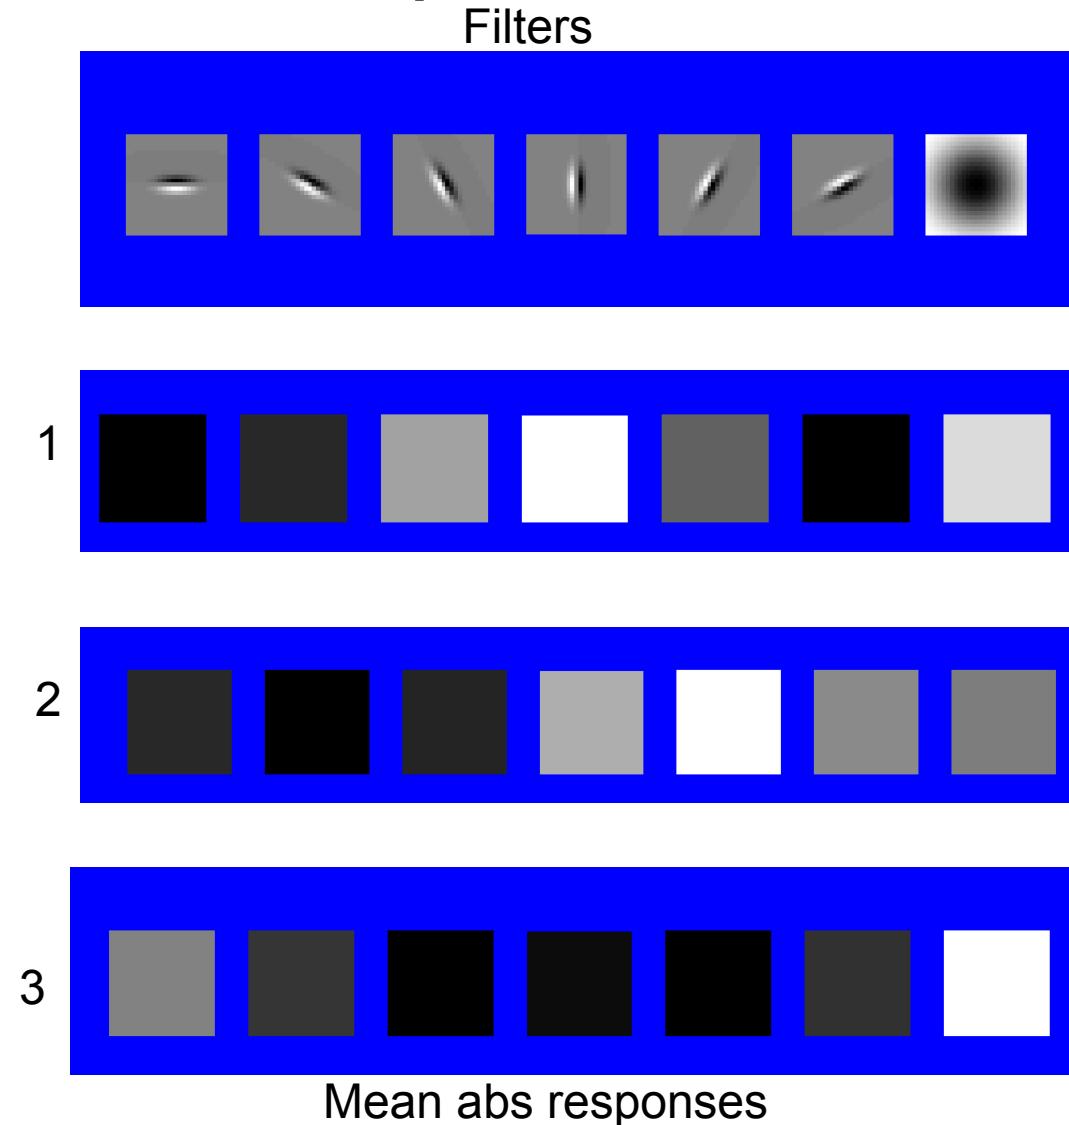




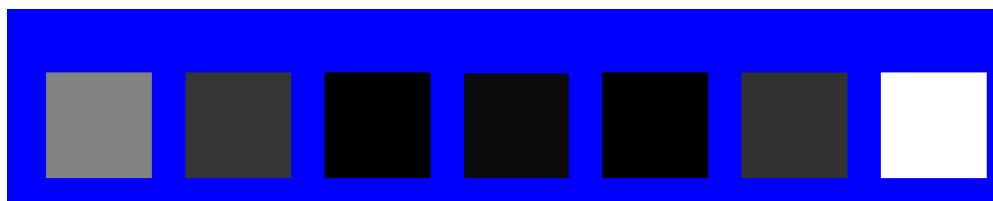
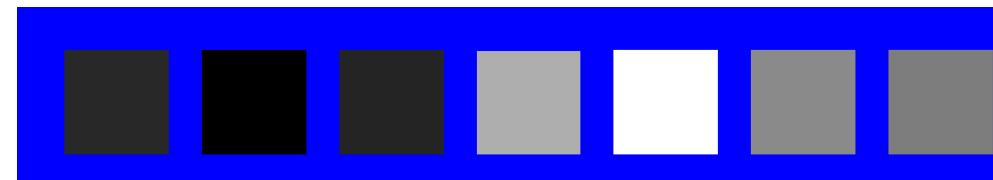
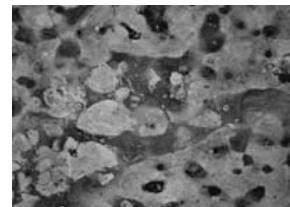
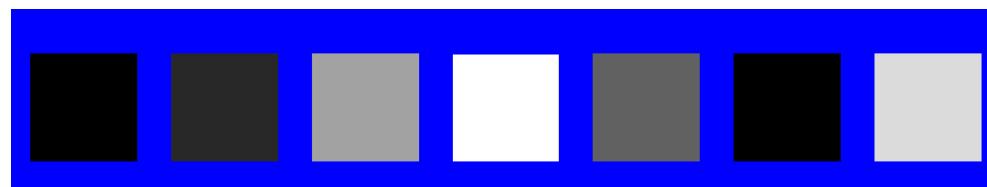
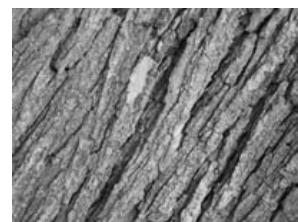
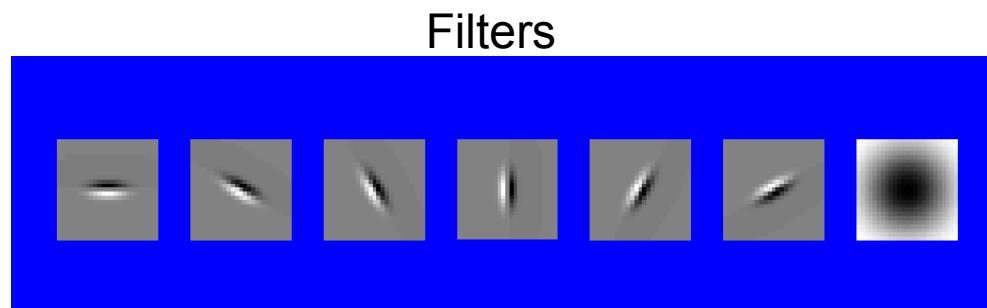




# You try: Can you match the texture to the response?

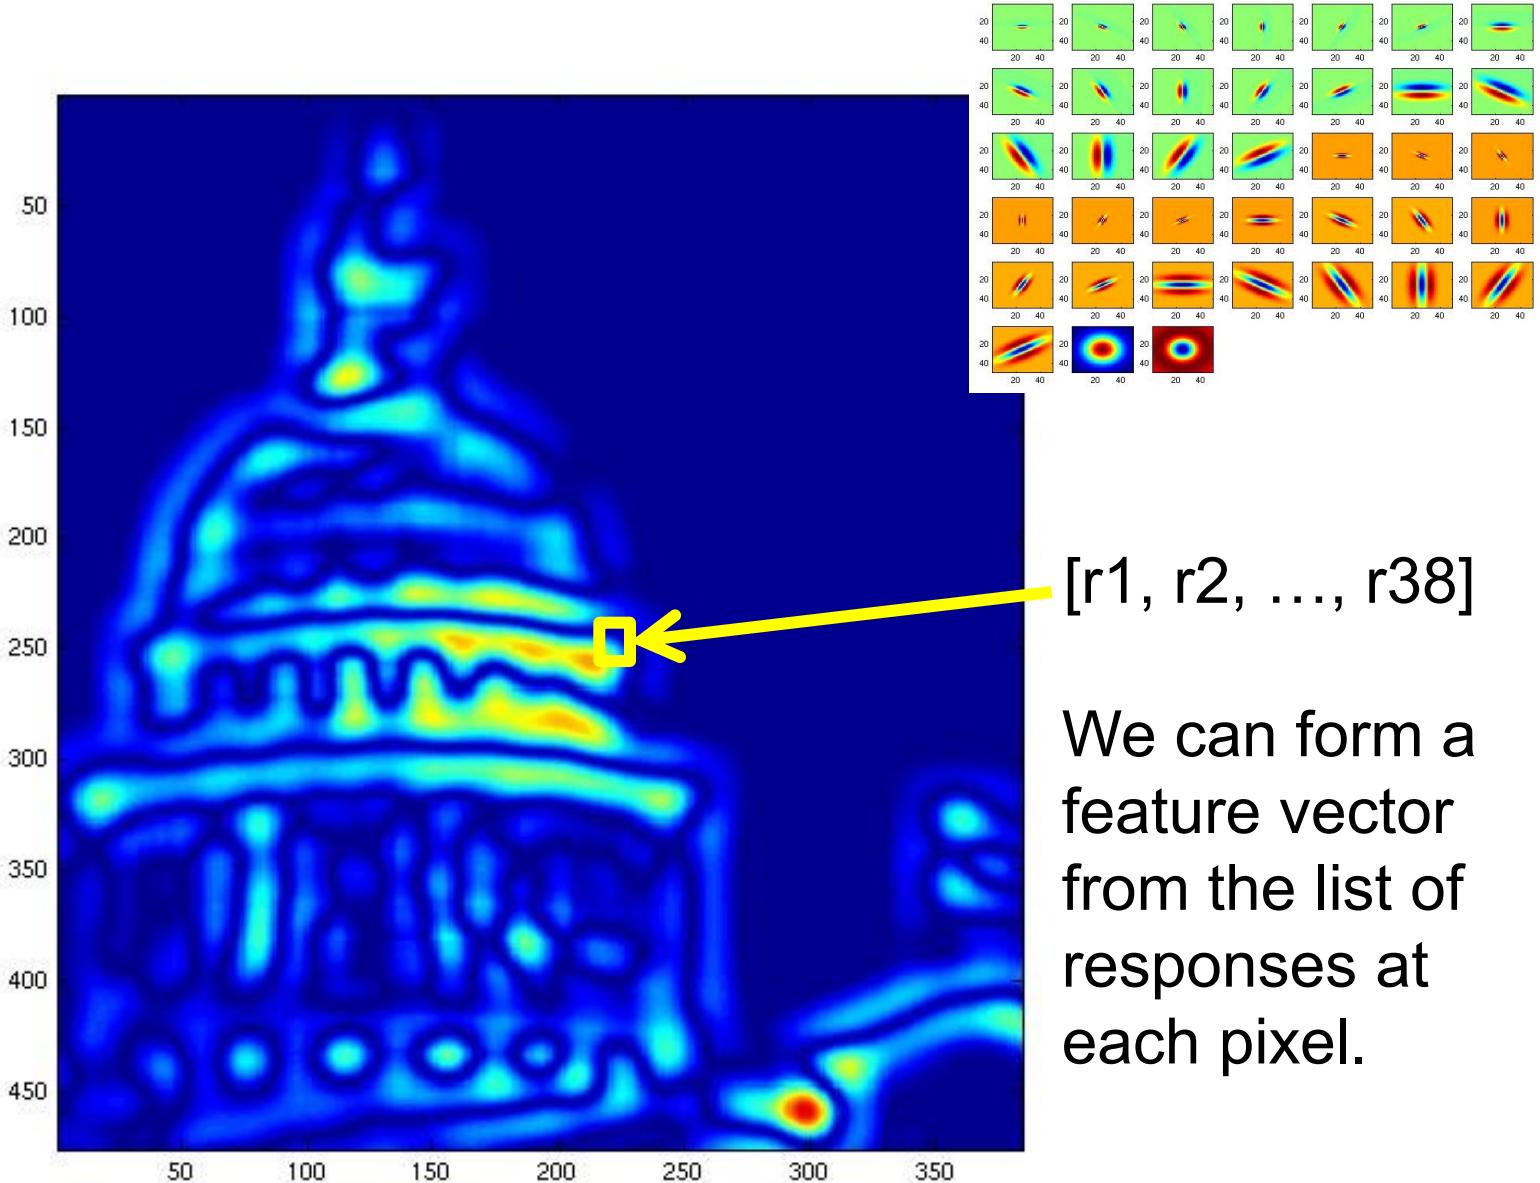


# Representing texture by mean abs response



Mean abs responses

Derek Hoiem



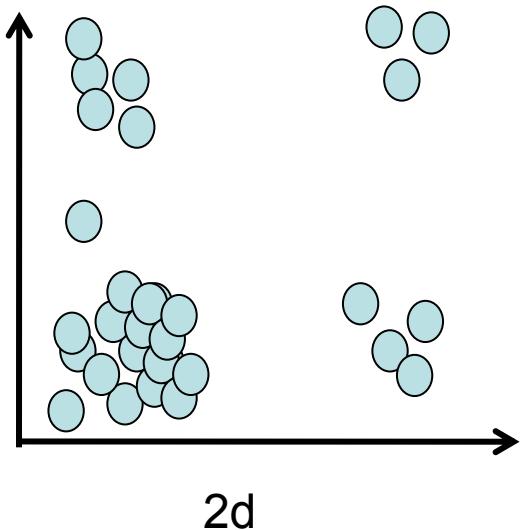
$[r_1, r_2, \dots, r_{38}]$

We can form a feature vector from the list of responses at each pixel.

# $d$ -dimensional features

$$D(a, b) = \sqrt{\sum_{i=1}^d (a_i - b_i)^2}$$

Euclidean distance ( $L_2$ )



Example uses of texture in  
vision: analysis

# Classifying materials, “stuff”

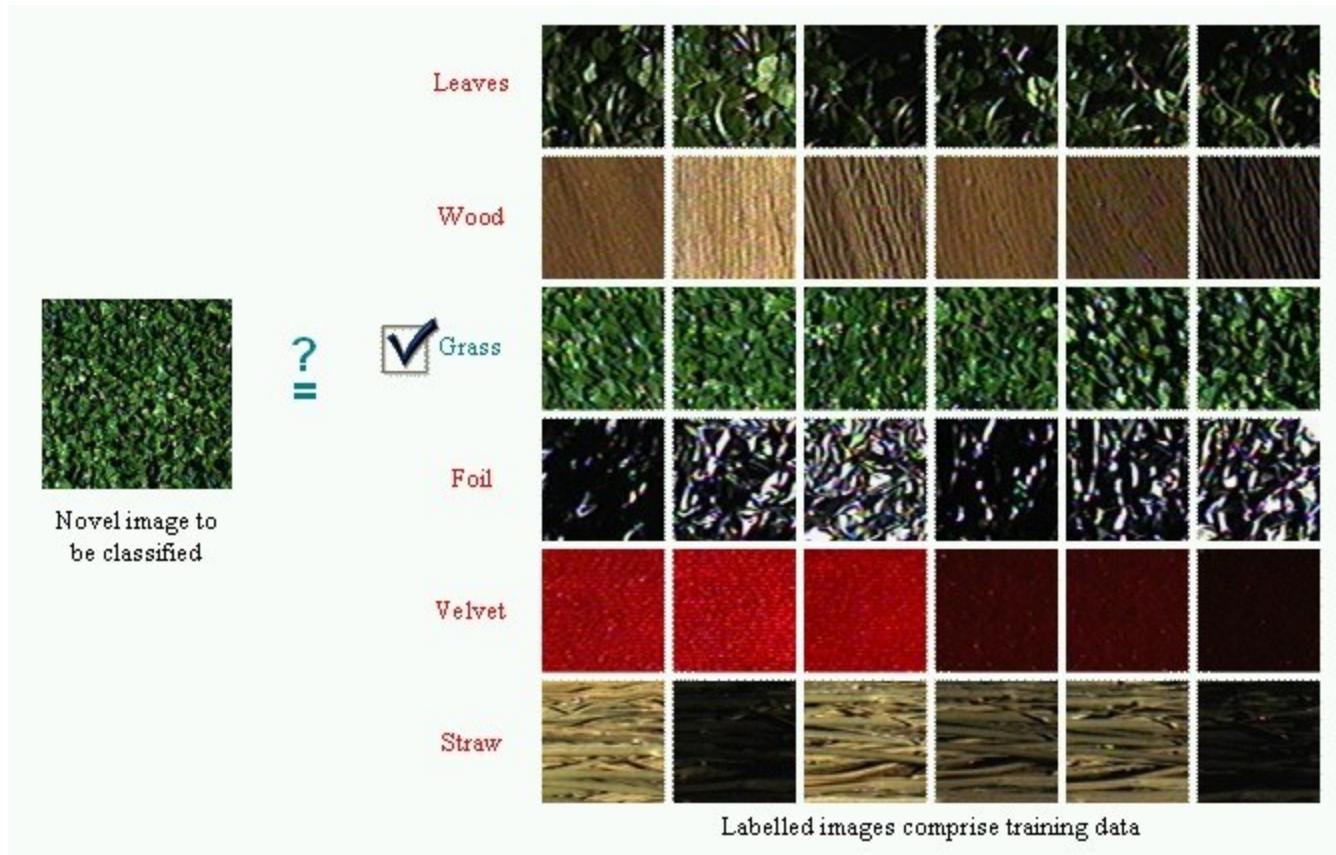
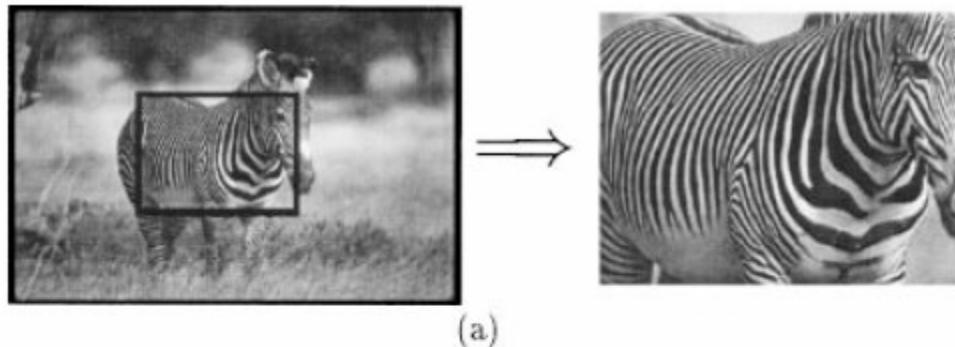
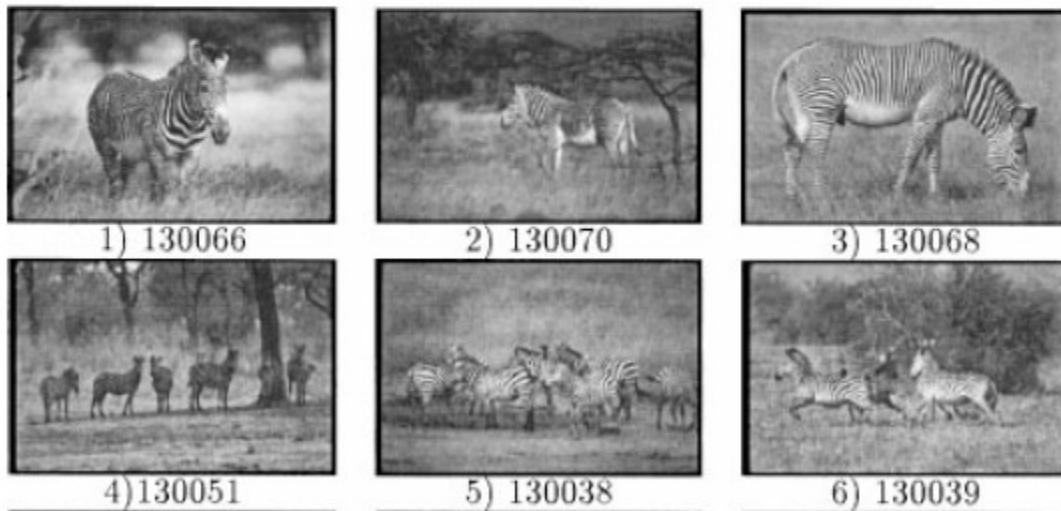


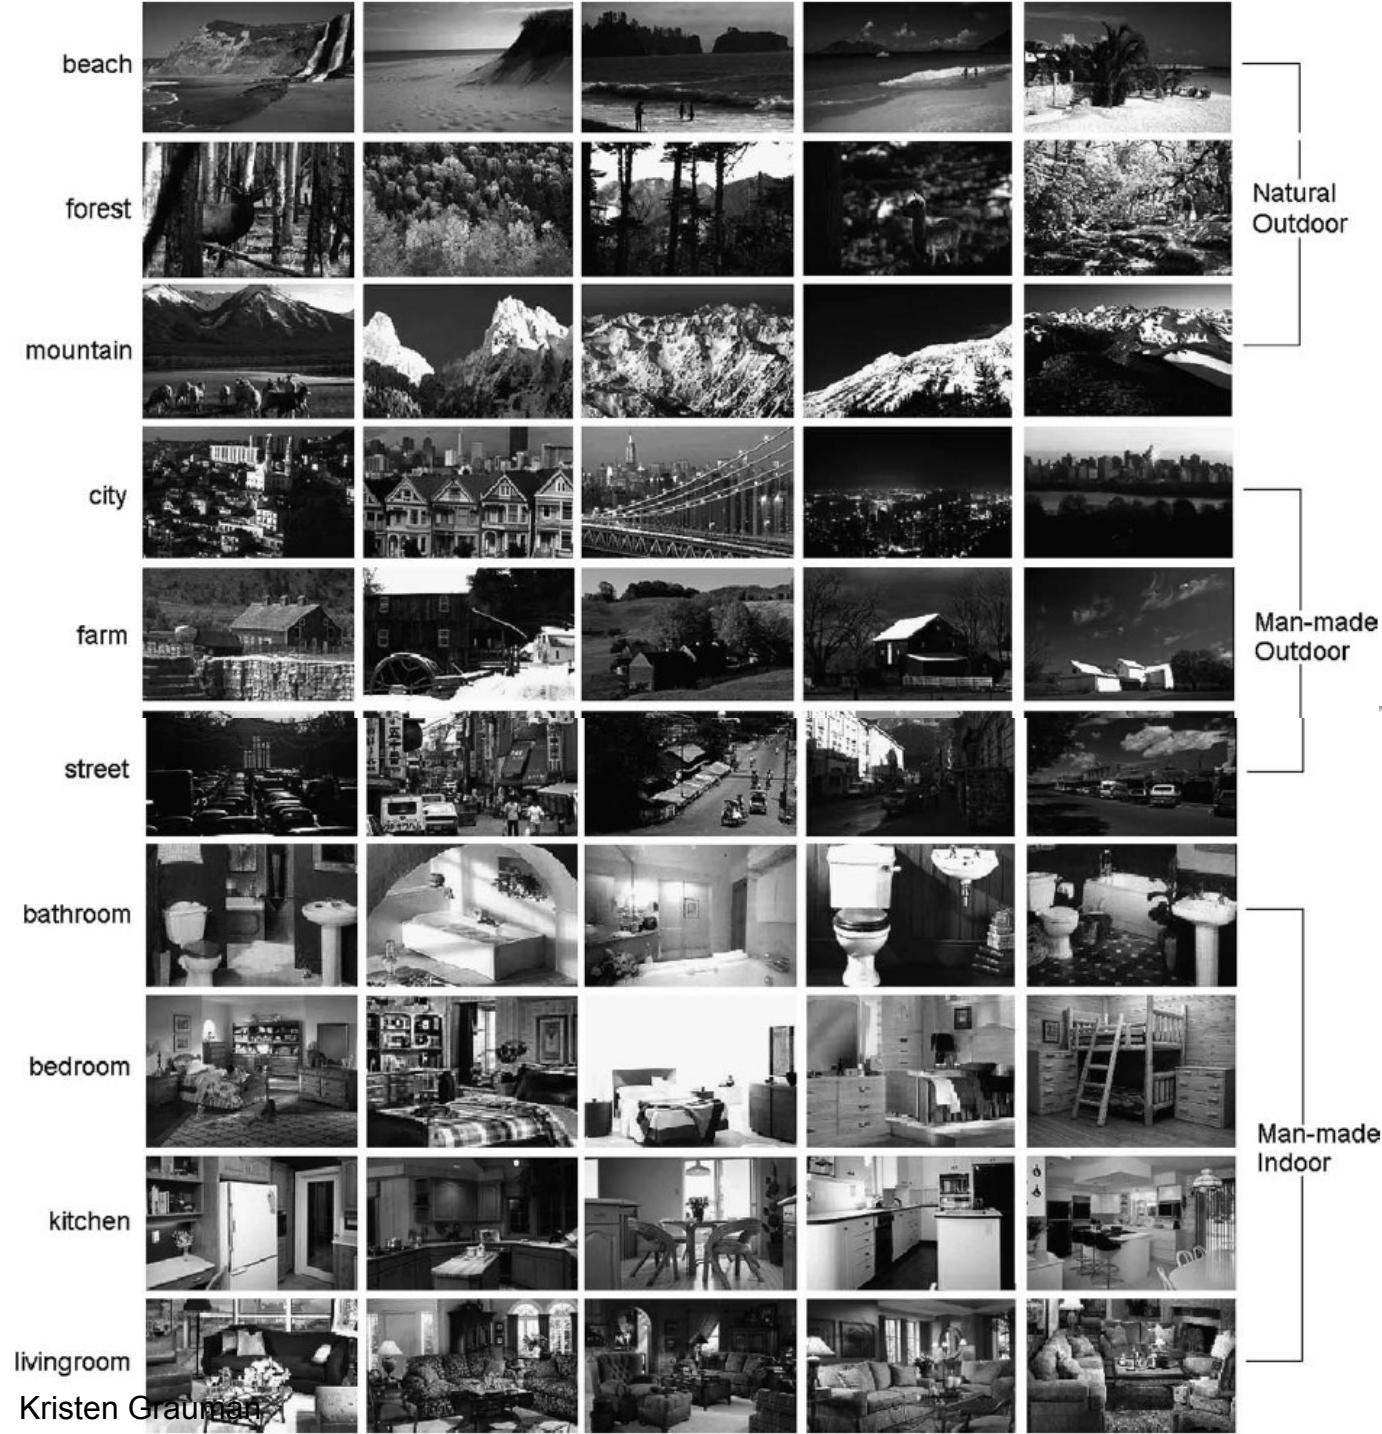
Figure by Varma  
& Zisserman



## Texture features for image retrieval



Y. Rubner, C. Tomasi, and L. J. Guibas. The earth mover's distance as a metric for image retrieval. *International Journal of Computer Vision*, 40(2):99-121, November 2000,



# Characterizing scene categories by texture

L. W. Renninger and  
J. Malik. When is  
scene identification  
just texture  
recognition? Vision  
Research 44 (2004)  
2301–2311



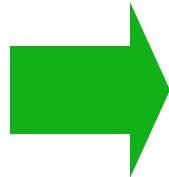
## Segmenting aerial imagery by textures

# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
  - Analyze, represent texture
  - Group image regions with consistent texture
- **Synthesis**
  - Generate new texture patches/images given some examples

# Texture synthesis

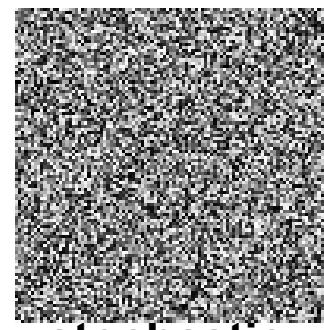
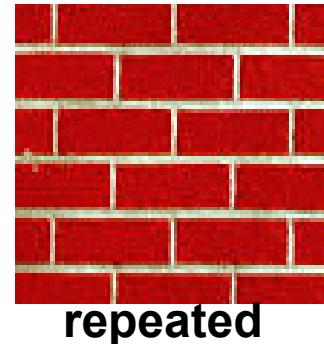
- Goal: create new samples of a given texture
- Many applications: virtual environments, hole-filling, texturing surfaces



# The Challenge

- Need to model the whole spectrum: from repeated to stochastic texture

Alexei A. Efros and Thomas K. Leung, “Texture Synthesis by Non-parametric Sampling,” Proc. International Conference on Computer Vision (ICCV), 1999.

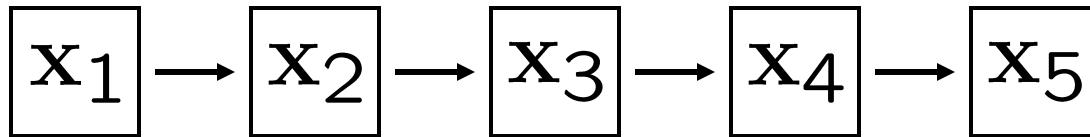


# Markov Chains

---

## Markov Chain

- a sequence of random variables  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$
- $\mathbf{x}_t$  is the **state** of the model at time t



# Markov Chain Example: Text

“A dog is a man’s best friend. It’s a dog eat dog world out there.”

$\mathbf{x}_{t-1}$

a		2/3		1/3							
dog			1/3				1/3	1/3			
is											
man’s	1				1						
best						1					
friend									1		
it’s	1										
eat		1									
world								1			
out									1		
there										1	
.											.

$\mathbf{x}_t$

$$p(\mathbf{x}_t | \mathbf{x}_{t-1})$$

# Text synthesis

---

Create plausible looking poetry, love letters, term papers, etc.

## Most basic algorithm

1. Build probability histogram
  - find all blocks of N consecutive words/letters in training documents
  - compute probability of occurrence  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_{t-(n-1)})$

**WE NEED TO EAT CAKE**

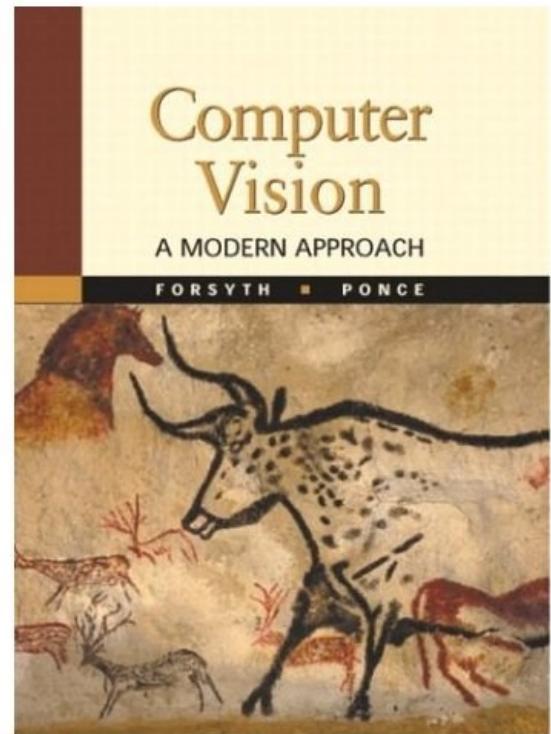
# Text synthesis

- Results:
  - “*As I've commented before, really relating to someone involves standing next to impossible.*”
  - “*One morning I shot an elephant in my arms and kissed him.*”
  - “*I spent an interesting evening recently with a grain of salt*”

Dewdney, “A potpourri of programmed prose and prosody” *Scientific American*, 1989.

# Synthesizing Computer Vision text

- What do we get if we extract the probabilities from a chapter on Linear Filters, and then synthesize new statements?



Check out Yisong Yue's website implementing text generation: build your own text Markov Chain for a given text corpus. <http://www.yisongyue.com/shaney/index.php>

# Synthesized text

- This means we cannot obtain a separate copy of the best studied regions in the sum.
- All this activity will result in the primate visual system.
- The response is also Gaussian, and hence isn't bandlimited.
- Instead, we need to know only its response to any data vector, we need to apply a low pass filter that strongly reduces the content of the Fourier transform of a very large standard deviation.
- It is clear how this integral exist (it is sufficient for all pixels within a  $2k + 1 \times 2k + 1 \times 2k + 1 \times 2k + 1$  — required for the images separately).

# Markov Random Field

---

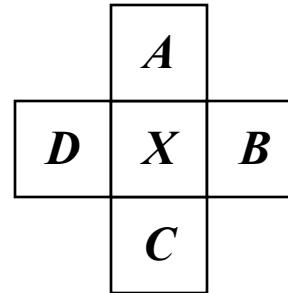
## A Markov random field (MRF)

- generalization of Markov chains to two or more dimensions.

### First-order MRF:

- probability that pixel  $X$  takes a certain value given the values of neighbors  $A$ ,  $B$ ,  $C$ , and  $D$ :

$$P(X|A, B, C, D)$$



# Texture Synthesis [\[Efros & Leung, ICCV 99\]](#)

---

Can apply 2D version of text synthesis

Texture corpus  
(sample)



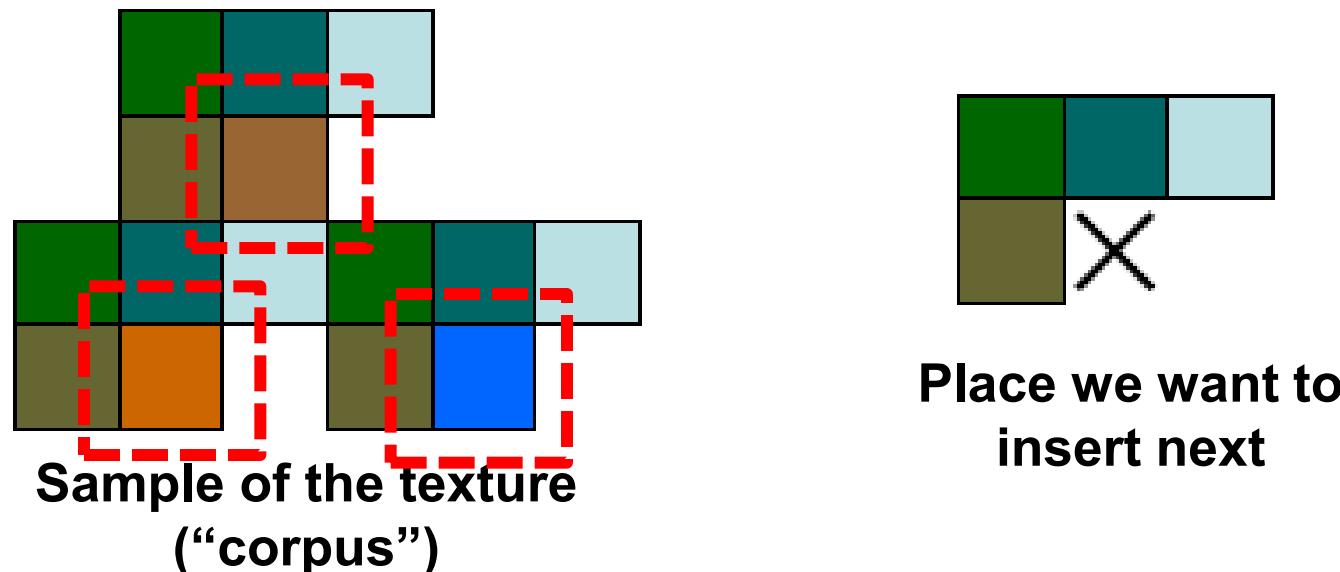
Output



# Texture synthesis: intuition

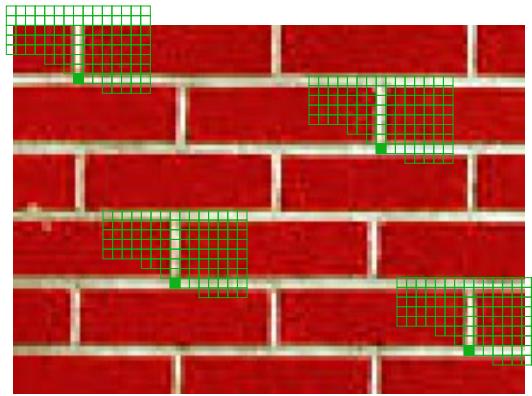
Before, we inserted the next word based on existing nearby words...

Now we want to insert **pixel intensities** based on existing nearby pixel values.

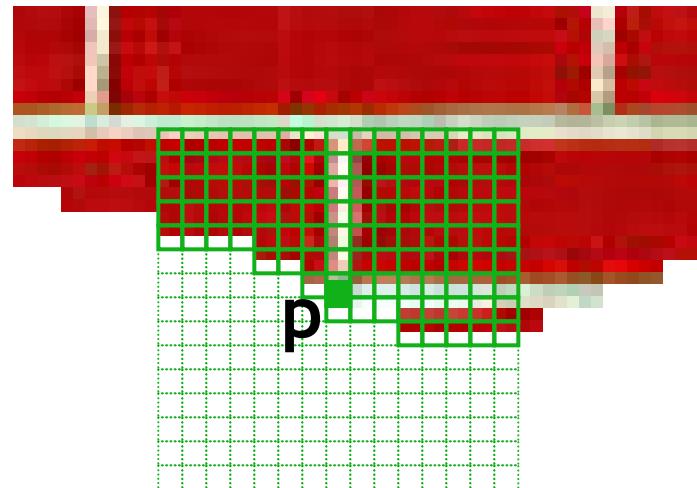


Distribution of a value of a pixel is conditioned on its neighbors alone.

# Synthesizing One Pixel



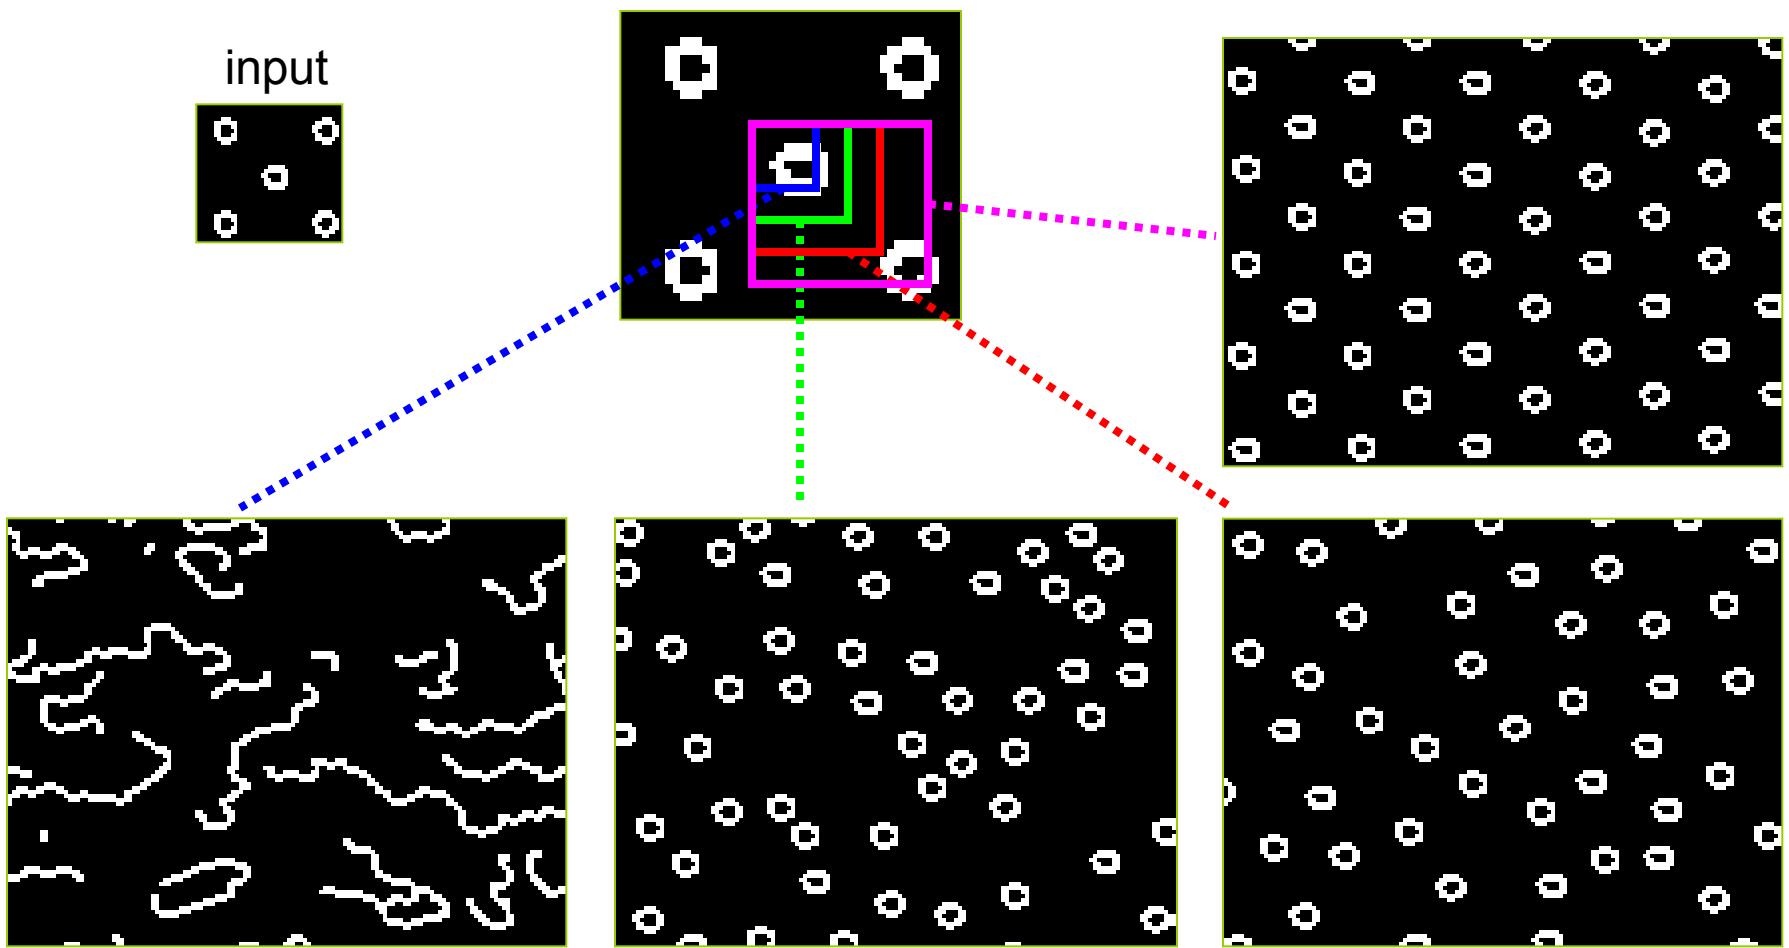
input image



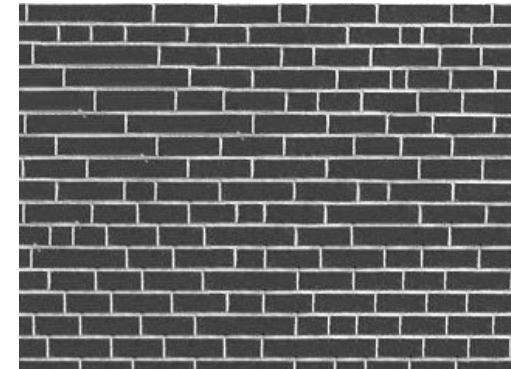
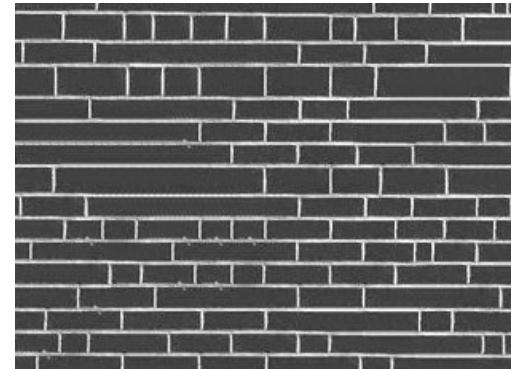
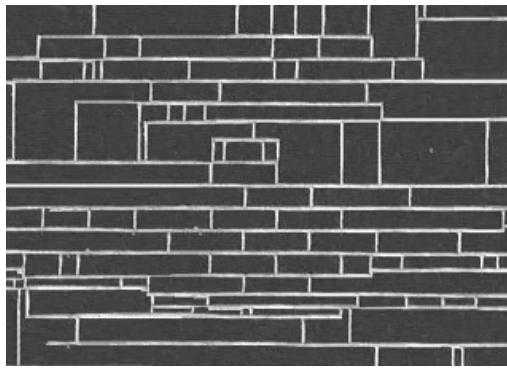
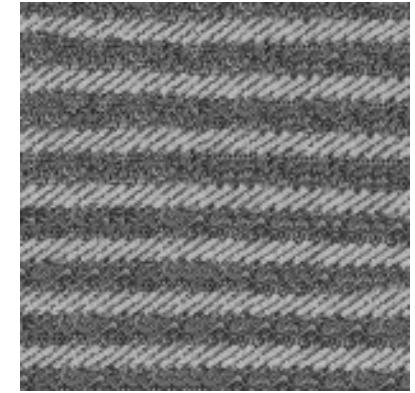
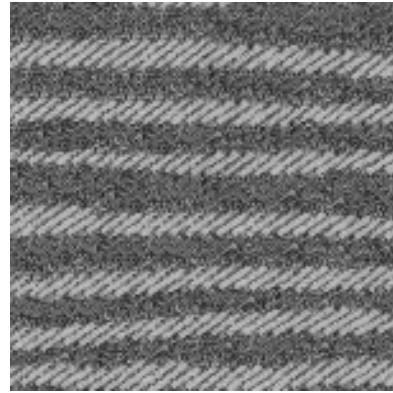
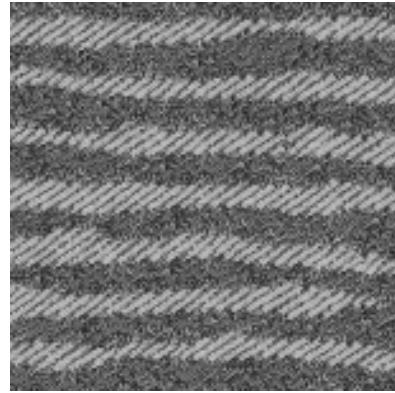
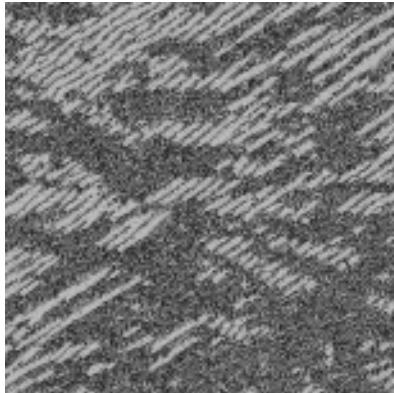
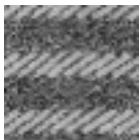
synthesized image

- What is  $P(x|\text{neighborhood of pixels around } x)$  ?
- Find all the windows in the image that match the neighborhood
- To synthesize  $x$ 
  - pick one matching window at random
  - assign  $x$  to be the center pixel of that window
- An **exact** neighbourhood match might not be present, so find the **best** matches using **SSD error** and randomly choose between them, preferring better matches with higher probability

# Neighborhood Window



# Varying Window Size

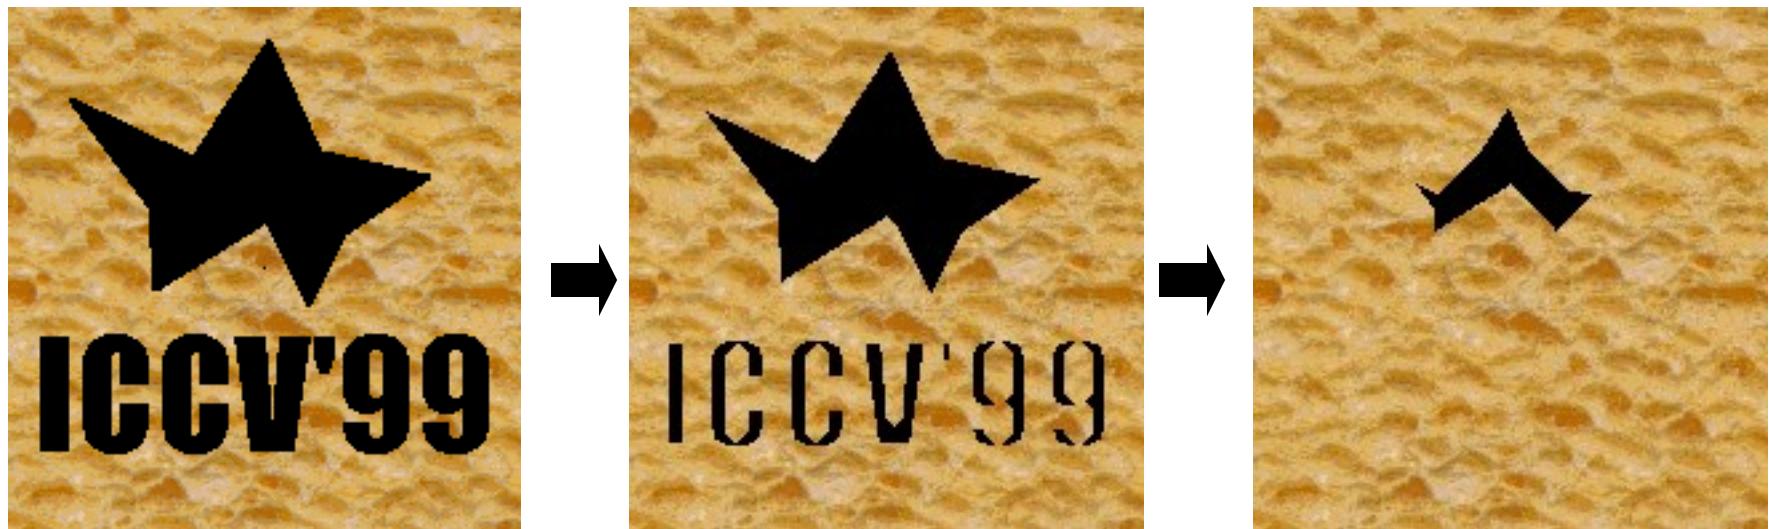


Increasing window size



# Growing Texture

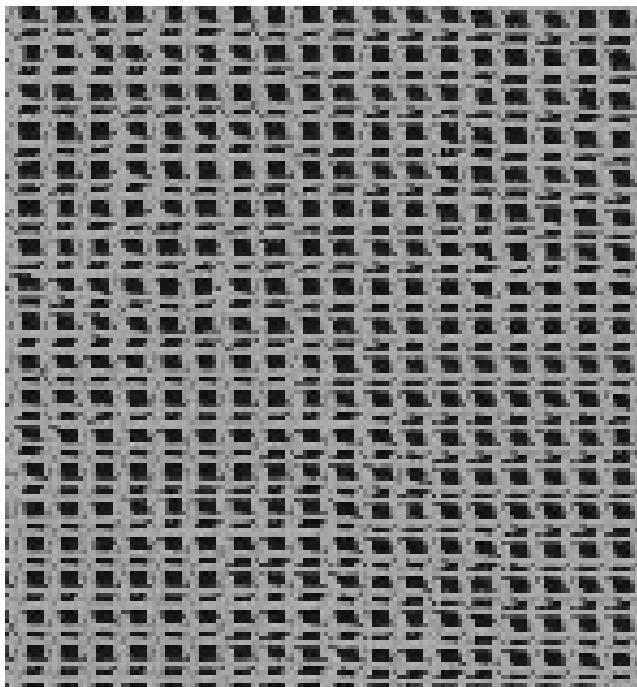
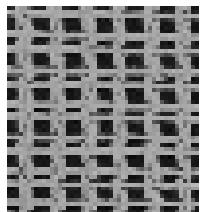
---



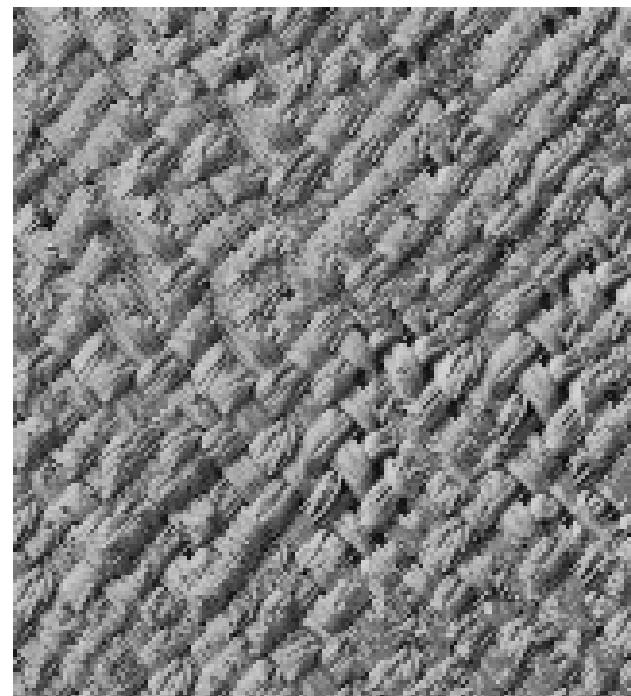
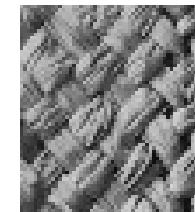
- Starting from the initial image, “grow” the texture one pixel at a time

# Synthesis results

french canvas



rafia weave

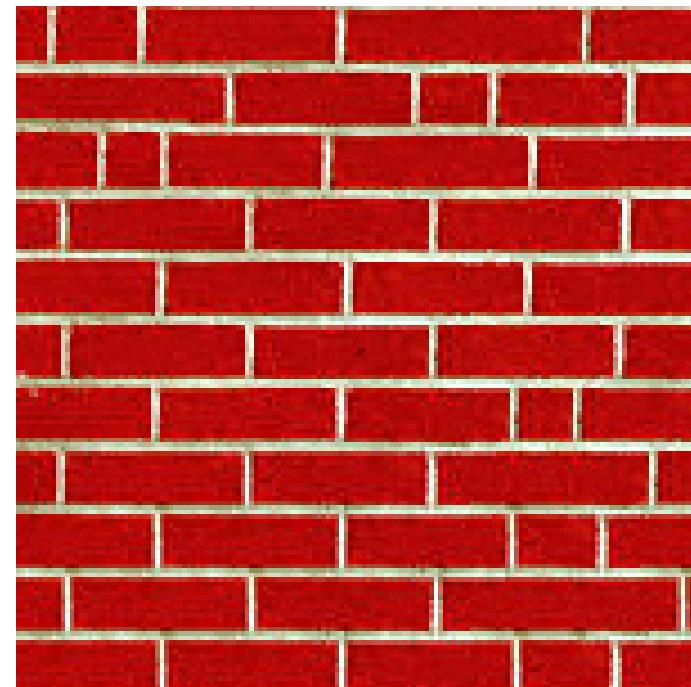
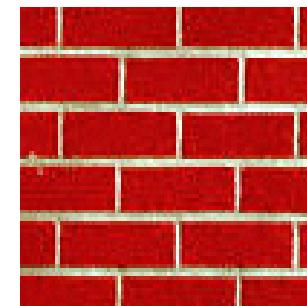


# Synthesis results

white bread



brick wall

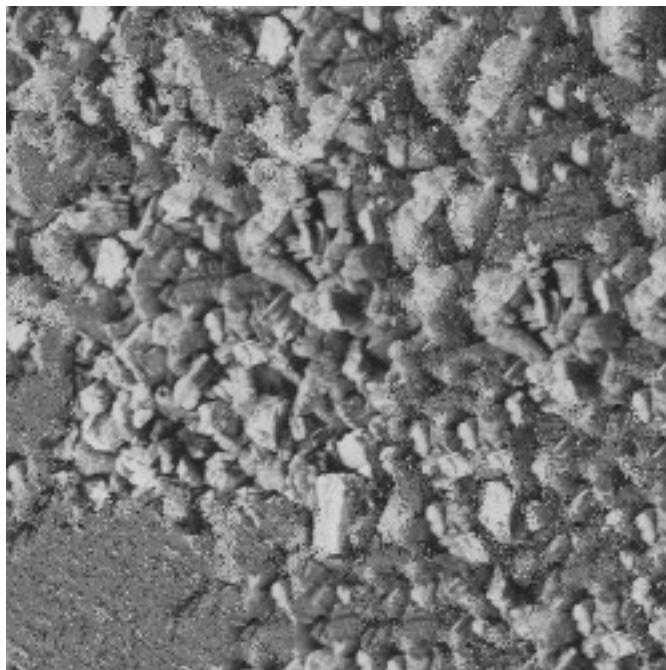


# Synthesis results

uring in the unsensato  
r Dick Gephardt was fai  
rful riff on the looming  
nly asked, "What's your  
tions?" A heartfelt sigh  
story about the emergenc  
es against Clinton. "Boy  
g people about continuin  
ardt began, patiently obs  
s, that the legal system h  
e with this latest tanger

ithairm, them . "Whnephartfe lartifelintomimen  
fel ck Clirtioout omaim thartfelins. f a ut s anetc  
the ry onst wartfe lck Gephntoomimeationl sigak  
Chiooufit Clinut Cll riff on, hat's yordn, parut tly  
ons ycontonsteht wasked, paim t sahe loo' riff on l  
nskoneploourtfeas leil A nst Clit, "Wleontongal s  
k Cirtioouirtfepe óng pme abegal fartfenstemem  
tiensteneltorydt telemeaphinsverdt was agemer  
ff ons artientont Cling peme asurtfe atih, "Boui s  
hal s fartfelt sig pedr tl'dt ske abounutie aboutioo  
itfaonewwas yous aboonthardt thatins fain, ped,  
ains, them, pabout wasy arfut coutly d, ln A h  
ole emthringbooreme agas fa bontinsyst Clinut  
ory about continst Clipeouinst Cloke agatiff out C  
stome zinemen tly ardt beoraboul n, thenly as t C  
cons faimeme Diontont wat coutlyohgans as fan  
ien, phrtfaul, "Wbaut cout congagal cómininga  
mifmst Cliiy abon al coountha ermungairt tf oun  
The looorystan loontieph. Intly on, theoplegatick  
ul tatteeontly atie Diontiomt wal s f thegæ ener  
mthahsat's enenhhmas fan. "intchthorw ahons v

# Failure Cases

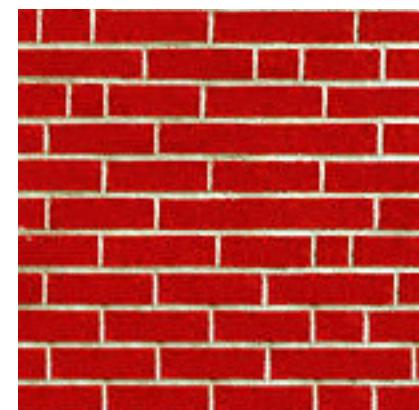
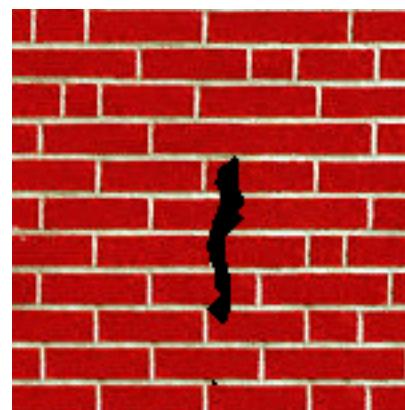


**Growing garbage**

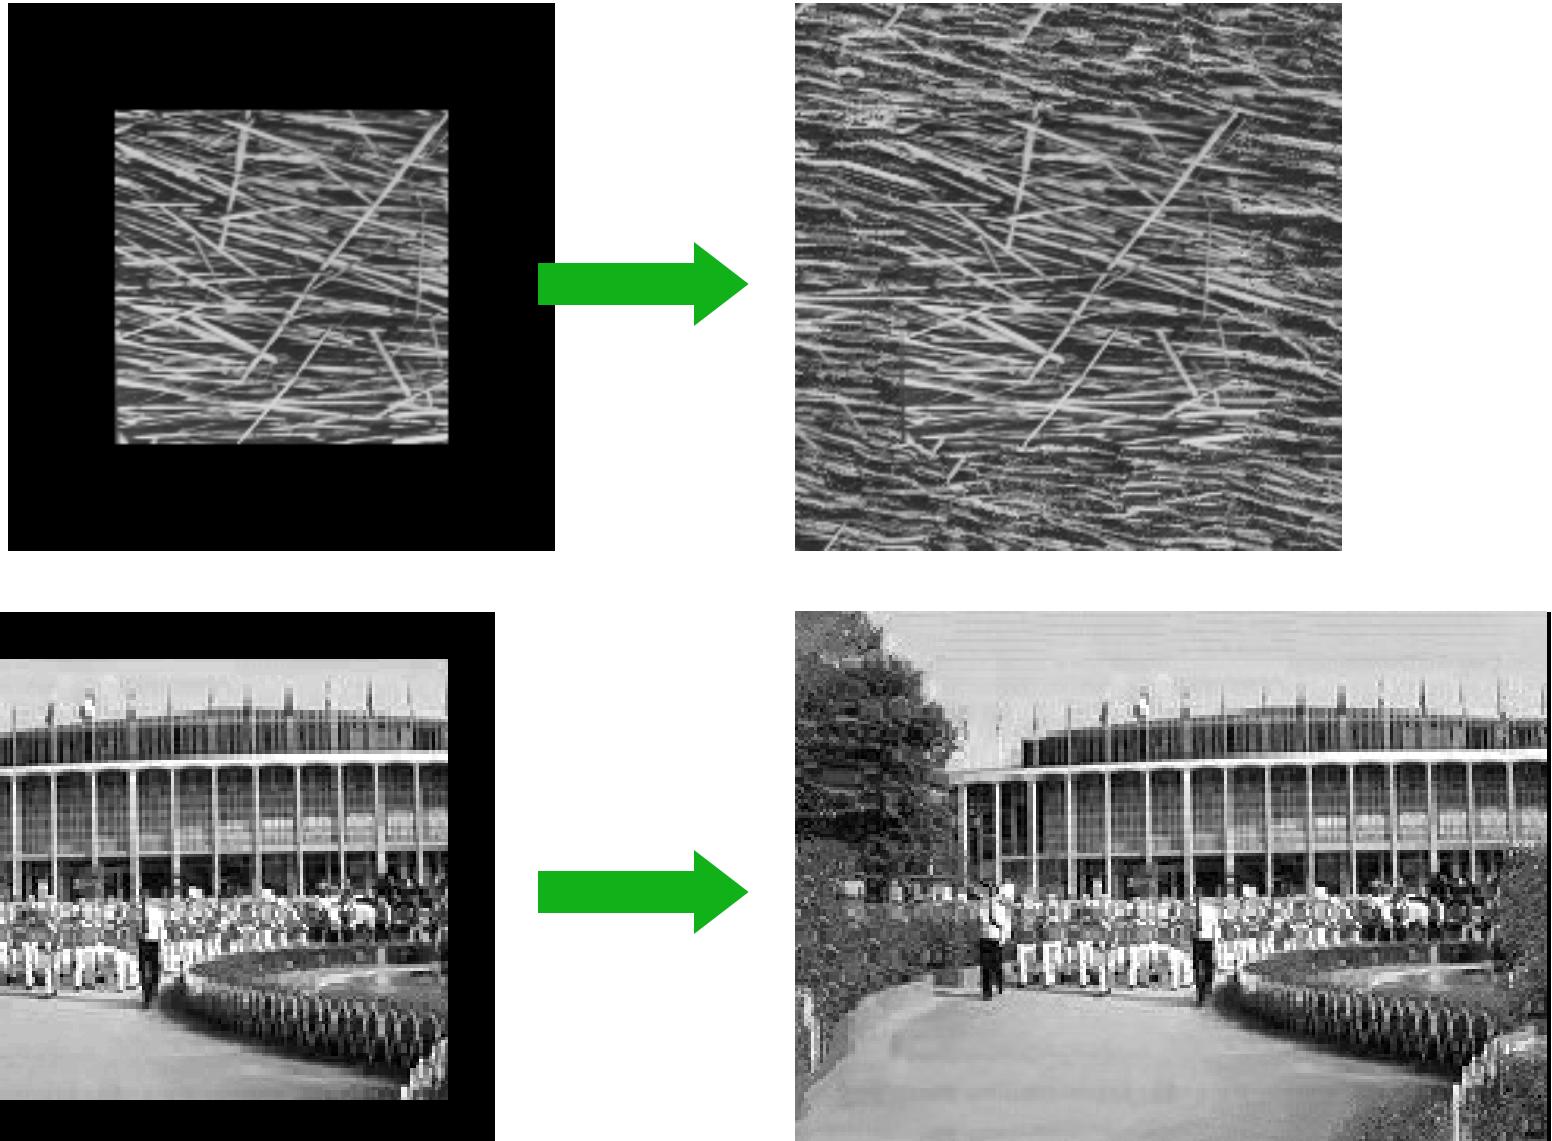


**Verbatim copying**

# Hole Filling

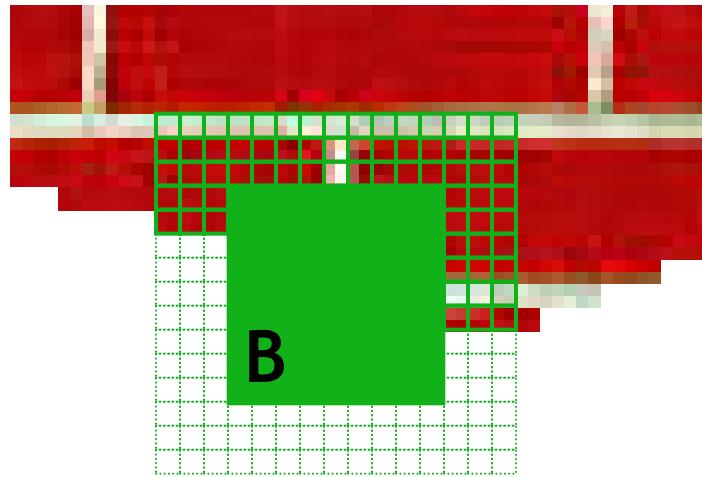


# Extrapolation



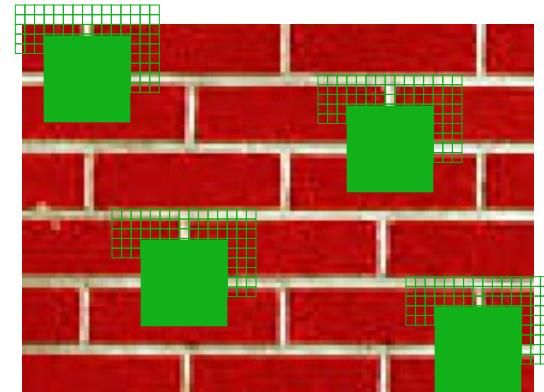
- The Efros & Leung algorithm
  - Simple
  - Surprisingly good results
  - Synthesis is easier than analysis!
  - ...but very slow

# Image Quilting [Efros & Freeman 2001]



Synthesizing a block

non-parametric  
sampling

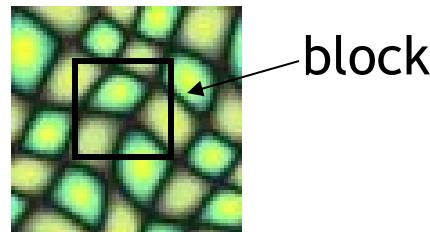


Input image

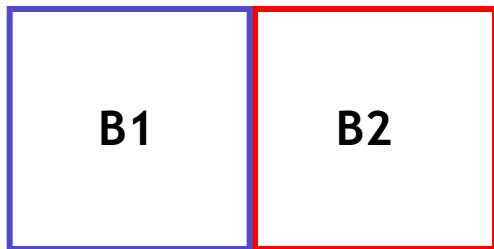
- Observation: neighbor pixels are highly correlated

Idea: unit of synthesis = block

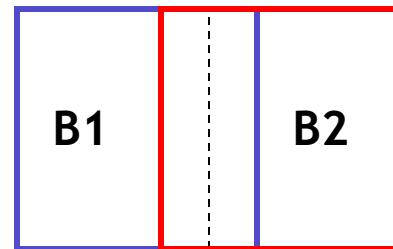
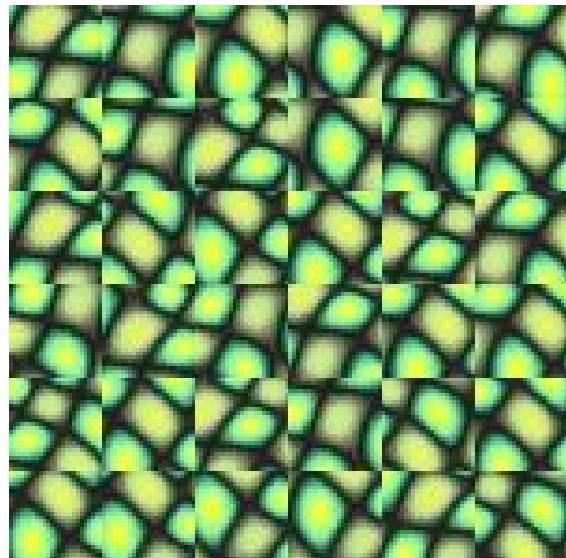
- Exactly the same but now we want  $P(B|N(B))$
- Much faster: synthesize all pixels in a block at once



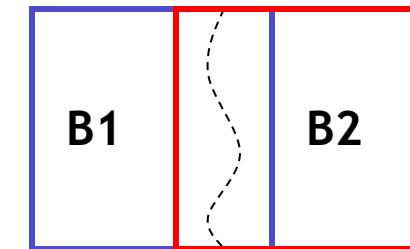
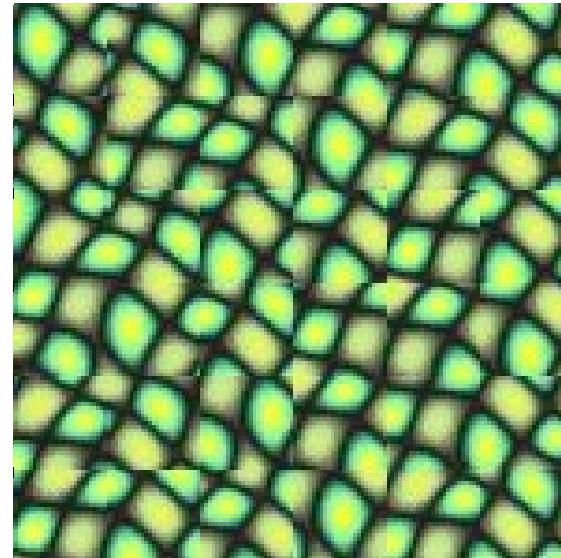
Input texture



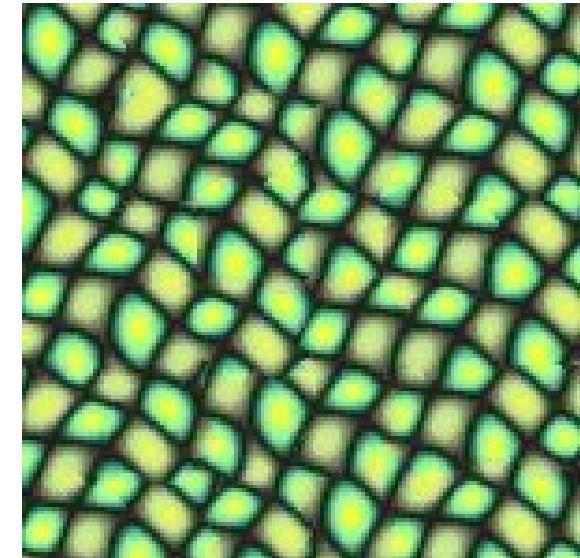
Random placement  
of blocks



Neighboring blocks  
constrained by overlap

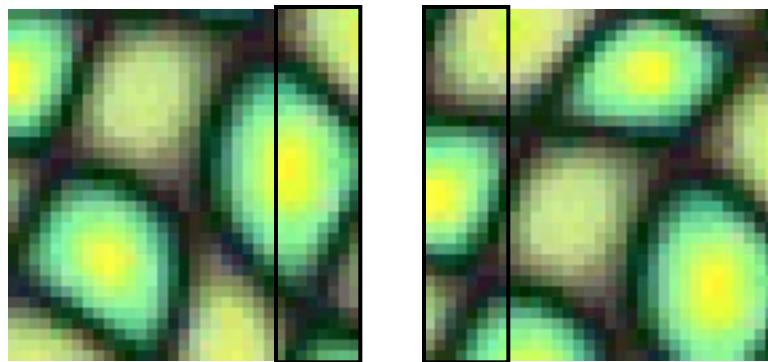


Minimal error  
boundary cut

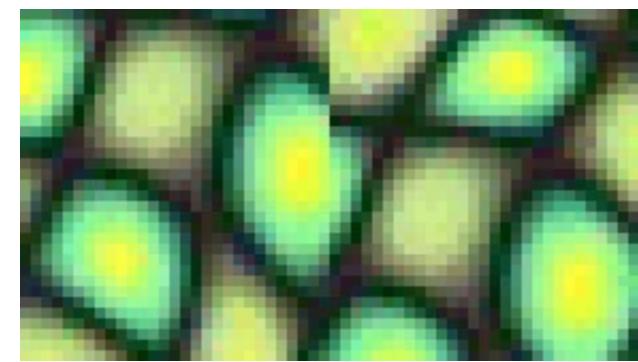


# Minimal error boundary

overlapping blocks

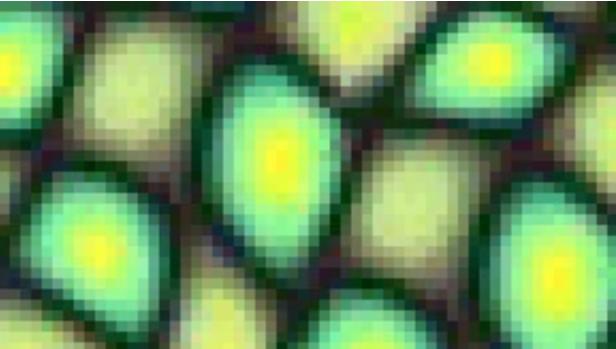


vertical boundary

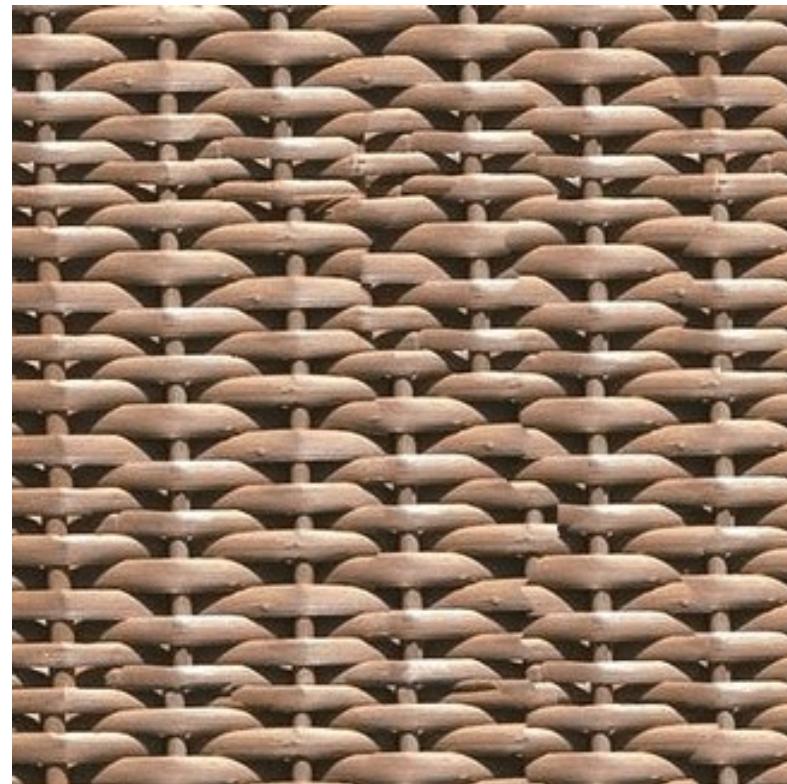
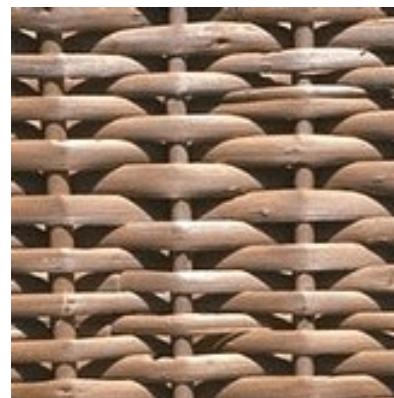


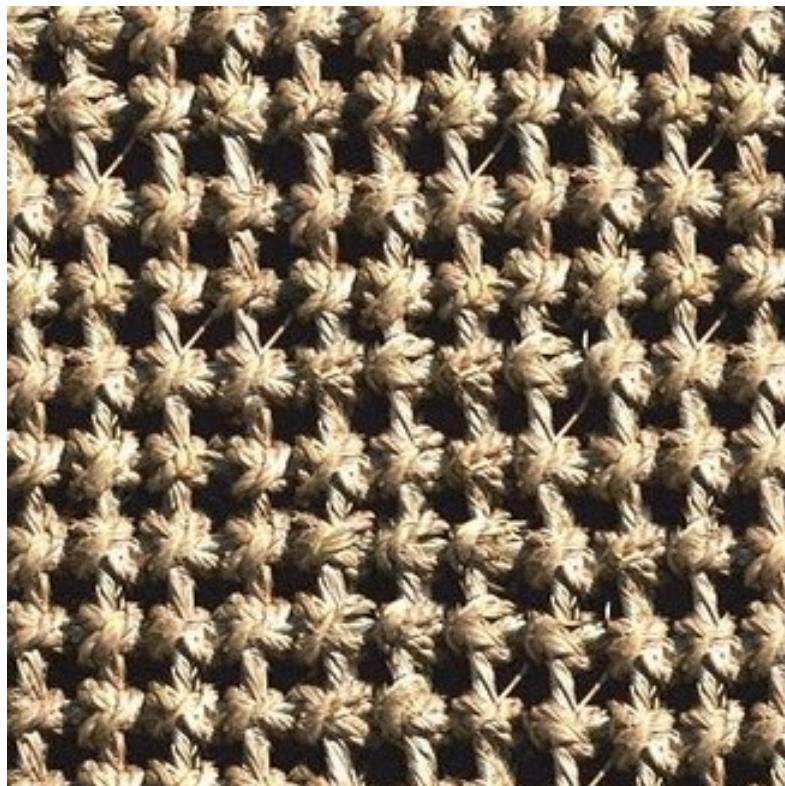
$$\left[ \begin{array}{c} \text{block 1} \\ - \\ \text{block 2} \end{array} \right]^2 = \text{overlap error}$$

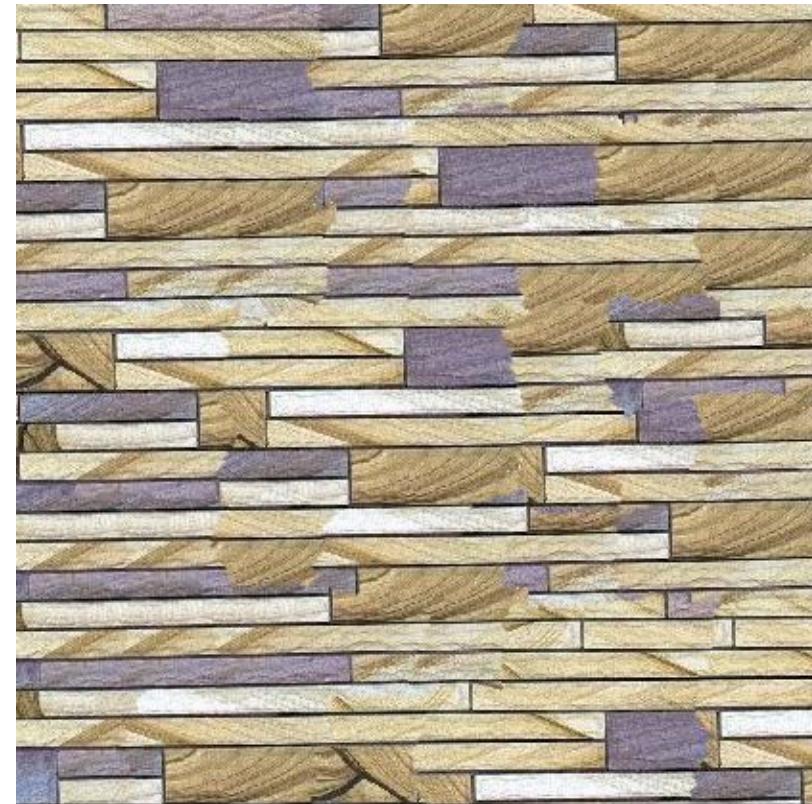
A diagram illustrating the calculation of the overlap error. It shows two overlapping blocks of a blurred image. The difference between them is squared to produce a red zigzag boundary, which is labeled as the "overlap error".



min. error boundary

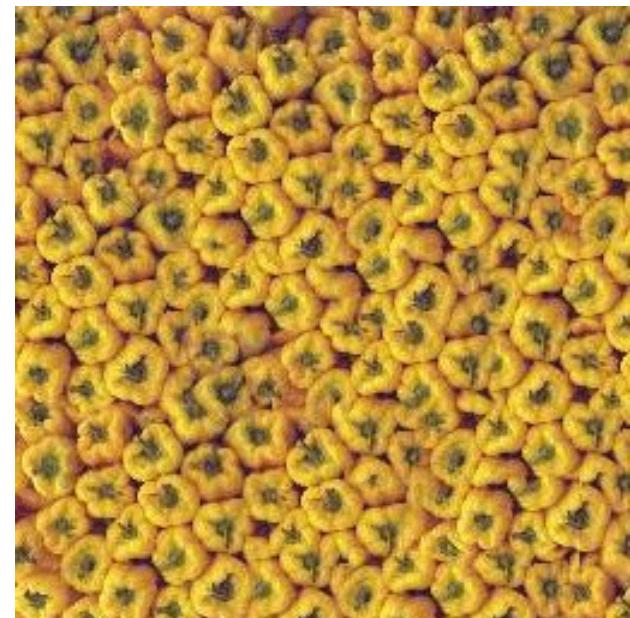
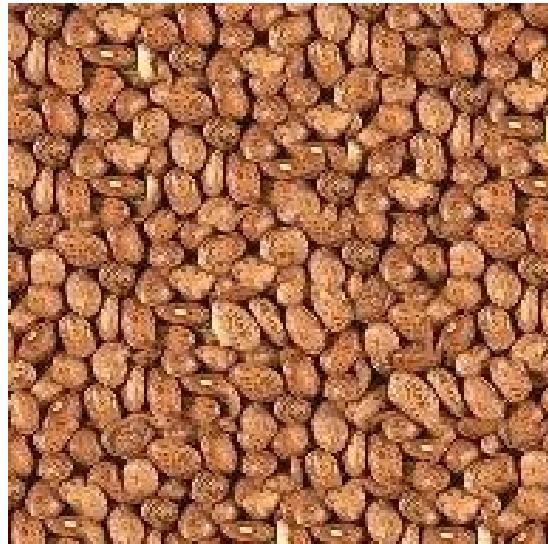
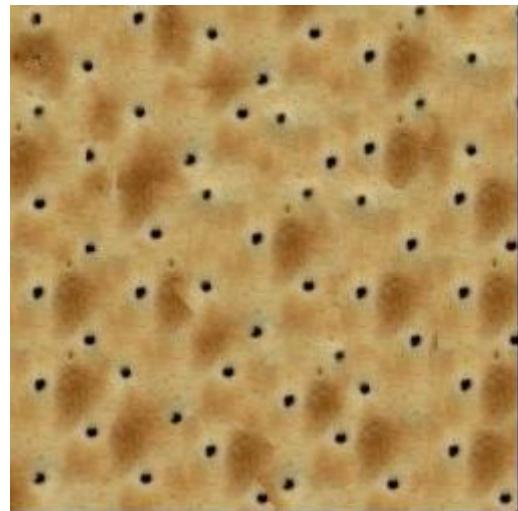
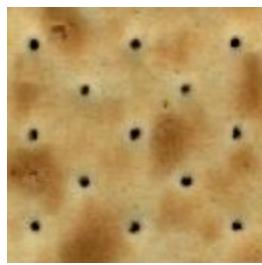














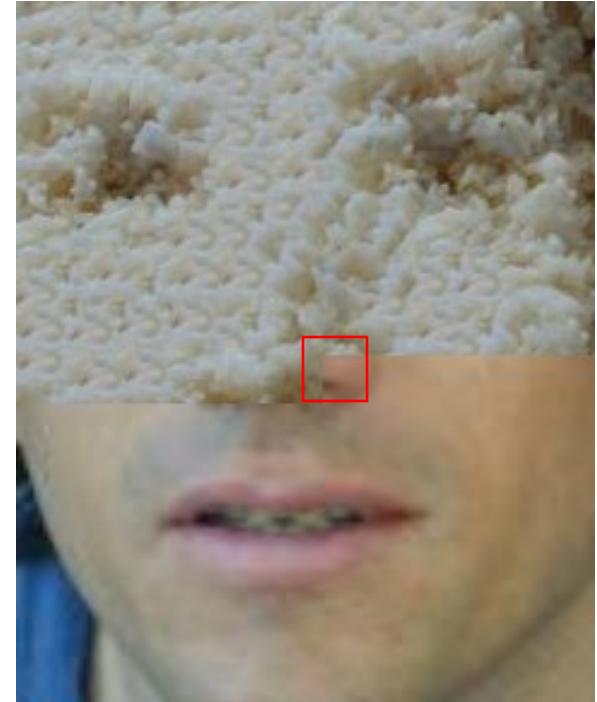
# Failures

(Chernobyl  
Harvest)



# Texture Transfer

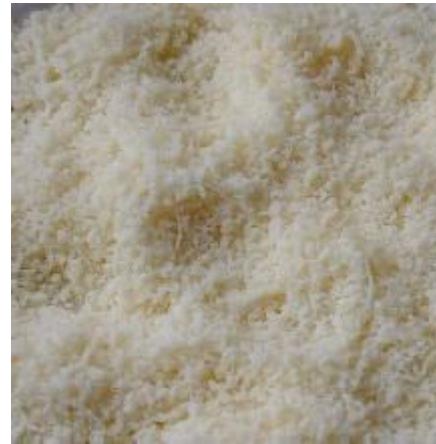
- Take the texture from one object and “paint” it onto another object
  - This requires separating texture and shape
  - That’s HARD, but we can cheat
  - Assume we can capture shape by boundary and rough shading
- Then, just add another constraint when sampling: similarity to underlying image at that spot





parmesan

+



=

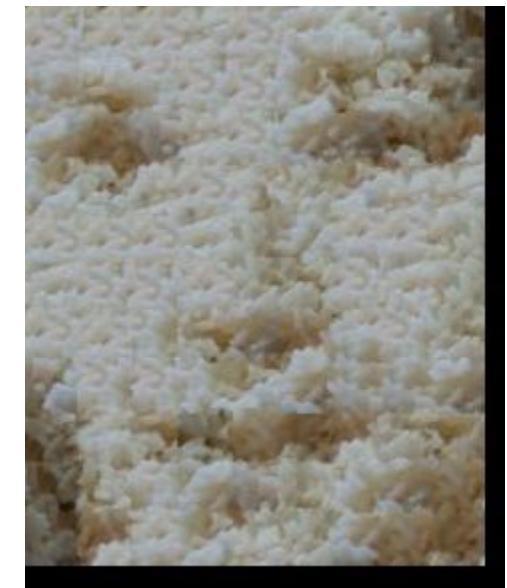


rice

+

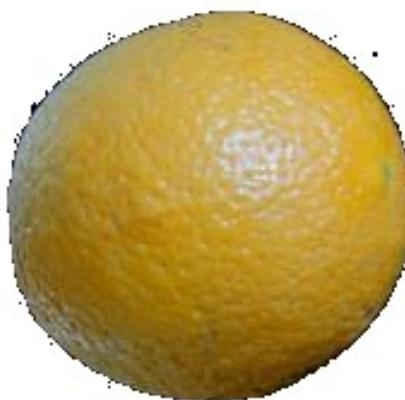


=

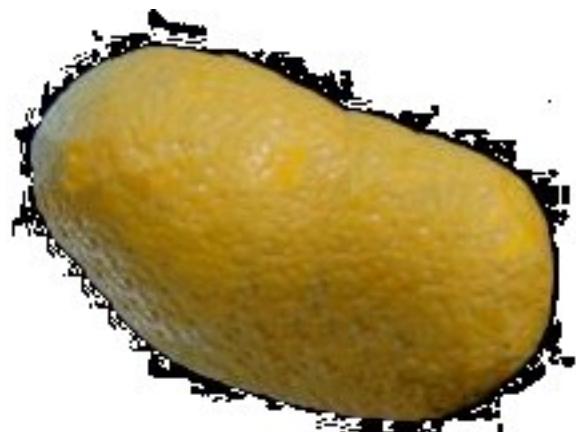


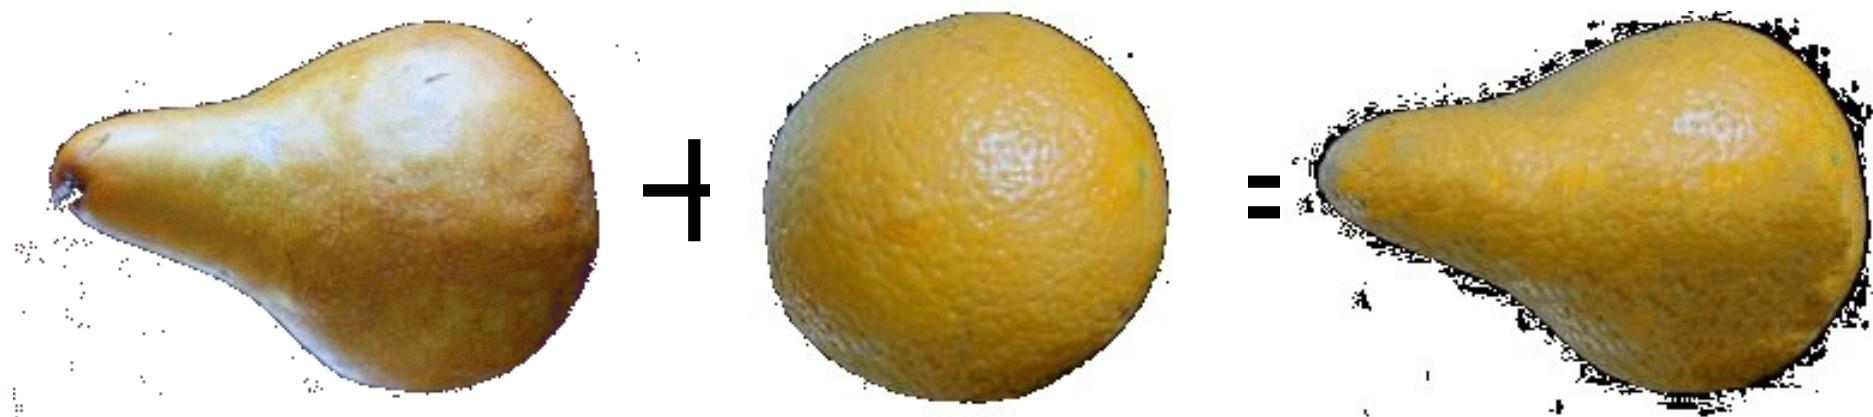


+



=





# (Manual) texture synthesis in the media



(WHITE HOUSE PHOTO/AP)

# (Manual) texture synthesis in the media



# WHATEVER IT TAKES



# Synthesizing textures when constructing 3d models of archaeological sites



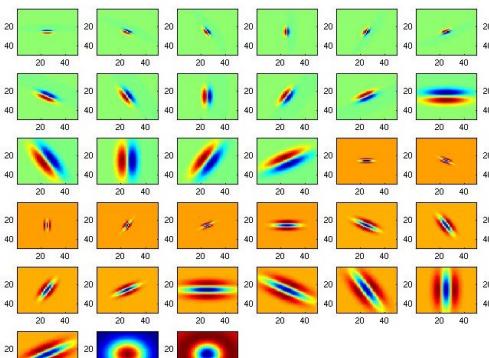
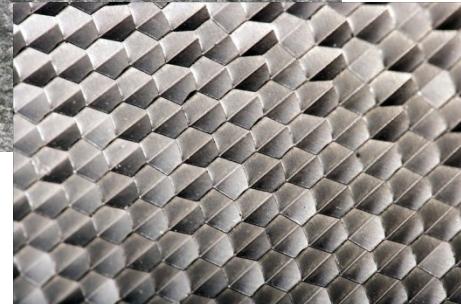
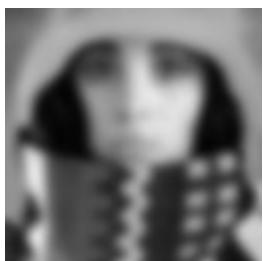
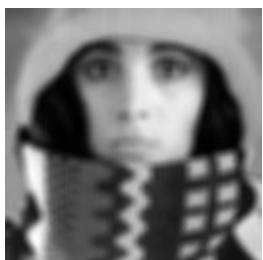
Figure 12. The Nymphaeum at the upper agora of Sagalassos with differently textured pillars. Overview of one half of the building (symmetric)



Figure 14. Nymphaeum pillars and back wall fragments in detail

A. Zalesny et al., Realistic Textures for Virtual Anastylosis

# So far: features and filters



# Transforming and describing images; textures, colors, edges