

Memory Coalescing Techniques

- 1 Accessing Global and Shared Memory
 - memory coalescing to global memory
 - avoiding bank conflicts in shared memory
- 2 Memory Coalescing Techniques
 - accessing global memory for a matrix
 - using shared memory for coalescing
- 3 Avoiding Bank Conflicts
 - computing consecutive powers

MCS 572 Lecture 34
Introduction to Supercomputing
Jan Verschelde, 9 April 2014

Memory Coalescing Techniques

1 Accessing Global and Shared Memory

- memory coalescing to global memory
- avoiding bank conflicts in shared memory

2 Memory Coalescing Techniques

- accessing global memory for a matrix
- using shared memory for coalescing

3 Avoiding Bank Conflicts

- computing consecutive powers

dynamic random access memories (DRAMs)

Accessing data in the global memory is critical to the performance of a CUDA application.

In addition to tiling techniques utilizing shared memories we discuss memory coalescing techniques to move data efficiently from global memory into shared memory and registers.

Global memory is implemented with dynamic random access memories (DRAMs). Reading one DRAM is a very slow process.

Modern DRAMs use a parallel process:

Each time a location is accessed, many consecutive locations that includes the requested location are accessed.

If an application uses data from consecutive locations before moving on to other locations, the DRAMs work close to the advertised peak global memory bandwidth.

memory coalescing

Recall that all threads in a warp execute the same instruction.

When all threads in a warp execute a load instruction, the hardware detects whether the threads access consecutive memory locations.

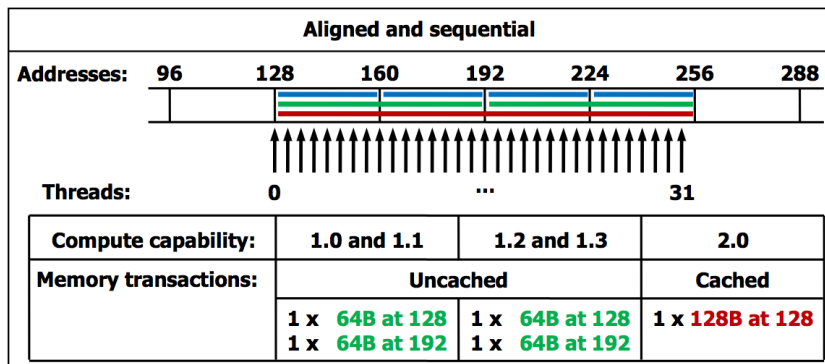
The most favorable global memory access is achieved when the same instruction for all threads in a warp accesses global memory locations.

In this favorable case, the hardware *coalesces* all memory accesses into a consolidated access to consecutive DRAM locations.

If thread 0 accesses location n , thread 1 accesses location $n + 1$, ... thread 31 accesses location $n + 31$, then all these accesses are *coalesced*, that is: combined into one single access.

The CUDA C Best Practices Guide gives a high priority recommendation to coalesced access to global memory.

an example of a global memory access by a warp



from Figure G-1 of the *NVIDIA Programming Guide*.

alignment in memory

In `/usr/local/cuda/include/vector_types.h`
we find the definition of the type `double2` as

```
struct __device_builtin__ __builtin_align__(16) double2
{
    double x, y;
};
```

The `__align__(16)` causes the doubles in `double2` to be 16-byte or 128-bit aligned.

Using the `double2` type for the real and imaginary part of a complex number allows for coalesced memory access.

exploring the effects of misaligned memory access

With a simple copy kernel we can explore what happens when access to global memory is misaligned:

```
__global__ void copyKernel  
( float *output, float *input, int offset )  
{  
    int i = blockIdx.x*blockDim.x + threadIdx.x + offset;  
    output[i] = input[i];  
}
```

The bandwidth will decrease significantly for `offset > 1`.

Memory Coalescing Techniques

1 Accessing Global and Shared Memory

- memory coalescing to global memory
- avoiding bank conflicts in shared memory

2 Memory Coalescing Techniques

- accessing global memory for a matrix
- using shared memory for coalescing

3 Avoiding Bank Conflicts

- computing consecutive powers

shared memory and memory banks

Shared memory has 32 banks that are organized such that successive 32-bit words are assigned to successive banks, i.e.: interleaved.


The bandwidth of shared memory is 32 bits per bank per clock cycle. Because shared memory is on chip, uncached shared memory latency is roughly $100\times$ lower than global memory.

A *bank conflict* occurs if two or more threads access any bytes within *different* 32-bit words belonging to the *same* bank.

If two or more threads access any bytes within the same 32-bit word, then there is no bank conflict between these threads.


The CUDA C Best Practices Guide gives a medium priority recommendation to shared memory access without bank conflicts.

examples of strided shared memory accesses



from Figure G-2 of the *NVIDIA Programming Guide*.

irregular and colliding shared memory accesses



from Figure G-3 of the *NVIDIA Programming Guide*.

Memory Coalescing Techniques

1 Accessing Global and Shared Memory

- memory coalescing to global memory
- avoiding bank conflicts in shared memory

2 Memory Coalescing Techniques

- accessing global memory for a matrix
- using shared memory for coalescing



3 Avoiding Bank Conflicts

- computing consecutive powers

accessing the elements in a matrix

Consider two ways of accessing the elements in a matrix:

- 1 elements are accessed row after row; or
- 2 elements are accessed column after column.



linear address system

Consider a 4-by-4 matrix:

| | | | |
|-----------|-----------|-----------|-----------|
| $a_{0,0}$ | $a_{0,1}$ | $a_{0,2}$ | $a_{0,3}$ |
| $a_{1,0}$ | $a_{1,1}$ | $a_{1,2}$ | $a_{1,3}$ |
| $a_{2,0}$ | $a_{2,1}$ | $a_{2,2}$ | $a_{2,3}$ |
| $a_{3,0}$ | $a_{3,1}$ | $a_{3,2}$ | $a_{3,3}$ |



| | | | | | | | | | | | | | | | |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| $a_{0,0}$ | $a_{0,1}$ | $a_{0,2}$ | $a_{0,3}$ | $a_{1,0}$ | $a_{1,1}$ | $a_{1,2}$ | $a_{1,3}$ | $a_{2,0}$ | $a_{2,1}$ | $a_{2,2}$ | $a_{2,3}$ | $a_{3,0}$ | $a_{3,1}$ | $a_{3,2}$ | $a_{3,3}$ |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |


In C, the matrix is stored row wise as a one dimensional array.

first access

Threads t_0 , t_1 , t_2 , and t_3 access the elements on the first two columns:

| | | | |
|-----------|-----------|-----------|-----------|
| $a_{0,0}$ | $a_{0,1}$ | $a_{0,2}$ | $a_{0,3}$ |
| $a_{1,0}$ | $a_{1,1}$ | $a_{1,2}$ | $a_{1,3}$ |
| $a_{2,0}$ | $a_{2,1}$ | $a_{2,2}$ | $a_{2,3}$ |
| $a_{3,0}$ | $a_{3,1}$ | $a_{3,2}$ | $a_{3,3}$ |


Two vertical arrows point down from the first and second columns of the table.





second access

Four threads t_0 , t_1 , t_2 , and t_3 access elements on the first two rows:

| | | | | |
|-----------|-----------|-----------|-----------|---|
| $a_{0,0}$ | $a_{0,1}$ | $a_{0,2}$ | $a_{0,3}$ | → |
| $a_{1,0}$ | $a_{1,1}$ | $a_{1,2}$ | $a_{1,3}$ | → |
| $a_{2,0}$ | $a_{2,1}$ | $a_{2,2}$ | $a_{2,3}$ | |
| $a_{3,0}$ | $a_{3,1}$ | $a_{3,2}$ | $a_{3,3}$ | |



uncoalesced versus coalesced access



Memory Coalescing Techniques

1 Accessing Global and Shared Memory

- memory coalescing to global memory
- avoiding bank conflicts in shared memory

2 Memory Coalescing Techniques

- accessing global memory for a matrix
- using shared memory for coalescing


3 Avoiding Bank Conflicts

- computing consecutive powers

tiled matrix-matrix multiplication

$$C_{i,j} = \sum_{k=1}^{m/w} A_{i,k} \cdot B_{k,j}$$


Diagram illustrating the structure of matrix B . The matrix is partitioned into a vertical tile of width w and height m . The tile is labeled B and the column index is j . The height m is indicated by a vertical double-headed arrow on the right.



tiling matrix multiplication with shared memory

For $C_{i,j} = \sum_{k=1}^{m/w} A_{i,k} \cdot B_{k,j}$, $A \in \mathbb{R}^{n \times m}$, $B \in \mathbb{R}^{m \times p}$, $A_{i,k}, B_{k,j}, C_{i,j} \in \mathbb{R}^{w \times w}$,

every warp reads one tile $A_{i,k}$ of A and one tile $B_{k,j}$ of B : every thread in the warp reads one element of $A_{i,k}$ and one element of $B_{k,j}$.

The number of threads equals w , the width of one tile, and threads are identified with `tx = threadIdx.x` and `ty = threadIdx.y`.

The `by = blockIdx.y` and `bx = blockIdx.x` correspond respectively to the first and the second index of each tile, so we have `row = by*w + ty` and `col = bx*w + tx`.

Row wise access to A uses $A[\text{row} * m + (k * w + tx)]$. For B :
 $B[(k * w + ty) * m + col] = B[(k * w + ty) * m + bx * w + tx]$.

Adjacent threads in a warp have adjacent `tx` values so we have coalesced access also to B .

tilted matrix multiplication kernel

```
__global__ void mul ( float *A, float *B, float *C, int m )
{
    __shared__ float As[w][w];
    __shared__ float Bs[w][w];
    int bx = blockIdx.x;          int by = blockIdx.y;
    int tx = threadIdx.x;        int ty = threadIdx.y;
    int col = bx*w + tx;         int row = by*w + ty;
    float Cv = 0.0;
    for(int k=0; k<m/w; k++)
    {
        As[ty][tx] = A[row*m + (k*w + tx)];
        Bs[ty][tx] = B[(k*w + ty)*m + col];
        __syncthreads();
        for(int ell=0; ell<w; ell++)
            Cv += As[ty][ell]*Bs[ell][tx];
        C[row][col] = Cv;
    }
}
```

Memory Coalescing Techniques

1 Accessing Global and Shared Memory

- memory coalescing to global memory
- avoiding bank conflicts in shared memory

2 Memory Coalescing Techniques

- accessing global memory for a matrix
- using shared memory for coalescing

3 Avoiding Bank Conflicts

- computing consecutive powers

consecutive powers

Consider the following problem:

Input : $x_0, x_1, x_2, \dots, x_{31}$, all of type float.

Output : $x_0^2, x_0^3, x_0^4, \dots, x_0^{33}, x_1^2, x_1^3, x_1^4, \dots, x_1^{33}, x_2^2, x_2^3, x_2^4, \dots, x_2^{33},$
 $\dots, x_{31}^2, x_{31}^3, x_{31}^4, \dots, x_{31}^{33}.$

This gives 32 threads in a warp 1,024 multiplications to do.

Assume the input and output resides in shared memory.

How to compute without bank conflicts?

writing with stride

Observe the order of the output sequence:

Input : $x_0, x_1, x_2, \dots, x_{31}$, all of type float.

Output : $x_0^2, x_0^3, x_0^4, \dots, x_0^{33}, x_1^2, x_1^3, x_1^4, \dots, x_1^{33}, x_2^2, x_2^3, x_2^4, \dots, x_2^{33},$
 $\dots, x_{31}^2, x_{31}^3, x_{31}^4, \dots, x_{31}^{33}.$

If thread i computes $x_i^2, x_i^3, x_i^4, \dots, x_i^{33}$, then after the first step, all threads write $x_0^2, x_1^2, x_2^2, \dots, x_{31}^2$ to shared memory.

If the stride is 32, all threads write into the same bank.

Instead of a simultaneous computation of 32 powers at once, the writing to shared memory will be serialized.

changed order of storage

If we alter the order in the output sequence:

Input : $x_0, x_1, x_2, \dots, x_{31}$, all of type float.

Output : $x_0^2, x_1^2, x_1^2, \dots, x_{31}^2, x_0^3, x_1^3, x_2^3, \dots, x_{31}^3, x_0^4, x_1^4, x_2^4, \dots, x_{31}^4,$
 $\dots, x_0^{33}, x_1^{33}, x_2^{33}, \dots, x_{31}^{33}.$

After the first step, thread i writes x_i^2 in adjacent memory, next to x_{i-1}^2 (if $i > 0$) and x_{i+1}^2 (if $i < 31$).

Without bank conflicts, the speedup will be close to 32.

summary and exercises

We covered §6.2 in the book of Kirk & Hwu.

- 1 Run `copyKernel` for large enough arrays for zero `offset` and an `offset` equal to two. Measure the timings and deduce the differences in memory bandwidth between the two different values for `offset`.
- 2 Consider the kernel of `matrixMul` in the GPU computing SDK. Is the loading of the tiles into shared memory coalesced? Justify your answer.
- 3 Write a CUDA program for the computation of consecutive powers, using coalesced access of the values for the input elements. Compare the two orders of storing the output sequence in shared memory: once with and once without bank conflicts.